

Cycle 2: Enterprise Networking

CS 436: Fall 2017


Matthew Caesar

<http://www.cs.illinois.edu/~caesar/cs436>


Lab 1 Recap

- Skills acquired: Network configuration, Layer 3 network administration, protocols (BGP, OSPF), Cisco/Cisco IOS
- Demos today and Wednesday – Sign up if you have not
- How to go deeper:
 - GNS3: play around with more protocols, vendors
 - Lurk at <https://www.reddit.com/r/networking/>
 - CCNA/CCIE certification

Conceptual network architecture (TCP/IP vs ISO)

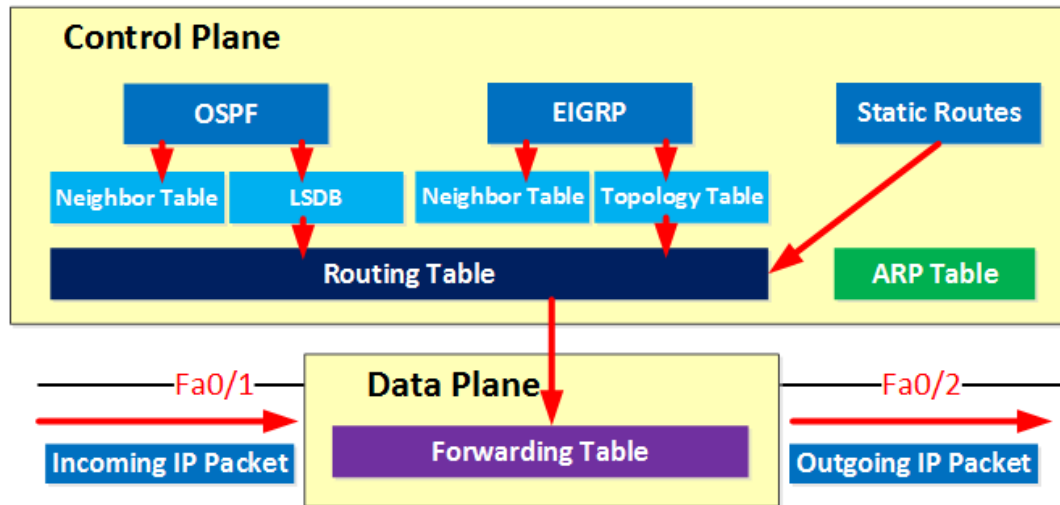


TCP/IP protocol stack



ISO/OSI layer

Data vs Control planes




- Two kinds of packets in the internet:
- **Data packets** – traffic computers send
 - Goes through the router, not to the router
 - Video streaming, VoIP, downloading web pages, bittorrent, etc.
- **Control packets** – packets used to “control” the internet
 - Destined to or originated from router itself
 - Routing protocol messages (BGP, OSPF, etc), session setup messages
- Can think of “planes” associated with each of these functions
 - Control vs data plane
 - Layer 2 control plane more plug-and-play, can focus more on data plane in Lab 2

Lab 2 Overview

- This lab is less about configuration
 - Layer 2 is “plug and play”!
- Focus on understanding traffic and network inputs
 - Netflow (large-scale packet monitoring)
 - Tcpdump/pcap (edge packet monitoring)
 - Routing traces (BGP feeds)

Lab 2 Overview: Netflow



Source IP	Destination IP	Application	Source Port	Dest . Port	Protocol	DSCP Name	TCP FLAGS	Traffic	No of Packets	NextHop
59.93.161.133	203.199.206.240	TCP_App	3593	2967	TCP	Default	S	2.73 KB	57	203.199.206.240
59.93.161.133	203.199.206.192	TCP_App	3543	2967	TCP	Default	S	2.73 KB	57	203.199.206.192
59.93.161.133	203.199.206.208	TCP_App	3560	2967	TCP	Default	S	2.73 KB	57	203.199.206.208
59.93.161.133	203.199.206.224	TCP_App	3577	2967	TCP	Default	S	2.73 KB	57	203.199.206.224
59.93.161.133	203.199.206.243	TCP_App	3596	2967	TCP	Default	S	2.73 KB	57	203.199.206.243
59.93.161.133	203.199.206.195	TCP_App	3546	2967	TCP	Default	S	2.73 KB	57	203.199.206.195

Cycle 2: Enterprise Networking

CS 436: Fall 2017

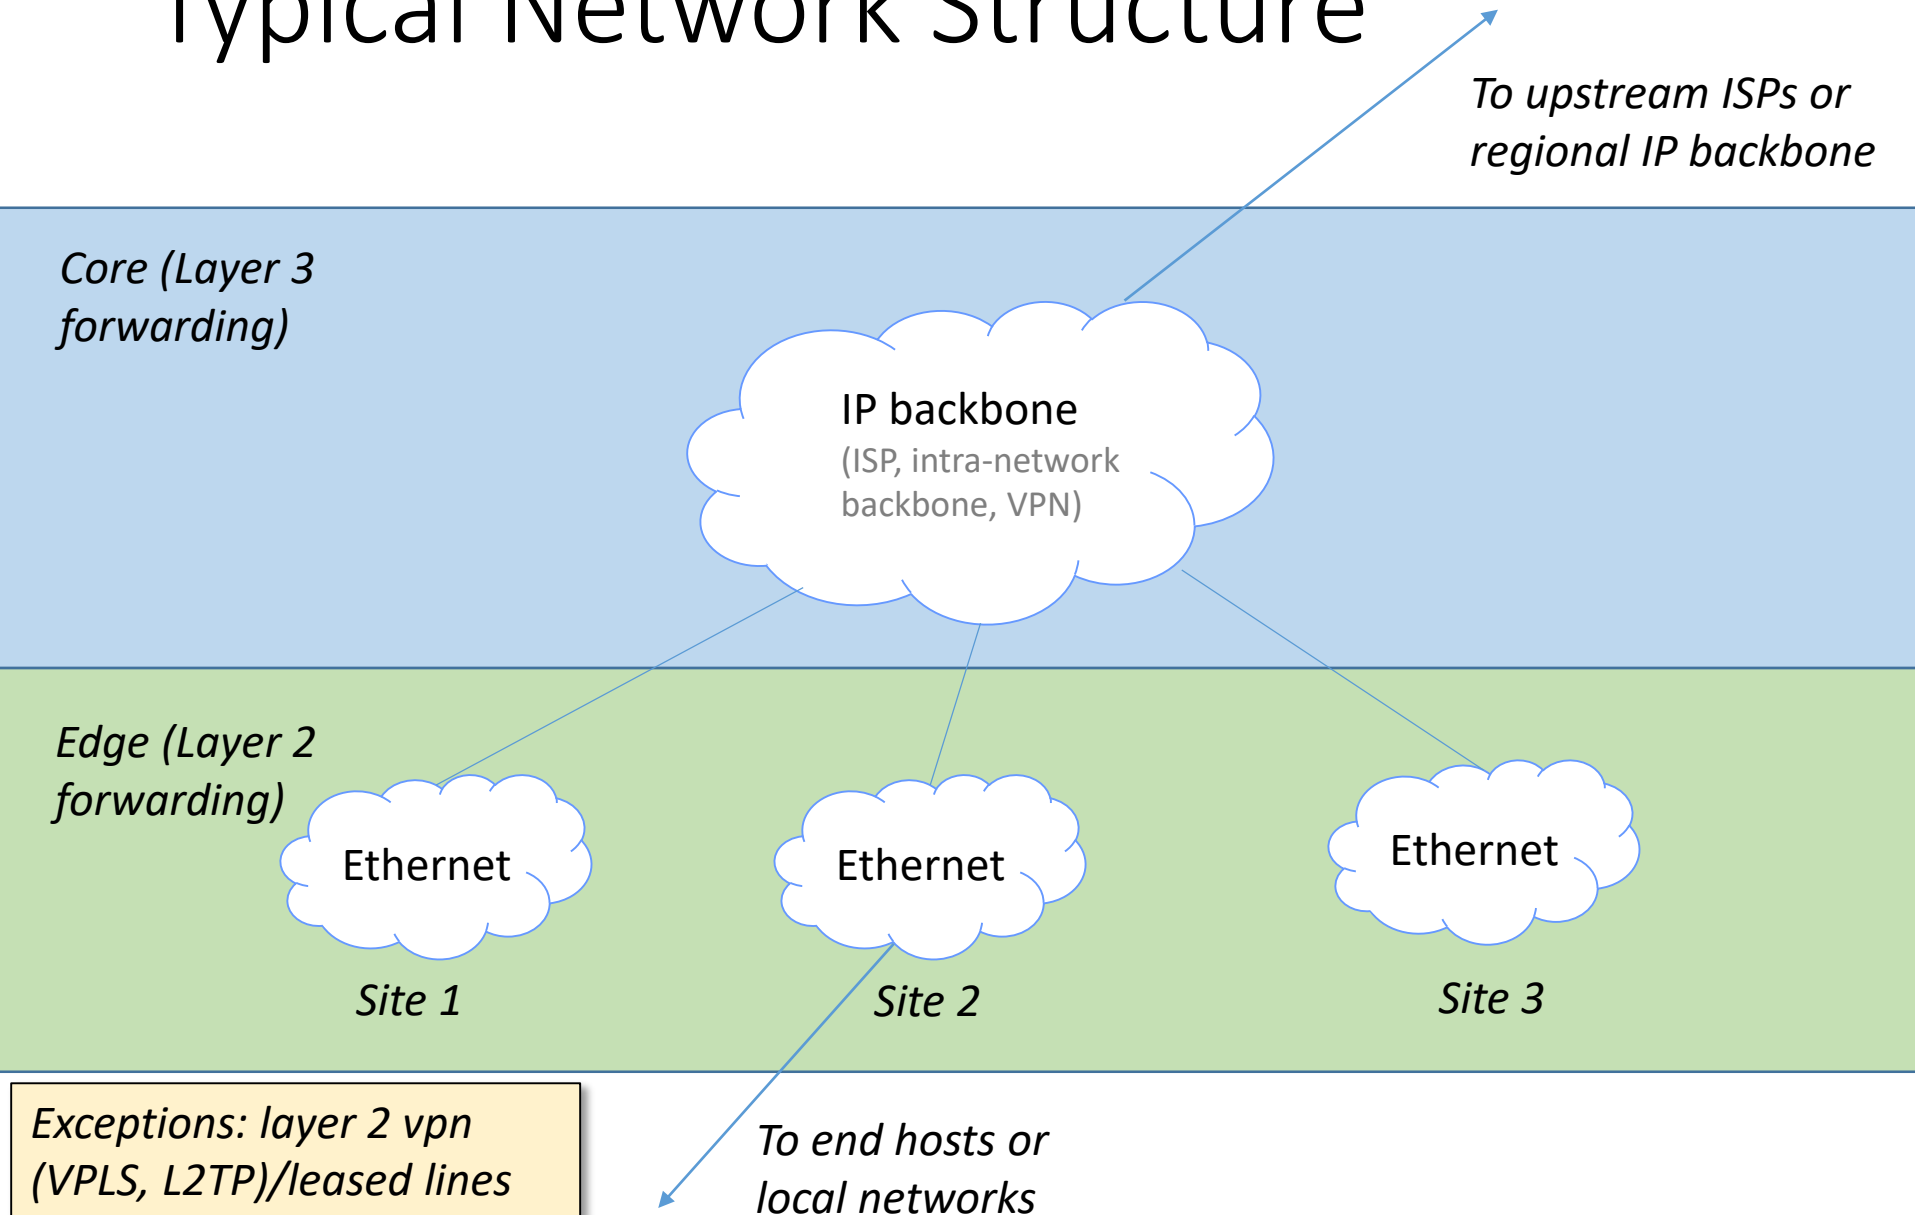
Matthew Caesar

<http://www.cs.illinois.edu/~caesar/cs436>

Motivating Questions

- We are running out of address space at our Miami branch office. What are our options?
- New legislation requires companies storing PCI data to log IP sources of all accesses – how to do that?
- Judge requires you to start sending warning notices to people bittorrenting TV shows on your network – how to do that?
- Here is a budget of \$100,000. Build a network for our Chicago branch office.

Typical Network Structure



Ethernet (layer 2) vs IP (layer 3) routing

- Ethernet is “plug and play”
 - Easy to build networks
 - May optionally configure ACLs, SSIDs (wireless), spanning tree properties, etc.
- Each host assigned a topology-independent *MAC address*
 - E.g., 00-14-22-01-23-45
- Uses “dumb” flooding (broadcast) to get data packets where they need to go
 - Less efficient than link-state (unicast)

C:\Users\mccae>ipconfig /all

Windows IP Configuration

Host Name : LAPTOP-M33LKCP0
 Primary Dns Suffix :
 Node Type : Hybrid
 IP Routing Enabled. : No
 WINS Proxy Enabled. : No
 DNS Suffix Search List. : SJC-WIFI-DHCP1

Ethernet adapter Ethernet:

Media State : Media disconnected
 Connection-specific DNS Suffix . :
 Description : Intel(R) Ethernet Connection (4) I219-LM
 Physical Address. : 54-EE-75-DB-6D-44
 DHCP Enabled. : Yes
 Autoconfiguration Enabled : Yes

Wireless LAN adapter Local Area Connection* 1:

Media State : Media disconnected
 Connection-specific DNS Suffix . :
 Description : Microsoft Wi-Fi Direct Virtual Adapter
 Physical Address. : 1C-4D-70-72-4F-12
 DHCP Enabled. : Yes
 Autoconfiguration Enabled : Yes

MAC address

Automatically discover host network configuration

Wireless LAN adapter Wi-Fi:

Connection-specific DNS Suffix . : SJC-WIFI-DHCP1
 Description : Intel(R) Dual Band Wireless-AC 8265
 Physical Address. : 1C-4D-70-72-4F-11
 DHCP Enabled. : Yes
 Autoconfiguration Enabled : Yes
 Link-local IPv6 Address : fe80::3df4:a06e:801d:3b1%15(Preferred)
 IPv4 Address. : 10.47.13.145(Preferred)
 Subnet Mask : 255.255.128.0
 Lease Obtained. : Thursday, September 14, 2017 6:35:18 AM
 Lease Expires : Thursday, September 14, 2017 9:42:30 AM
 Default Gateway : 10.47.1.1
 DHCP Server : 10.47.1.100
 DHCPv6 IAID : 85740912
 DHCPv6 Client DUID. : 00-01-00-01-21-2B-72-CC-54-EE-75-DB-6D-44
 DNS Servers : 8.8.8.8
 8.8.4.4
 NetBIOS over Tcpip. : Enabled

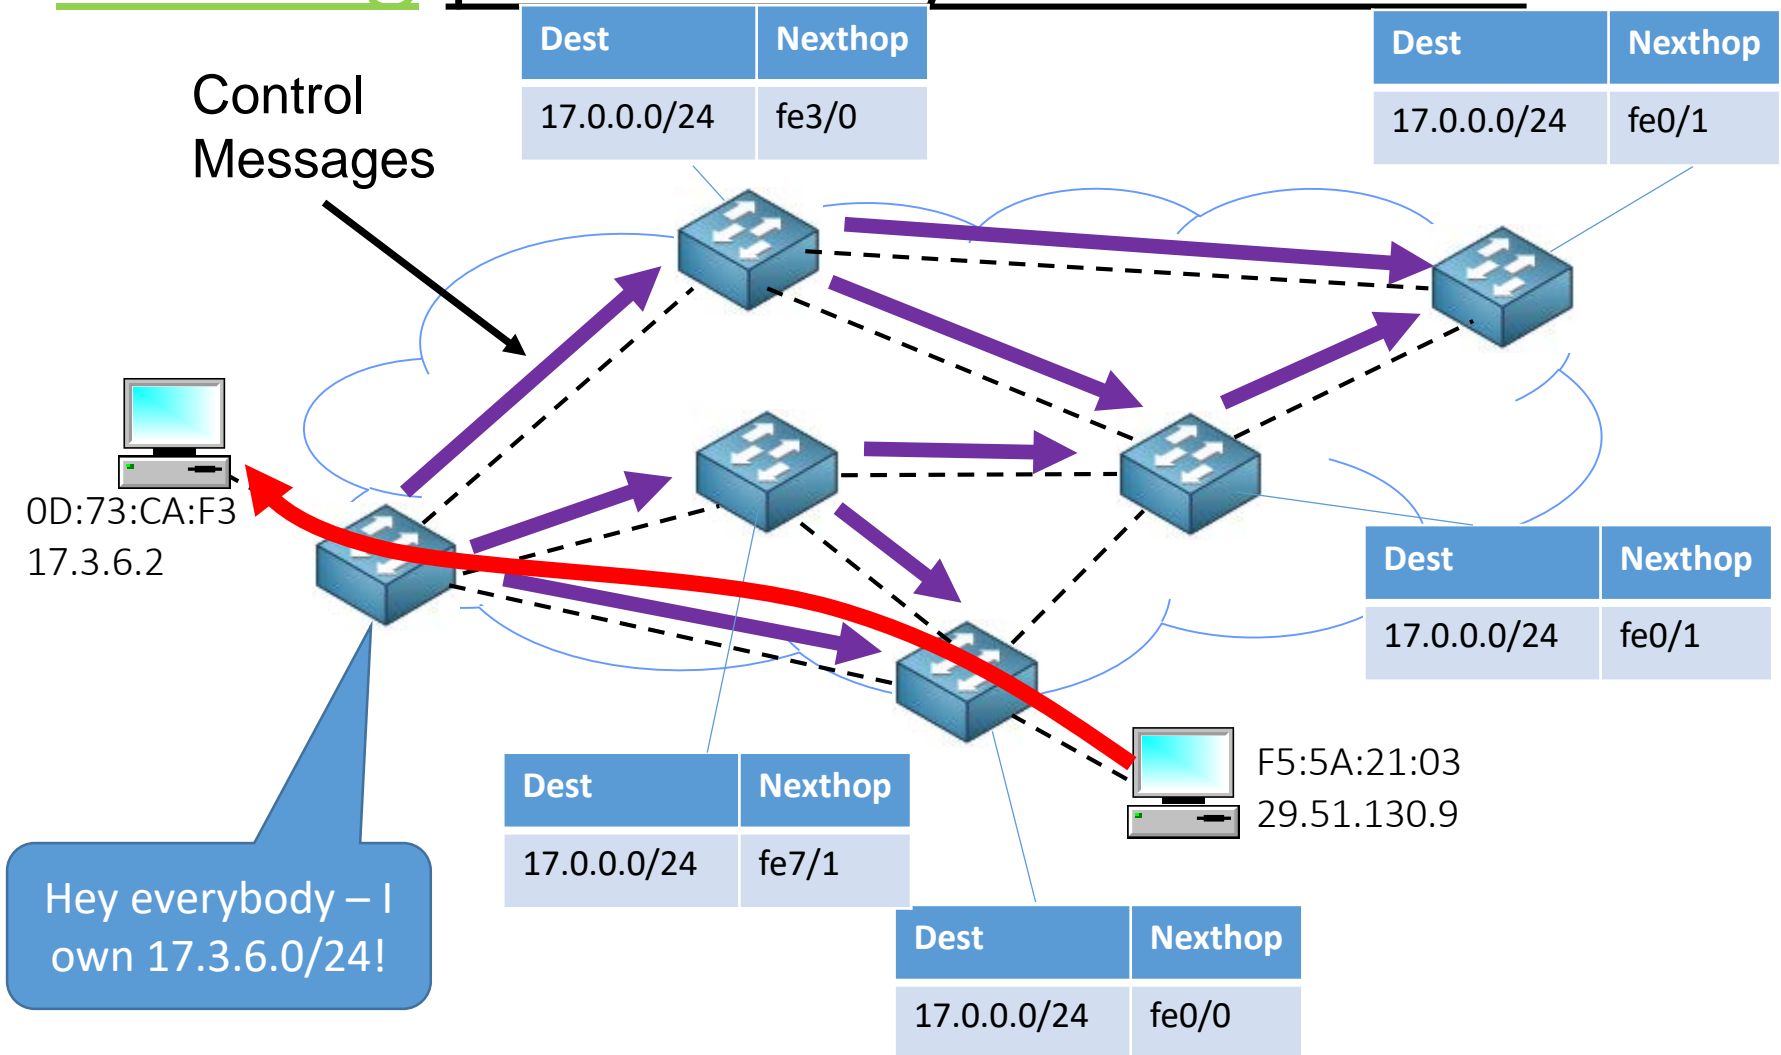
Which IP addresses are on my subnet

Where to forward packets if destination is off my subnet

Ethernet adapter Bluetooth Network Connection:

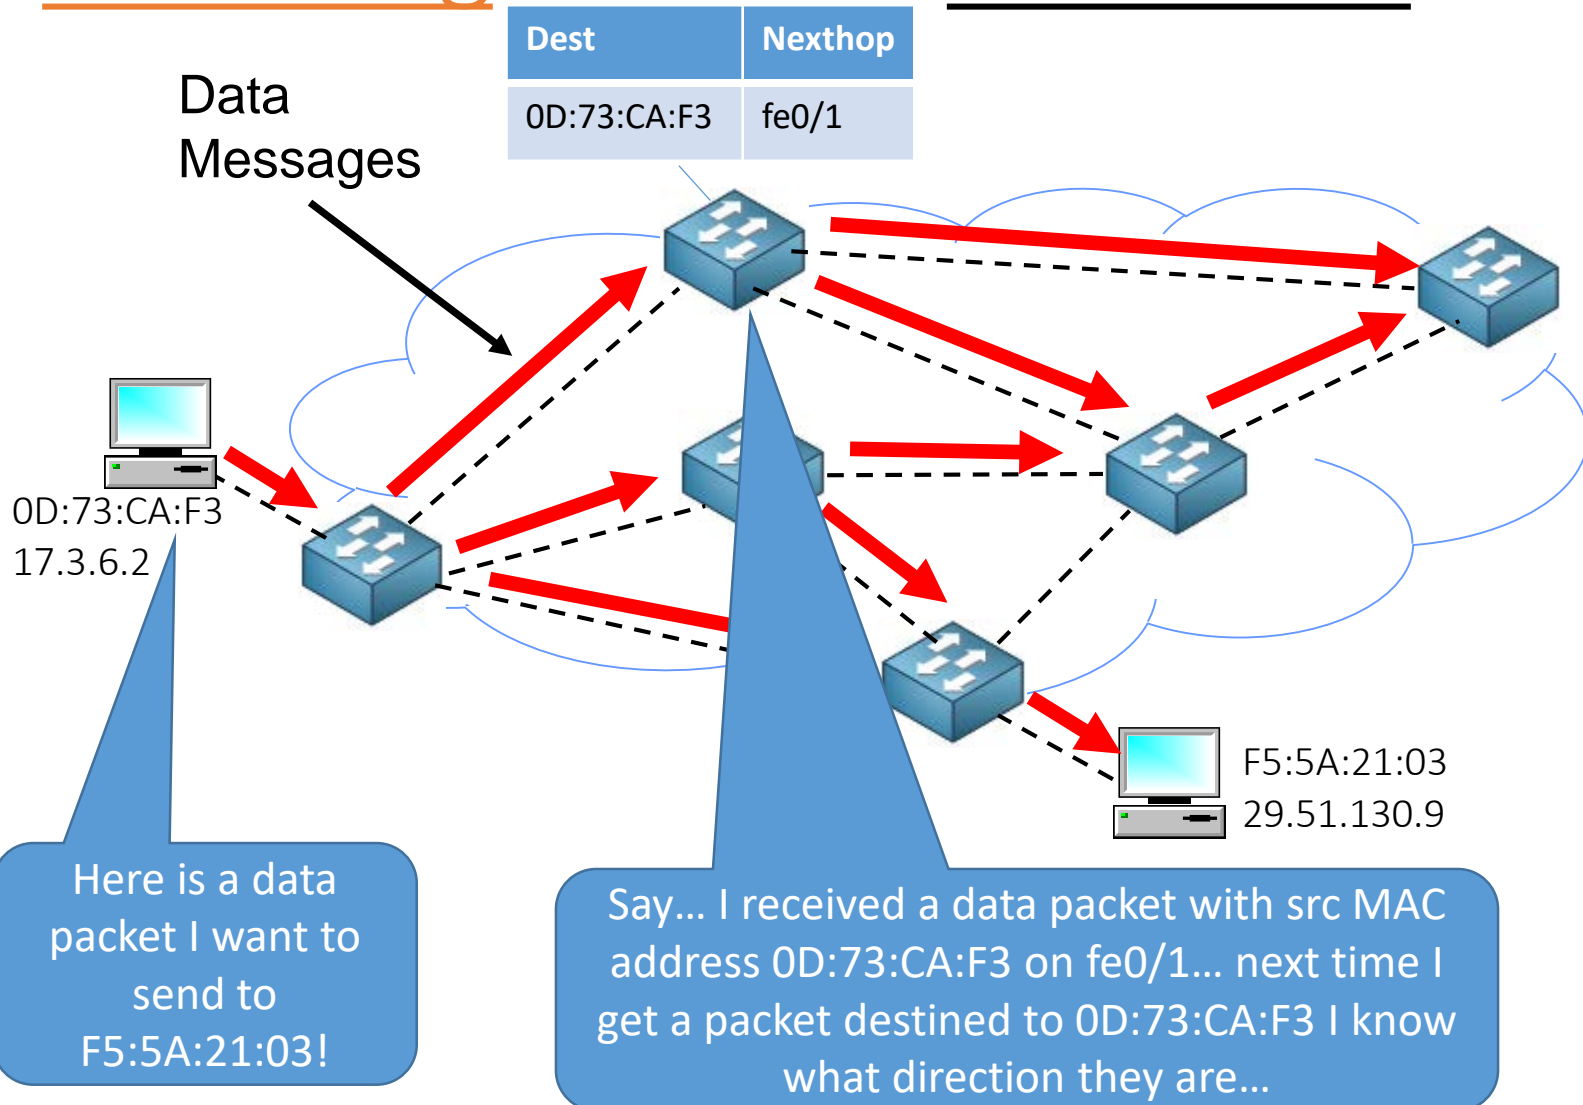
L2 Switching vs L3 Routing:

Routing proactively builds state

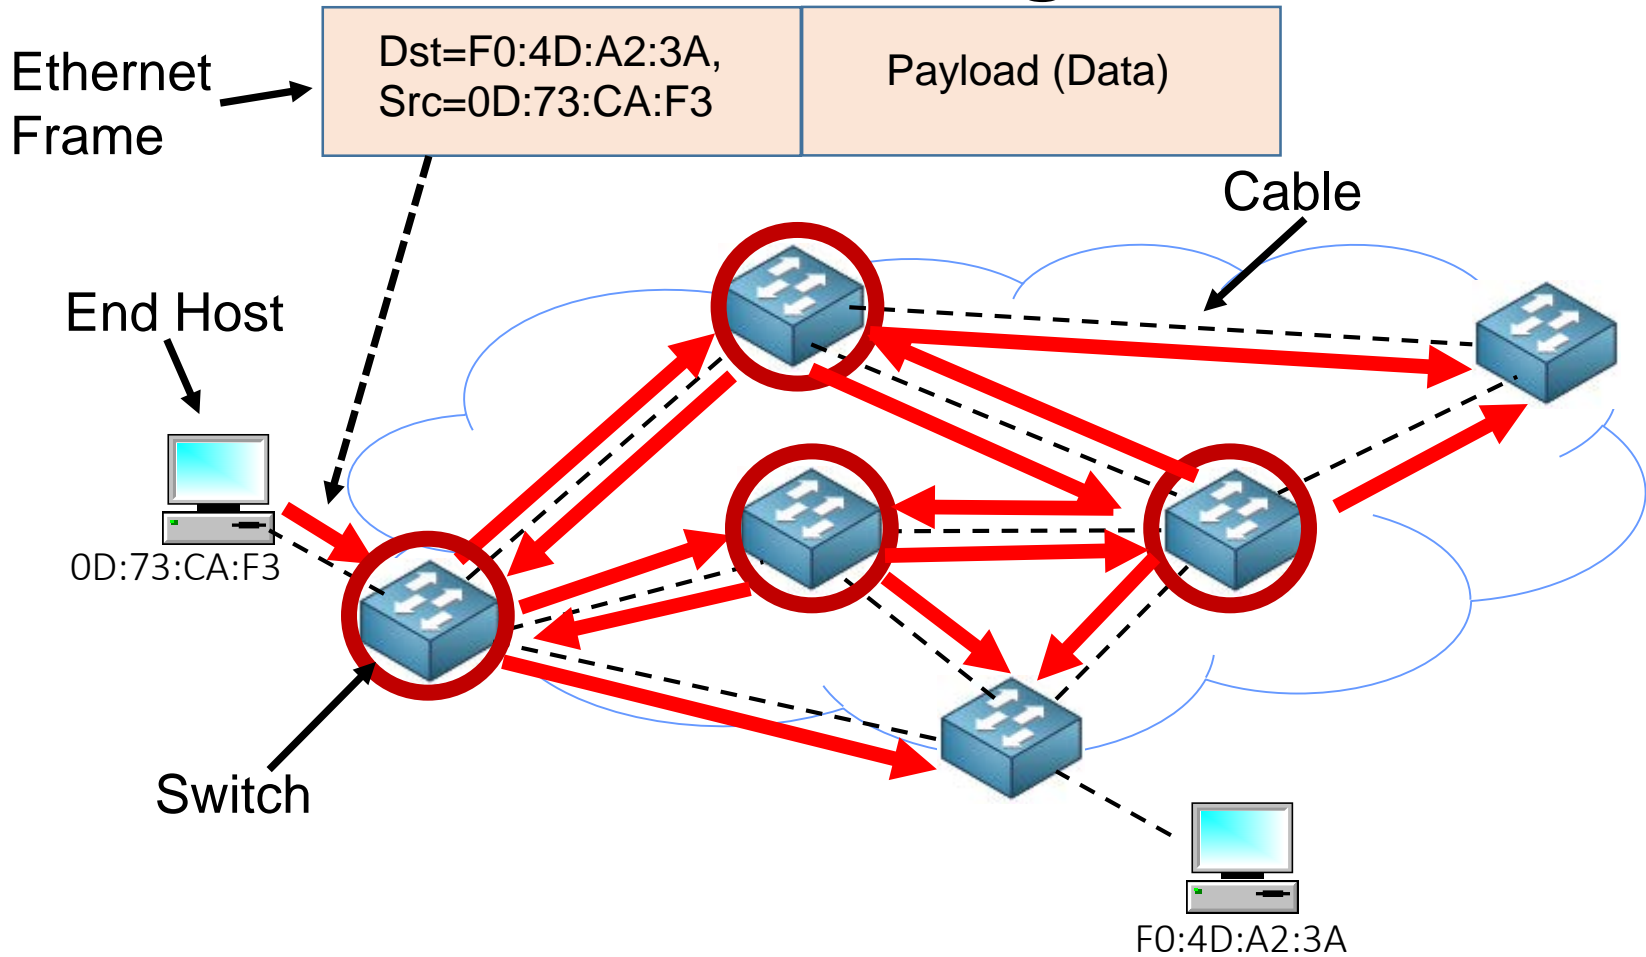


L2 Switching vs L3 Routing:

Switching relies on broadcast




Ethernet Forwarding



- Problem: Broadcast Storms
- How to flood with stateless switches?


Ethernet Forwarding



- **Solution: Construct a Spanning Tree**
 - Elect a “root” switch
 - Root-facing ports are active, others disabled
 - Improvement: Per-VLAN spanning trees


Avoiding Flooding

- Flooding packets throughout network introduces problems
 - Scalability, privacy, resource isolation, lack of access control
- Scalability requirement is growing very fast
 - Large enterprises: 50k end hosts
 - Data centers: 100k servers, 5k switches
 - Metro-area Ethernet: over 1M subscribers




Avoiding Flooding

- Suppose source sends a frame to a destination
 - Which LANs should a frame be forwarded on?
- Trivial algorithm
 - Forward all frames on all (other) LAN's
 - Potentially heavy traffic and processing overhead
- Optimize by using address information
 - “Learn” which hosts live on which LAN
 - Maintain forwarding table
 - Only forward when necessary
 - Reduces bridge workload



Learning Bridges

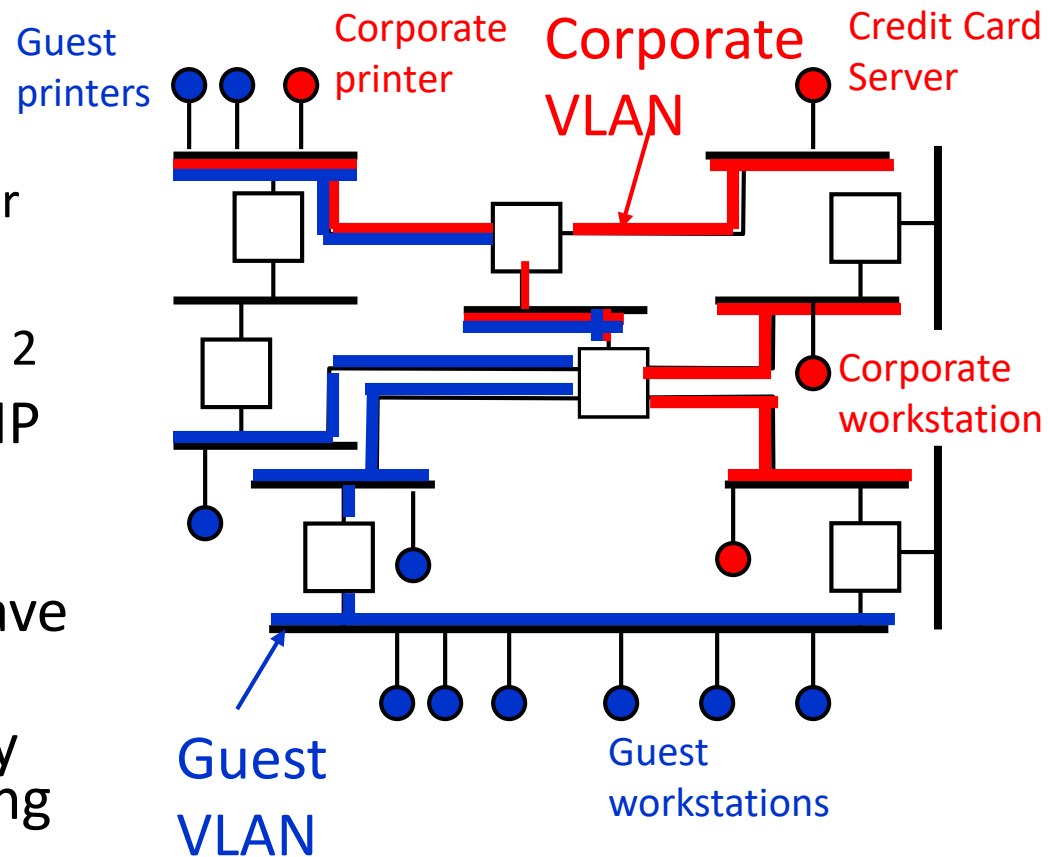
- Bridge learns table entries based on source address
 - When receive frame from A on port 1
add A to list of hosts on port 1
 - Time out entries to allow movement of hosts
- Table is an “optimization”, meaning it helps performance but is not mandatory
- Always forward broadcast frames



Host	Port
A	1
B	1
C	1
X	2
Y	2
Z	2

Scaling Ethernet with VLANs


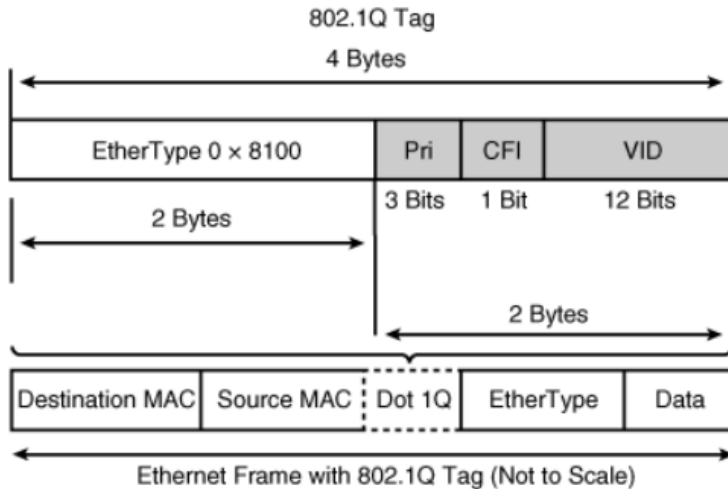
- Divide up hosts into logical groups called **VLANs**
 - Like virtual machines, but for LANs (creates “virtual networks”)
 - VLANs isolate traffic at layer 2
- Each VLAN corresponds to IP subnet, single broadcast domain
- Ethernet packet headers have VLAN tag
- Bridges forward packet only on subnets on corresponding VLAN



Virtual LANs



- Downsides of VLANs
 - Are (usually) manually configured, complicates network management
 - Hard to seamlessly migrate across VLAN boundaries due to addressing restrictions
- Upsides of VLANs
 - Limits scope of broadcasts
 - Logical separation improves isolation, security
 - Can change virtual topology without changing physical topology
 - E.g., used in data centers for VM migration

How VLANs are implemented




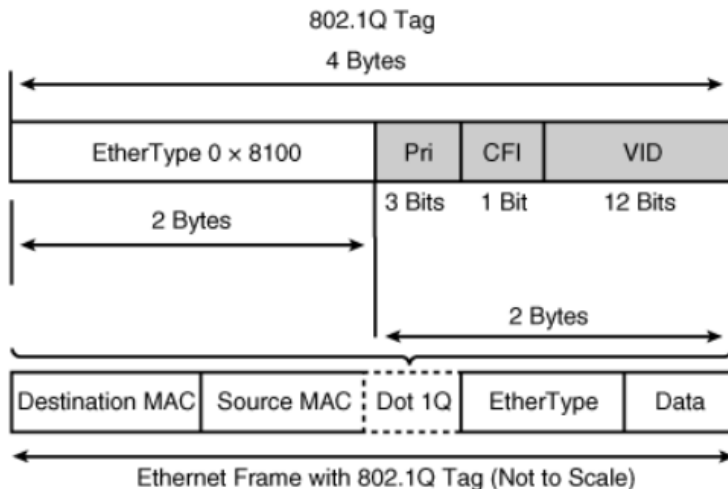
- Packets are annotated with 12-bit **VLAN tags**
 - Up to 4096 VLANs can be encapsulated within a single VLAN ID
- LAN switches can configure ports as access ports or trunk ports
 - **Access ports** append tags on packets
 - VLAN membership almost always statically encoded in access switch's configuration file
 - **Trunk ports** can multiplex several VLANs

How VLANs are implemented



- 802.1Q (VLAN spec) defines a few other fields too
 - **Ethertype** of 0x8100 instructs switch to decode next 2 bytes as VLAN header
 - 3 bits of priority (like IP ToS)
 - 1 bit for compatibility with token ring
- What if 4096 VLANs isn't enough?
 - **QinQ** (802.1ad) – can encapsulate VLANs within VLANs by stacking VLAN tags
 - Up to 4096 VLANs can be multiplexed within a single VLAN ID → 4096^2 combinations

How VLANs are implemented




- Native mode

- IEEE likes to make specs that are backwards compatible
- 802.1Q allows trunk ports to carry both tagged and untagged frames
- Frames with no tags are said to be part of the switch's native VLAN

Dynamic Trunking Protocol

Ok.

Can we create
VLAN 7?

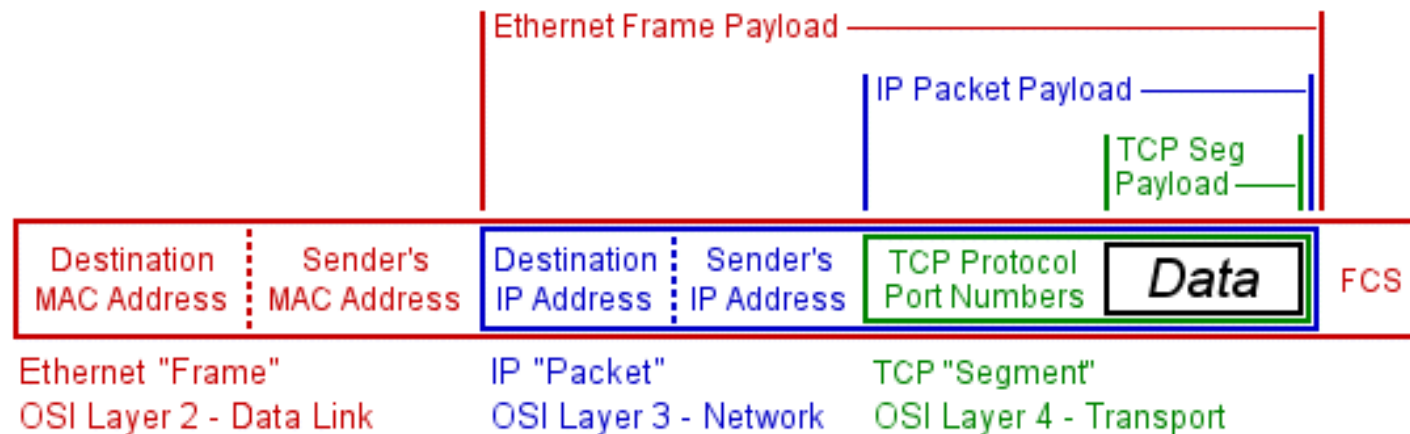


- Protocol to automate certain aspects of VLAN configuration
 - Determines whether two connected switches want to create a trunk
 - Automatically sets parameters such as encapsulation and VLAN range
- DTP transitions port through a set of states
 - Auto (port is willing to be trunked), On/Off (permanently forces link into/from trunking, even if neighbor disagrees), Desirable (attempts to make port a trunk; pursues agreement with neighbor)

Traffic

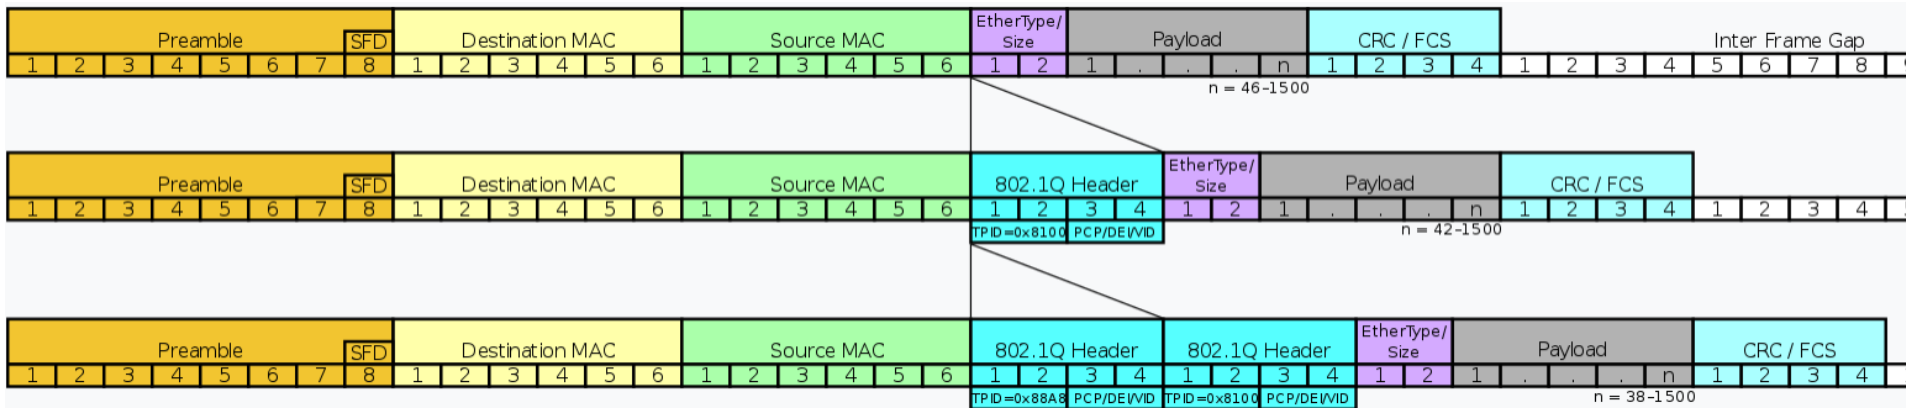
Encapsulation

Encapsulation Payloads



- Each layer of protocol stack encapsulates data passed to it.
- Each forwarding layer inspects data only at that encapsulation layer
 - Switching only looks at Ethernet header, Routing only looks at IP header, etc.
 - Terminology: "Layer-3 switch", "Layer-4 load balancer", "Layer-7 load balancer"

Ethernet Header



- Preamble – 56-bit pattern used to sync clocks
- Protocol ID: set to 0x8100 to identify frame as a 802.1Q-tagged frame
- Priority (PCP): 1 (background), 0 (best-effort, default), 2 (excellent effort), 3 (critical application), 4 (video), 5 (voice), 6 (internetwork control), 7 (network control)
- DEI – indicates if frame can be dropped during congestion
- VID – VLAN identifier. 12 bits allows for 4096 VLANs.
 - 802.1ad allows double tagging; 24 bits for VLANs

IPv4 Header

Offsets	Octet	0								1								2								3							
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31
0	0	Version				IHL				DSCP						ECN		Total Length															
4	32	Identification															Flags		Fragment Offset														
8	64	Time To Live								Protocol							Header Checksum																
12	96	Source IP Address																															
16	128	Destination IP Address																															
20	160	Options (if IHL > 5)																															
24	192																																
28	224																																
32	256																																

- Version: 4; IHL: specifies size of header (words)
- DSCP: Diffserv marking (priority/QoS)
- Identification/Frag offset: fragment counter/offset. Flags: don't fragment/more fragments/fragment ok
- TTL: hop counter; Protocol: TCP, ICMP, UDP, etc.
- Options: source routing, record route, etc. typically filtered.

TCP header

Offsets	Octet	0								1								2								3																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																	
Octet	Bit	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18	19	20	21	22	23	24	25	26	27	28	29	30	31																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																										
0	0	Source port																Destination port																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																									
4	32	Sequence number																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																									
8	64	Acknowledgment number (if ACK set)																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																									
12	96	Data offset				Reserved 0 0 0			N	C	E	U	A	P	R	S	F	Window Size																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																									
																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																																											</

- Ports: assigned to application. Local may be ephemeral.
- Seq num: used for reassembly; Initial seq num.
- Ack num: next seqno the sender is expecting
- Data offset: size of TCP header
- Flags: SYN, ACK, FIN are main used ones
- Window size: size of sender's receive window
- Options: Selective Acks Permitted, Max Segment Size announcement, etc

Addressing

Medium Access Control Address

- MAC address
 - Numerical address associated with an adapter
 - Flat name space of 48 bits (e.g., 00-15-C5-49-04-A9 in HEX)
 - Unique, hard-coded in the adapter when it is built
- Hierarchical Allocation
 - **Blocks**: assigned to vendors (e.g., Dell) by the IEEE
 - First 24 bits (e.g., 00-15-C5-**-**-**)
 - **Adapter**: assigned by the vendor from its block
 - Last 24 bits
- Broadcast address (FF-FF-FF-FF-FF-FF)
 - Send the frame to *all* adapters

MAC Address vs. IP Address

- MAC addresses (used in link-layer)
 - **Hard-coded** (often) in read-only memory when adapter is built
 - Like a social security number
 - **Flat** name space of 48 bits (e.g., 00-0E-9B-6E-49-76)
 - Portable, and can stay the same as the host moves
 - Used to get packet between interfaces on same network
- IP addresses
 - **Configured**, or learned dynamically
 - Like a postal mailing address
 - **Hierarchical** name space of 32 bits (e.g., 12.178.66.9)
 - Not portable, and depends on where the host is attached
 - Used to get a packet to destination IP subnet

Naming


- Application layer: URLs and domain names
 - names “resources” -- hosts, content, program
 - *(recall: mixes the what and where of an object)*
- Network layer: IP addresses
 - host’s network location
- Link layer: MAC addresses
 - host identifier
- Use all three for end-to-end communication!

Discovery

- A host is “born” knowing only its MAC address
- Must discover lots of information before it can communicate with a remote host B
 - what is my IP address?
 - what is B’s IP address? (remote)
 - what is B’s MAC address? (if B is local)
 - what is my first-hop router’s address? (if B is not local)
 - ...

ARP and DHCP

- Link layer discovery protocols
 - “Address Resolution Protocol”, “Dynamic Host Configuration Protocol”
 - confined to a single local-area network (LAN)
 - rely on broadcast capability of a LAN



ARP and DHCP

- Link layer discovery protocols
- Serve two functions
 - Discovery of local end-hosts
 - for communication between hosts on the same LAN

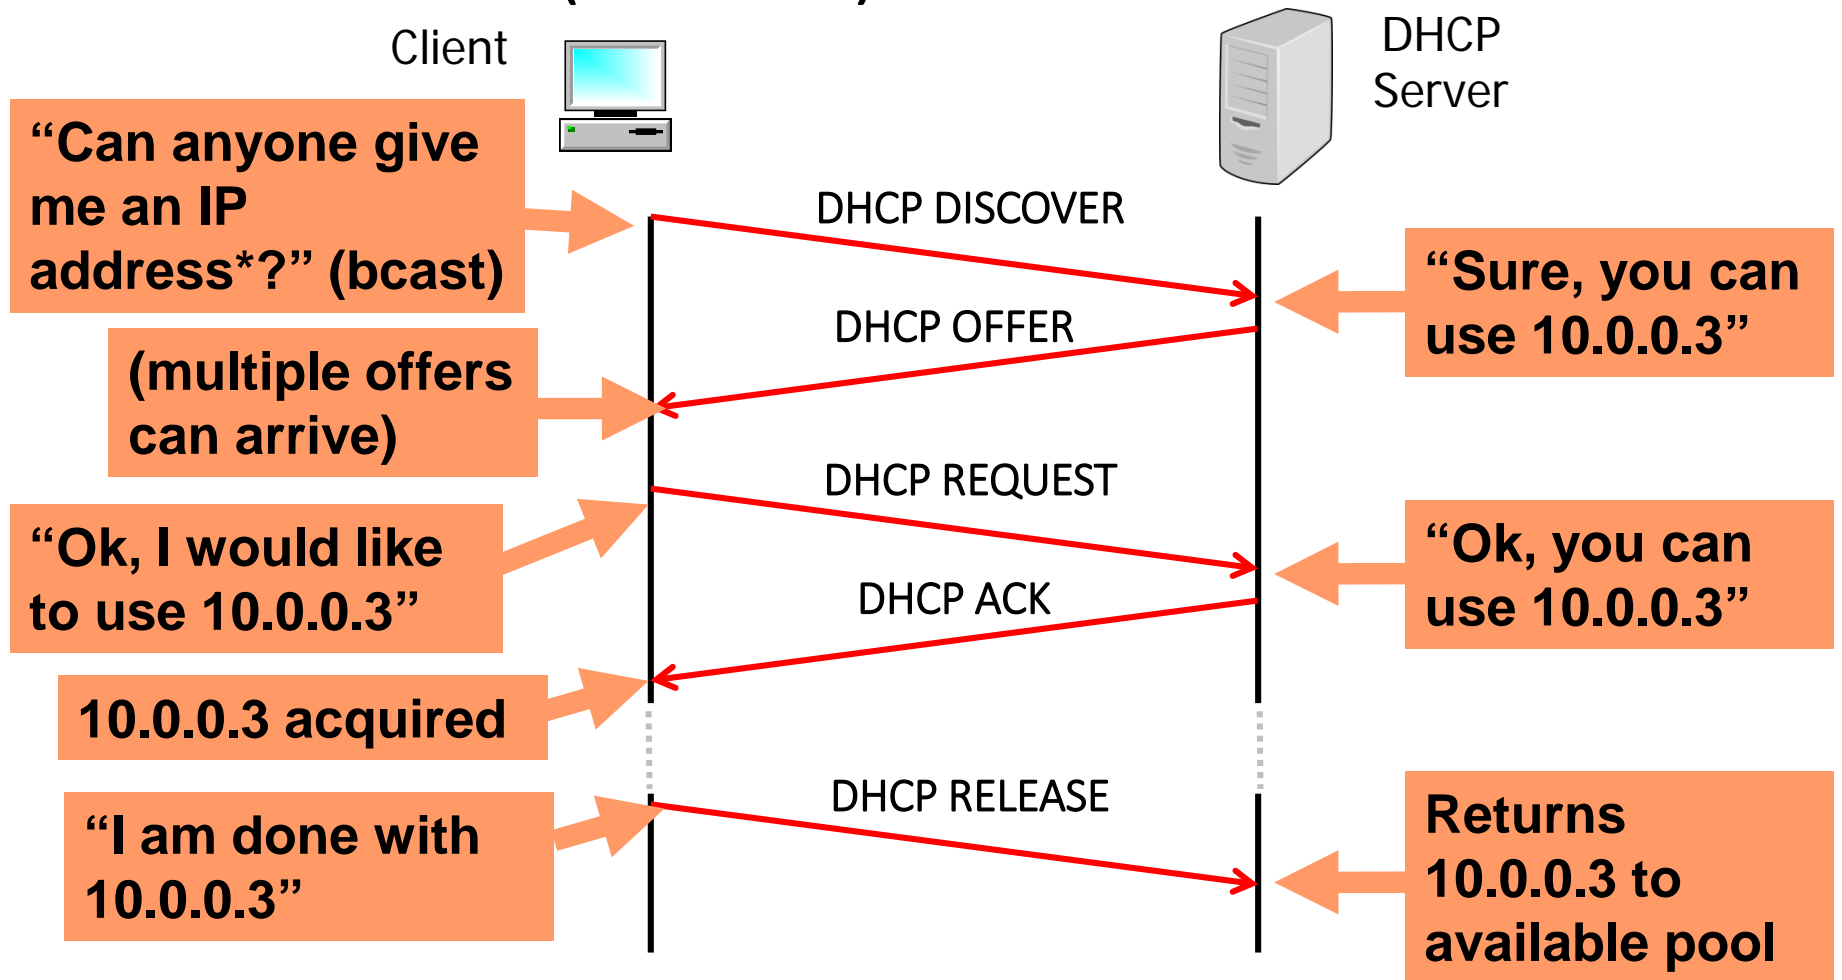
ARP and DHCP

- Link layer discovery protocols
- Serve two functions
 - Discovery of local end-hosts
 - Bootstrap communication with remote hosts
 - what's my IP address?
 - who/where is my local DNS server?
 - who/where is my first hop router?

Dynamic Host Configuration Protocol (DHCP)

- Automatically configure hosts
 - Assign IP addresses, DNS server, default gateway, etc.
 - Client listen on UDP port 68, servers on 67
- Very common LAN protocol
 - Rare to find a device that doesn't support it
- Address is assigned for a **lease time**

Dynamic Host Configuration Protocol (DHCP)



*and other config information

DHCP

- “Dynamic Host Configuration Protocol”
 - defined in RFC 2131
- A host uses DHCP to discover
 - its own IP address
 - its netmask
 - IP address(es) for its DNS name server(s)
 - IP address(es) for its first-hop “default” router(s)

DHCP: operation

1. One or more local DHCP servers maintain required information
 - IP address pool, netmask, DNS servers, *etc.*
 - application that listens on UDP port 67

DHCP: operation

1. One or more local DHCP servers maintain required information
2. Client **broadcasts** a DHCP discovery message
 - L2 broadcast, to MAC address FF:FF:FF:FF:FF:FF

DHCP: operation

1. One or more local DHCP servers maintain required information
2. Client **broadcasts** a DHCP discovery message
3. One or more DHCP servers responds with a DHCP “offer” message
 - proposed IP address for client, lease time
 - other parameters

DHCP: operation

1. One or more local DHCP servers maintain required information
2. Client **broadcasts** a DHCP discovery message
3. One or more DHCP servers responds with a DHCP “offer” message
4. Client **broadcasts** a DHCP request message
 - specifies which offer it wants
 - echoes accepted parameters
 - other DHCP servers learn they were not chosen

DHCP: operation

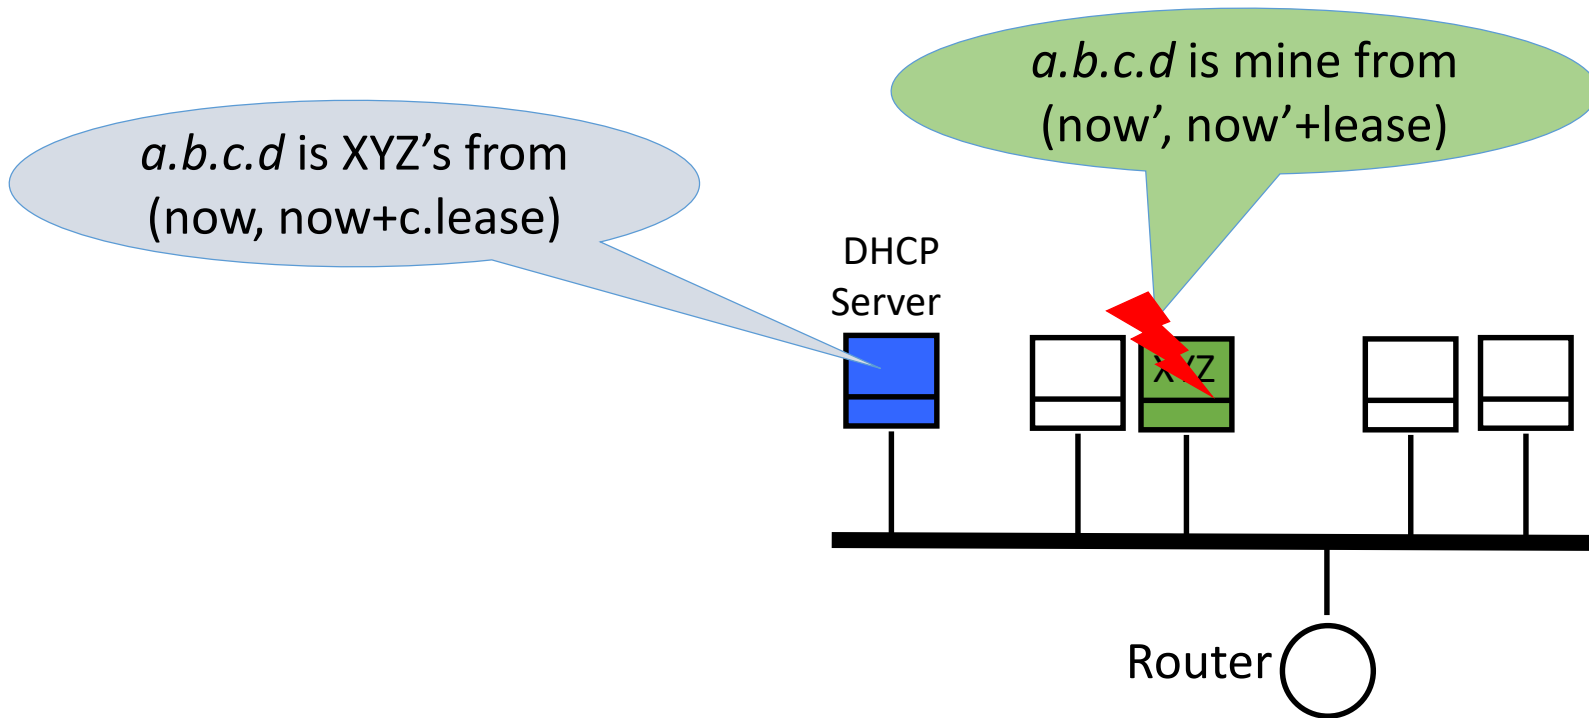
1. One or more local DHCP servers maintain required information
2. Client **broadcasts** a DHCP discovery message
3. One or more DHCP servers responds with a DHCP “offer” message
4. Client **broadcasts** a DHCP request message
5. Selected DHCP server responds with an ACK

(DHCP “relay agents” used when the DHCP server isn’t on the same broadcast domain -- see text)

DHCP uses “soft state”

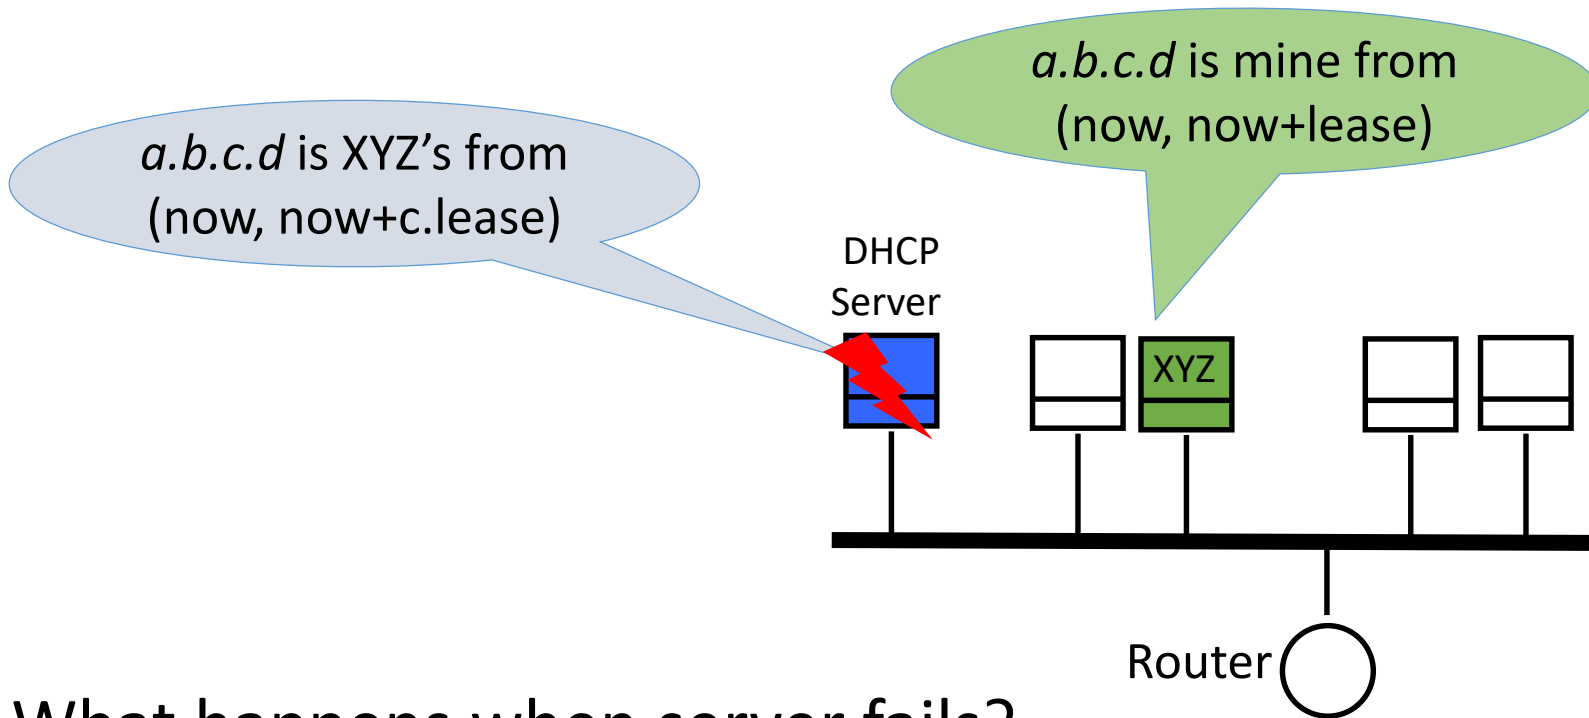
- Soft state: if not refreshed, state is forgotten
 - hard state: allocation is deliberately returned/withdrawn
 - used to track address allocation in DHCP
- Implementation
 - address allocations are associated with a lease period
 - server: sets a timer associated with the record of allocation
 - client: must request a refresh before lease period expires
 - server: resets timer when a refresh arrives; sends ACK
 - server: reclaims allocated address when timer expires
- Simple, yet robust under failure
 - state always fixes itself in (small constant of) lease time

Soft state under failure



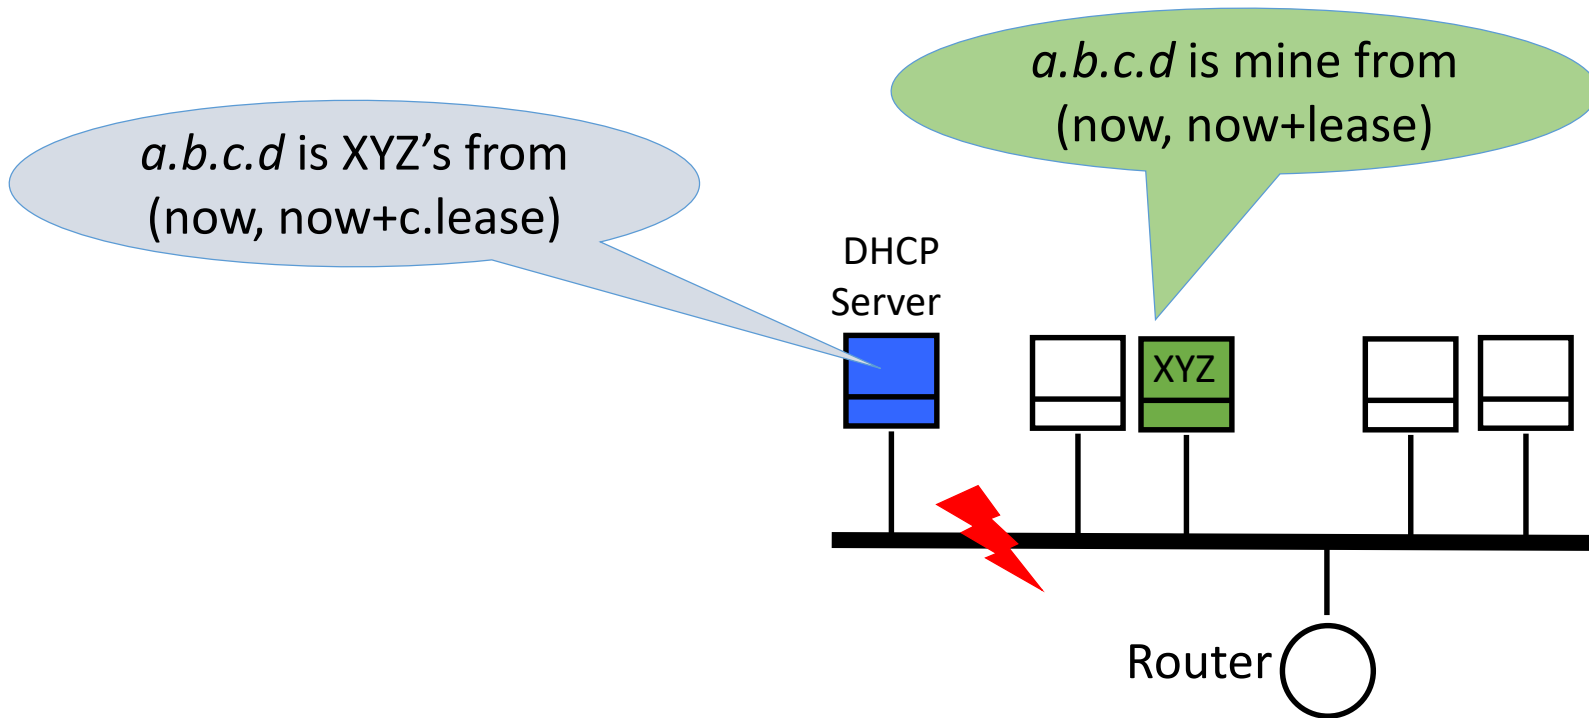
- What happens when host XYZ fails?
 - refreshes from XYZ stop
 - server reclaims *a.b.c.d* after $O(\text{lease period})$

Soft state under failure



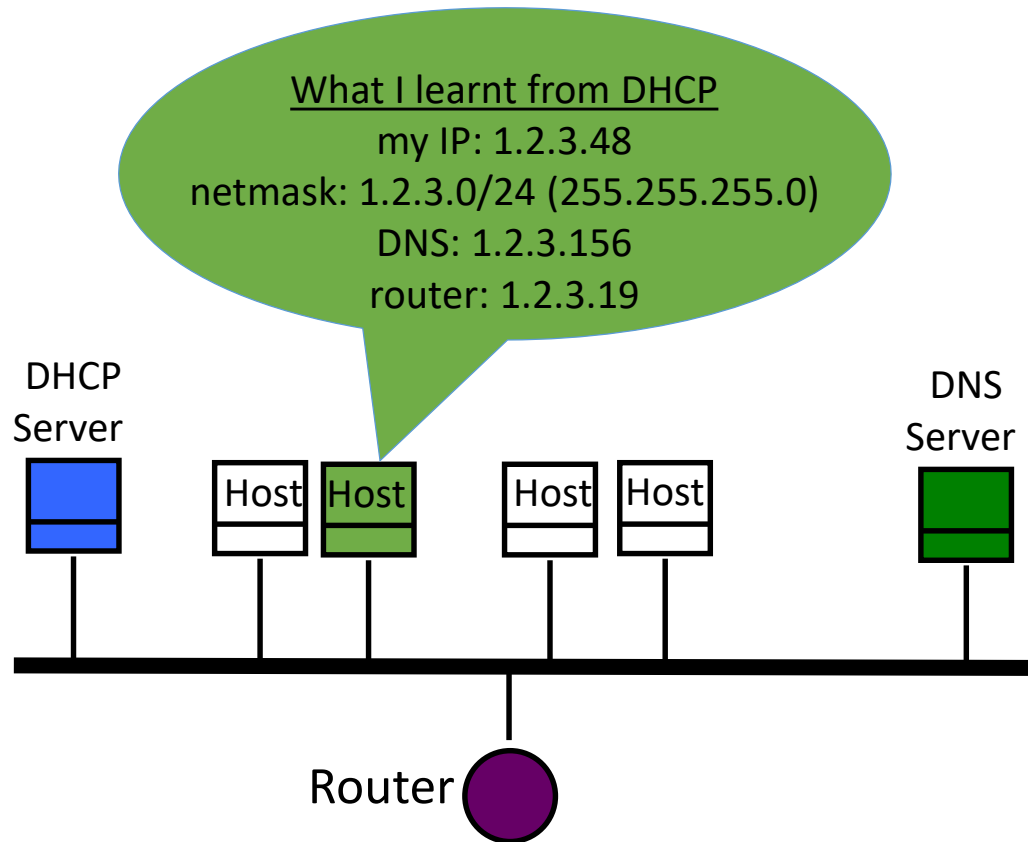
- What happens when server fails?
 - ACKs from server stop
 - XYZ releases address after $O(\text{lease period})$; send new request
 - A new DHCP server can come up from a 'cold start' and we're back on track in $\sim \text{lease time}$

Soft state under failure




- What happens if the network fails?
 - refreshes and ACKs don't get through
 - XYZ release address; DHCP server reclaims it

Are we there yet?



Sending Packets Over Link-Layer



- Link layer only understands MAC addresses
 - Translate the destination IP address to MAC address
 - Encapsulate the IP packet inside a link-level frame

ARP: Address Resolution Protocol

```
C:\Users\Matthew Caesar>arp -a

Interface: 192.168.1.84 --- 0x8
    Internet Address      Physical Address      Type
    192.168.1.76          84-d6-d0-a1-05-84     dynamic
    192.168.1.87          00-80-92-cb-aa-0c     dynamic
    192.168.1.254         f8-2c-18-94-96-7d     dynamic
    192.168.1.255         ff-ff-ff-ff-ff-ff     static
    224.0.0.22            01-00-5e-00-00-16     static
    224.0.0.251           01-00-5e-00-00-fb     static
    224.0.0.252           01-00-5e-00-00-fc     static
    239.255.255.250       01-00-5e-7f-ff-fa     static
    255.255.255.255       ff-ff-ff-ff-ff-ff     static
```

- Every host maintains an **ARP** table
 - list of (IP address → MAC address) pairs
- Consult the table when sending a packet
 - Map destination IP address to destination MAC address
 - Encapsulate the (IP) data packet with MAC header; transmit
- But: what if IP address **not** in the table?
 - Sender broadcasts: **“Who has IP address 1.2.3.156?”**
 - Receiver responds: **“MAC address 58-23-D7-FA-20-B0”**
 - Sender caches result in its ARP table

Address Resolution Protocol (ARP)

- Networked applications are programmed to deal with IP addresses
- But Ethernet forwards to MAC address
- How can OS know the MAC address corresponding to a given IP address?
- Solution: **Address Resolution Protocol**
 - Broadcasts **ARP request** for MAC address owning a given IP address

Broadcast ARP request:
“Who owns IP address 4.4.4.4?”

IP=2.2.2.2
MAC=AA:AA:AA:AA:AA

IP=3.3.3.3
MAC=BB:BB:BB:BB:BB

<u>IP</u>	<u>MAC</u>
4.4.4.4	CC:CC:CC:CC:CC
5.5.5.5	DD:DD:DD:DD:DD

Broadcast ARP reply:
“I own 4.4.4.4, and my MAC address is CC:CC:CC:CC:CC”

IP=4.4.4.4
MAC=CC:CC:CC:CC:CC


Broadcast *Gratuitous* ARP reply:
“I own 5.5.5.5, and my MAC address is DD:DD:DD:DD:DD”

IP=5.5.5.5
MAC=DD:DD:DD:DD:DD

- ARP: determine mapping from IP to MAC address
- What if IP address not on subnet?
 - Each host configured with “default gateway”, use ARP to resolve its IP address
- **Gratuitous ARP:** tell network your IP to MAC mapping
 - Used to detect IP conflicts, IP address changes; update other machines' ARP tables, update bridges' learned information

What if the destination is remote?

- Look up the MAC address of the first hop router
 - 1.2.3.48 uses ARP to find MAC address for first-hop router **1.2.3.19** rather than ultimate destination IP address
- How does the red host know the destination is not local?
 - Uses netmask (entered manually or discovered via DHCP)
- How does the red host know about 1.2.3.19?
 - Also entered manually or DHCP (assigned as the “gateway”)



Security Analysis of ARP

- Impersonation
 - Any node that hears request can answer ...
 - ... and can say whatever they want
- Actual legit receiver never sees a problem
 - Because even though later packets carry its IP address, its NIC doesn't capture them since not its MAC address
- Solutions
 - Dynamic Arp Inspection (DAI) – disallow ARP requests from MAC address not manually whitelisted on port
 - Port Security – configure port to disallow data traffic from MAC address not manually whitelisted to port

Steps in Sending a Packet

What do hosts need to know?

And how do they find out?


Steps in reaching a Host

- First look up destination's IP address
- Need to know where local DNS server is
 - DHCP
- Also needs to know its own IP address
 - DHCP

Sending a Packet


- On same subnet:
 - Use MAC address of destination.
 - *ARP*
- On some other subnet:
 - Use MAC address of first-hop router.
 - *DHCP + ARP*
- And how can a host tell whether destination is on same or other subnet?
 - Use the *netmask*
 - *DHCP*

Example: A Sending a Packet to B



How does host **A** send an IP packet to host **B**?


Example: A Sending a Packet to B



1. **A** sends packet to **R**.
2. **R** sends packet to **B**.


Host A Decides to Send Through R

- Host **A** constructs an IP packet to send to **B**
 - Source 111.111.111.111, destination 222.222.222.222
- Host **A** has a gateway router **R**
 - Used to reach destinations outside of 111.111.111.0/24
 - Address 111.111.111.110 for R learned via **DHCP**




Host A Sends Packet Through R

- Host **A** learns the MAC address of **R**'s interface
 - **ARP** request: broadcast request for 111.111.111.110
 - **ARP** response: **R** responds with E6-E9-00-17-BB-4B
- Host **A** encapsulates the packet and sends to **R**




R Decides how to Forward Packet

- Router **R**'s adapter receives the packet
 - **R** extracts the IP packet from the Ethernet frame
 - **R** sees the IP packet is destined to 222.222.222.222
- Router **R** consults its forwarding table
 - Packet matches 222.222.222.0/24 via other adapter (port)



R Sends Packet to B

- Router **R** learns the MAC address of host **B**
 - **ARP** request: broadcast request for 222.222.222.222
 - **ARP** response: **B** responds with 49-BD-D2-C7-56-2A
- Router **R** encapsulates the packet and sends to **B**



Key Ideas in Both ARP and DHCP

- **Broadcasting**: used for initial bootstrap
- **Caching**: remember the past for a while
 - Store the information you learn to reduce overhead
 - Remember your own address & other host's addresses
 - *Key optimization for performance*
- **Soft state**: eventually forget the past
 - Associate a time-to-live field with the information
 - ... and either refresh or discard the information
 - *Key for robustness*

Discovery mechanisms

We've seen two broad approaches

- Broadcast (ARP, DHCP)
 - flooding doesn't scale
 - no centralized point of failure
 - zero configuration
- Directory service (DNS)
 - no flooding
 - root of the directory is vulnerable (caching is key)
 - needs configuration to bootstrap (local, root servers, *etc.*)

Can we get the best of both?

- Internet-scale yet zero config?