

Cycle 4: Cloud Networking and Services

CS 436: Spring 2018


Matthew Caesar

<http://www.cs.illinois.edu/~caesar/cs436>

What is a cloud?

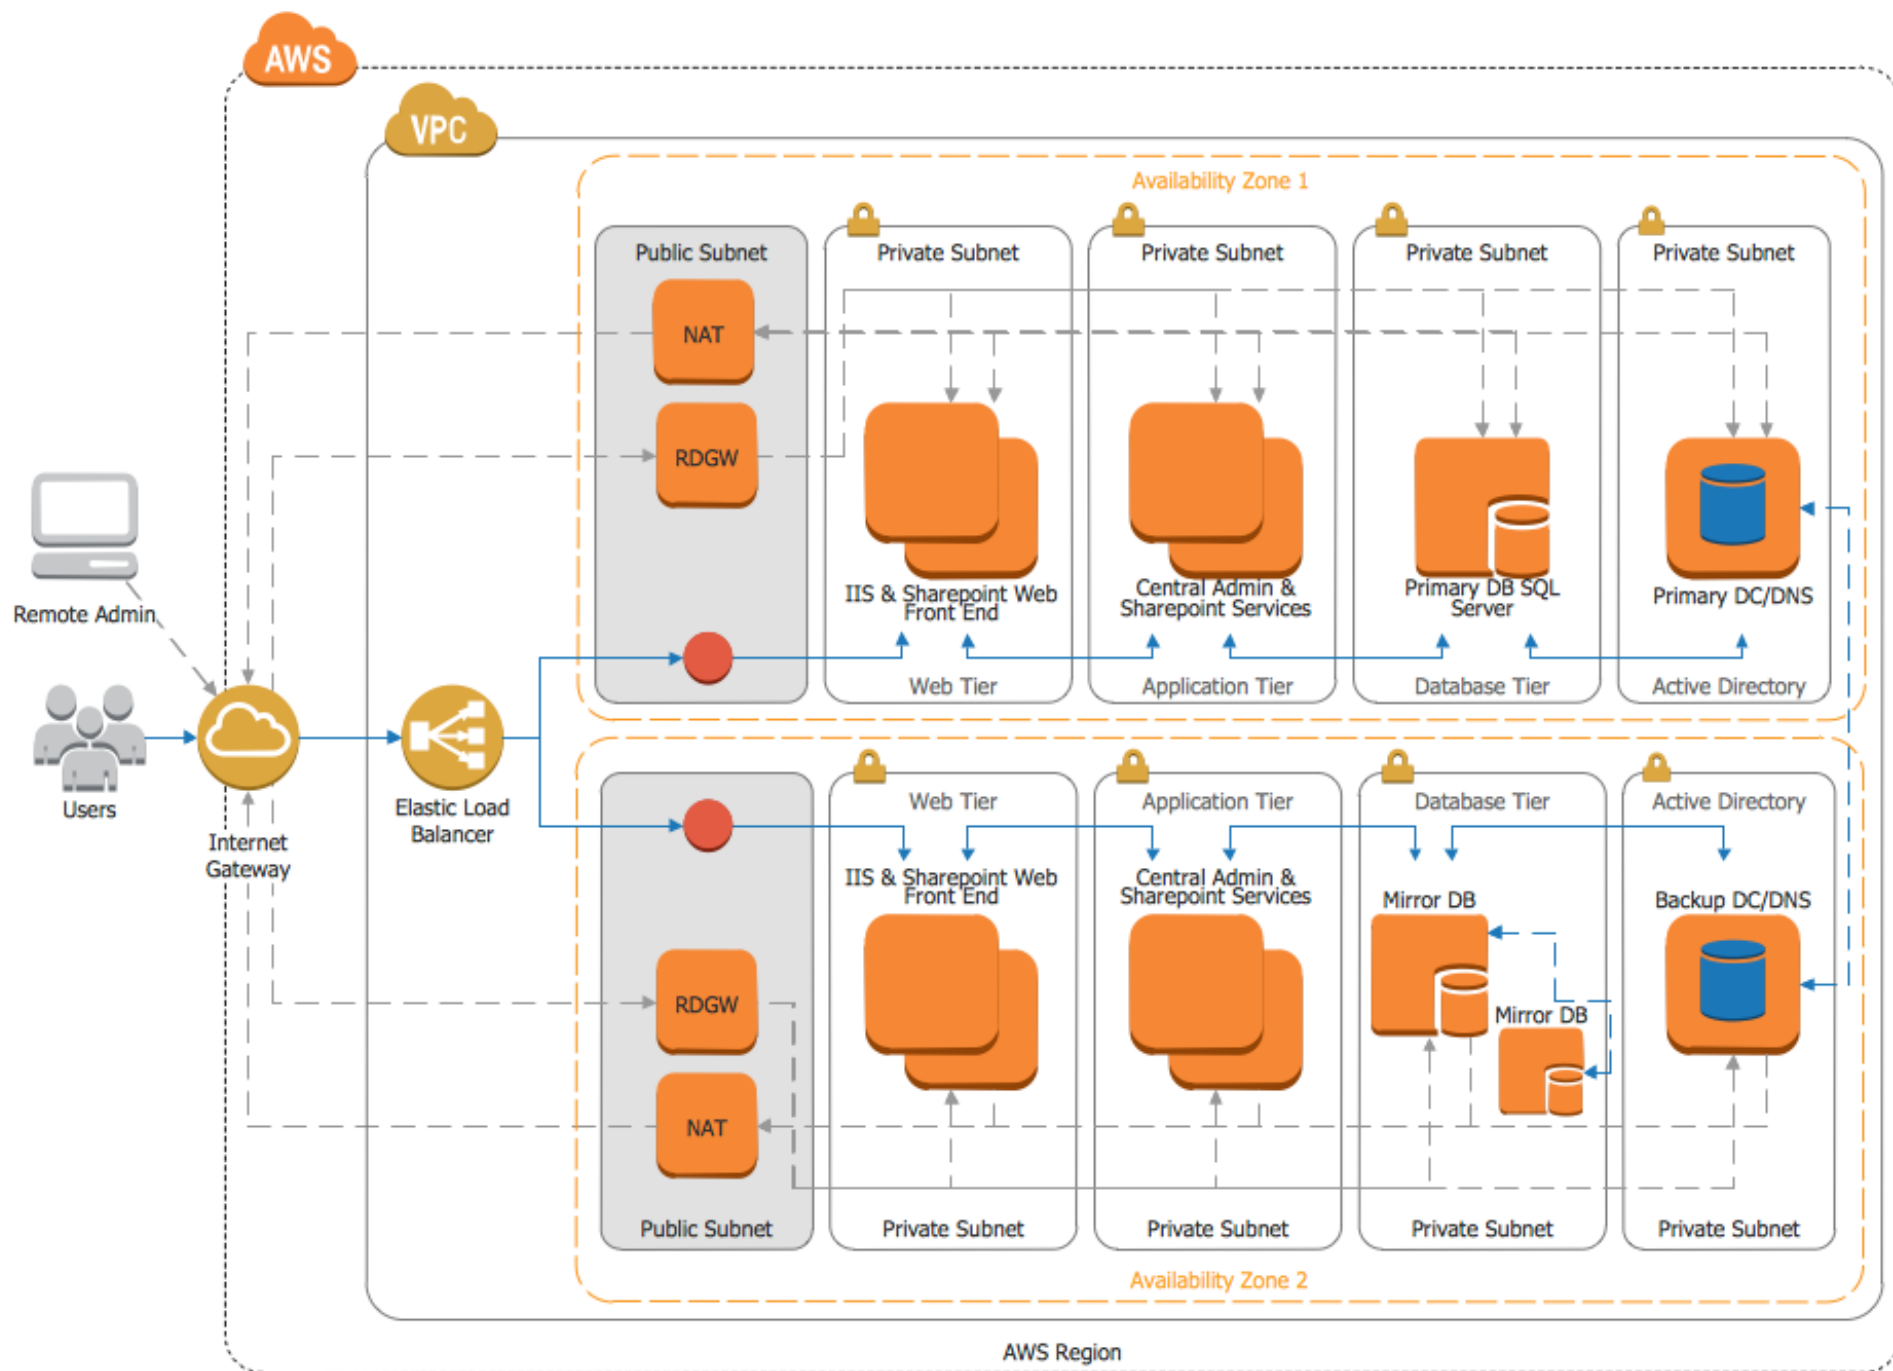
- IT model enabling shared, ubiquitous access to compute resources
- Cloud networking: access of networking resources in the cloud
- Examples:
 - Cloud-hosted networking (AWS, Azure)
 - Cloud-hosted management (Meraki, Aerohive)

Types of clouds



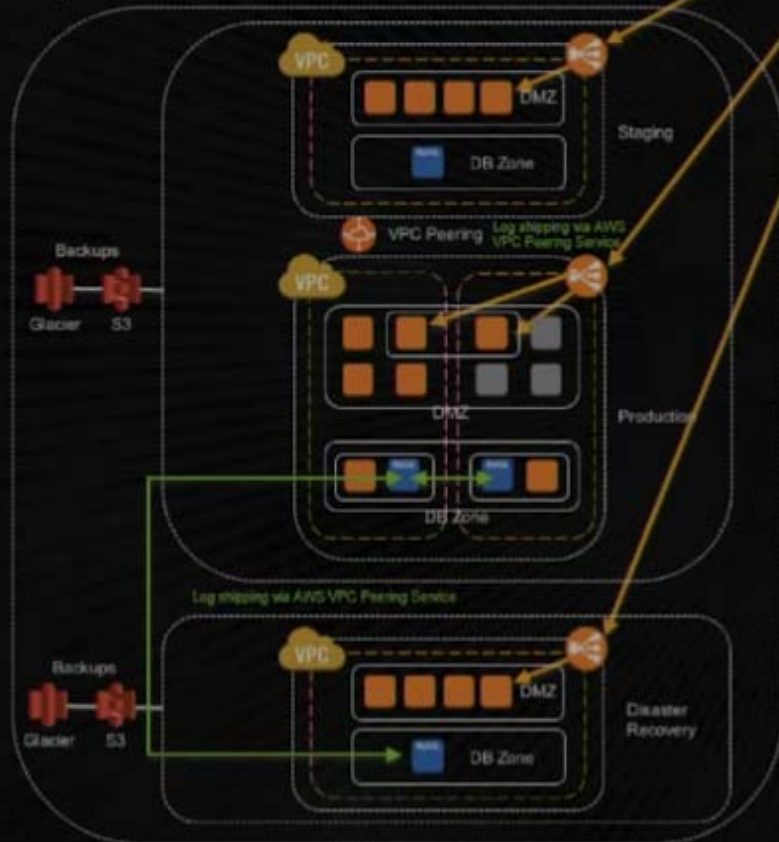
Types of clouds

- Software as a Service (SaaS)
 - Provider provides software, which is centrally hosted
 - Examples: Google Apps, Salesforce, Cisco WebEx
- Infrastructure as a Service (IaaS)
 - Customers lease virtual machines they configure themselves
 - Examples: AWS, Azure, IBM SmartCloud
- Platform as a Service (PaaS)
 - Customers work with services created/maintained by provider
 - Examples: AWS, Azure, RedHat OpenShift
- Function as a Service (FaaS)
 - Customer provides code, Provider provides functions and event triggers
 - Examples: Google Cloud Functions, AWS Lambda, MS Azure Functions

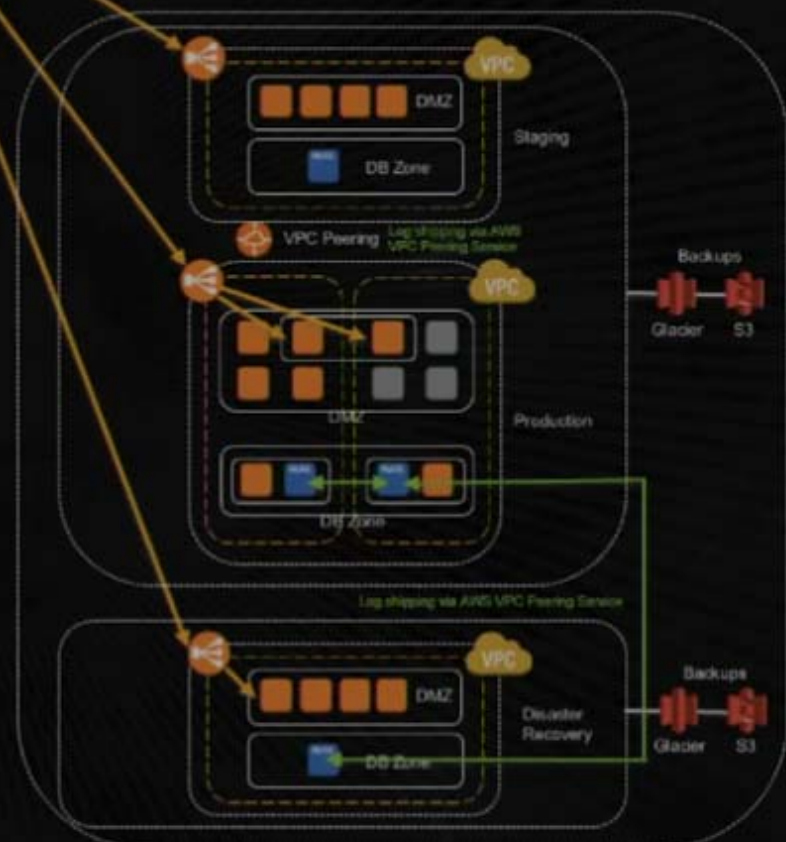




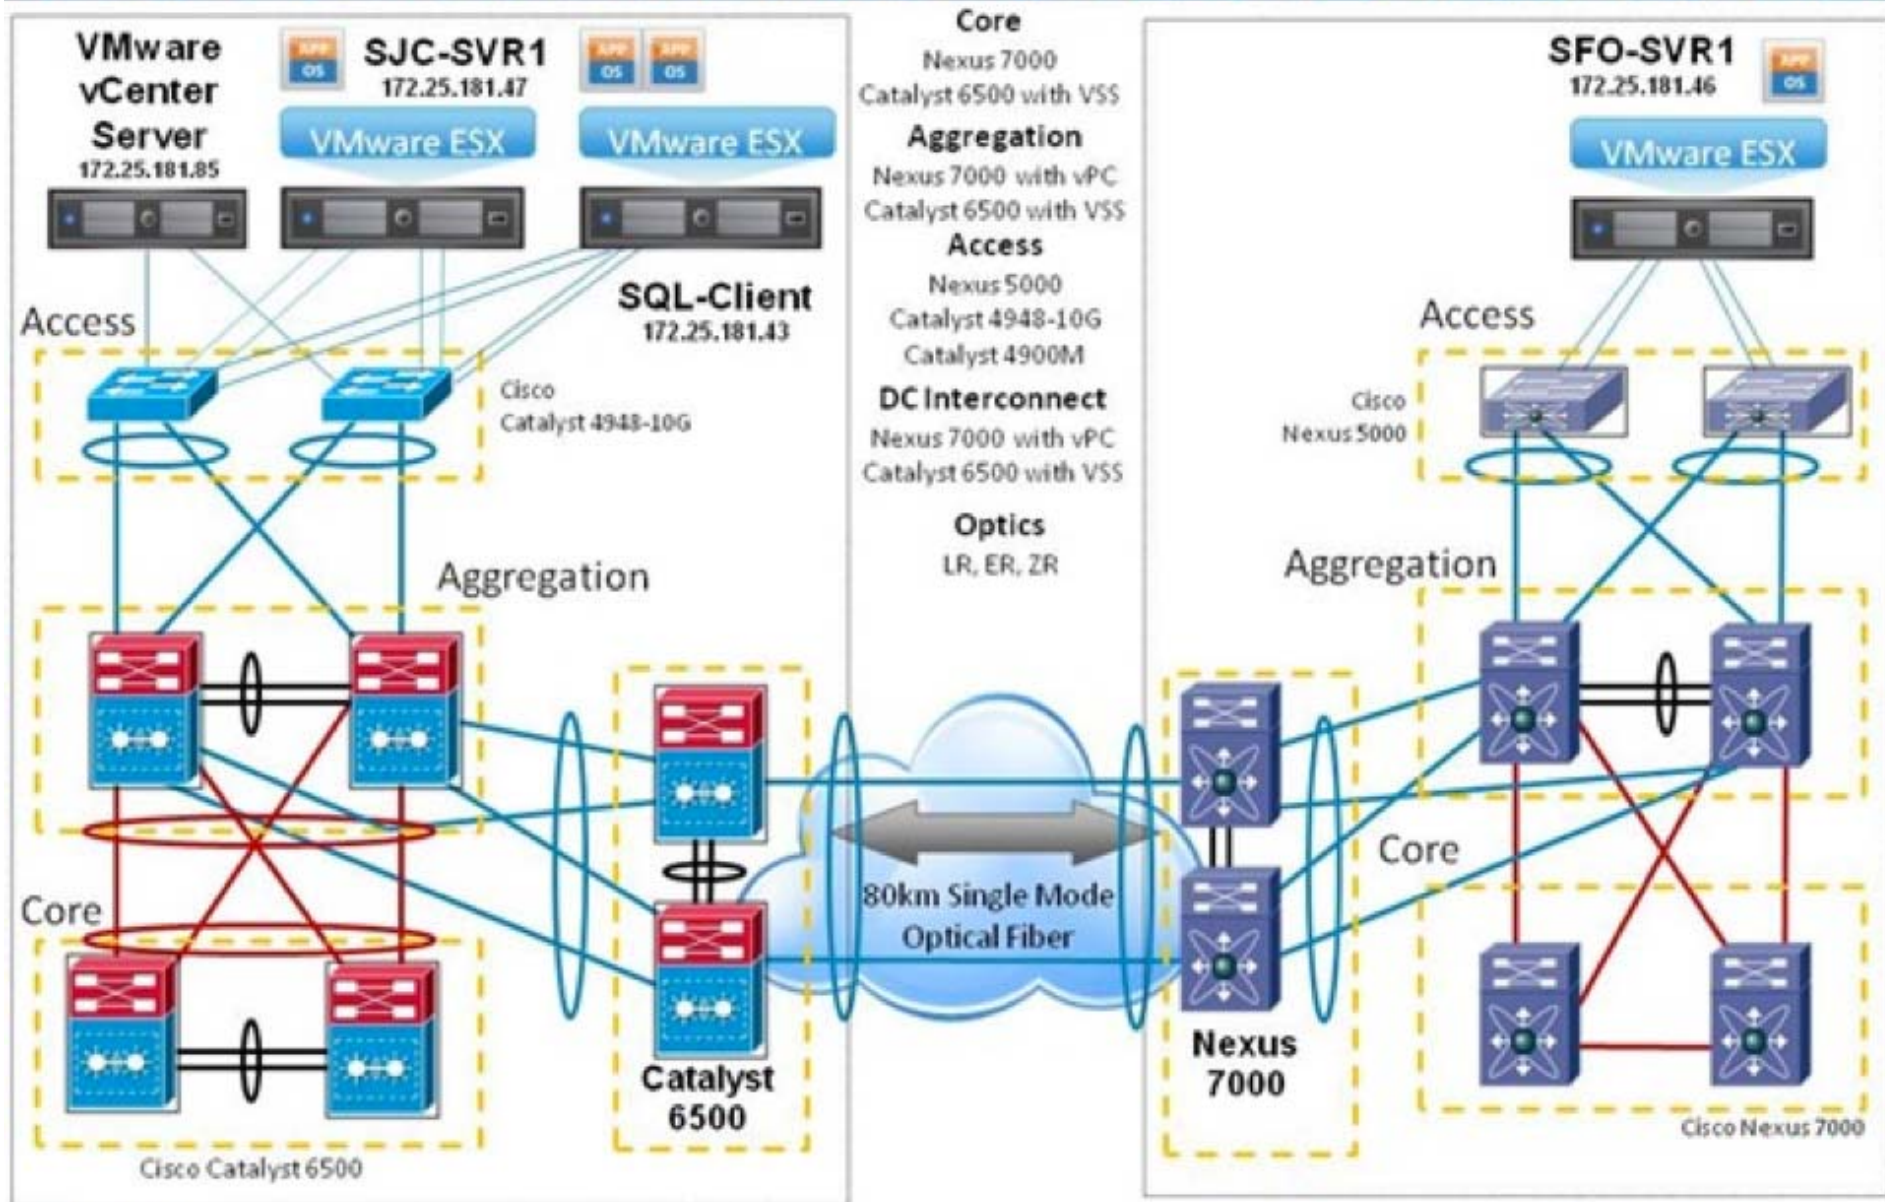
Application one



Application two



VMware and Cisco: Inter-Datacenter VMotion Test Topology



San Jose

San Francisco



Public Services

Directing Traffic to Services

- Need to provide services to public users (or internal users)

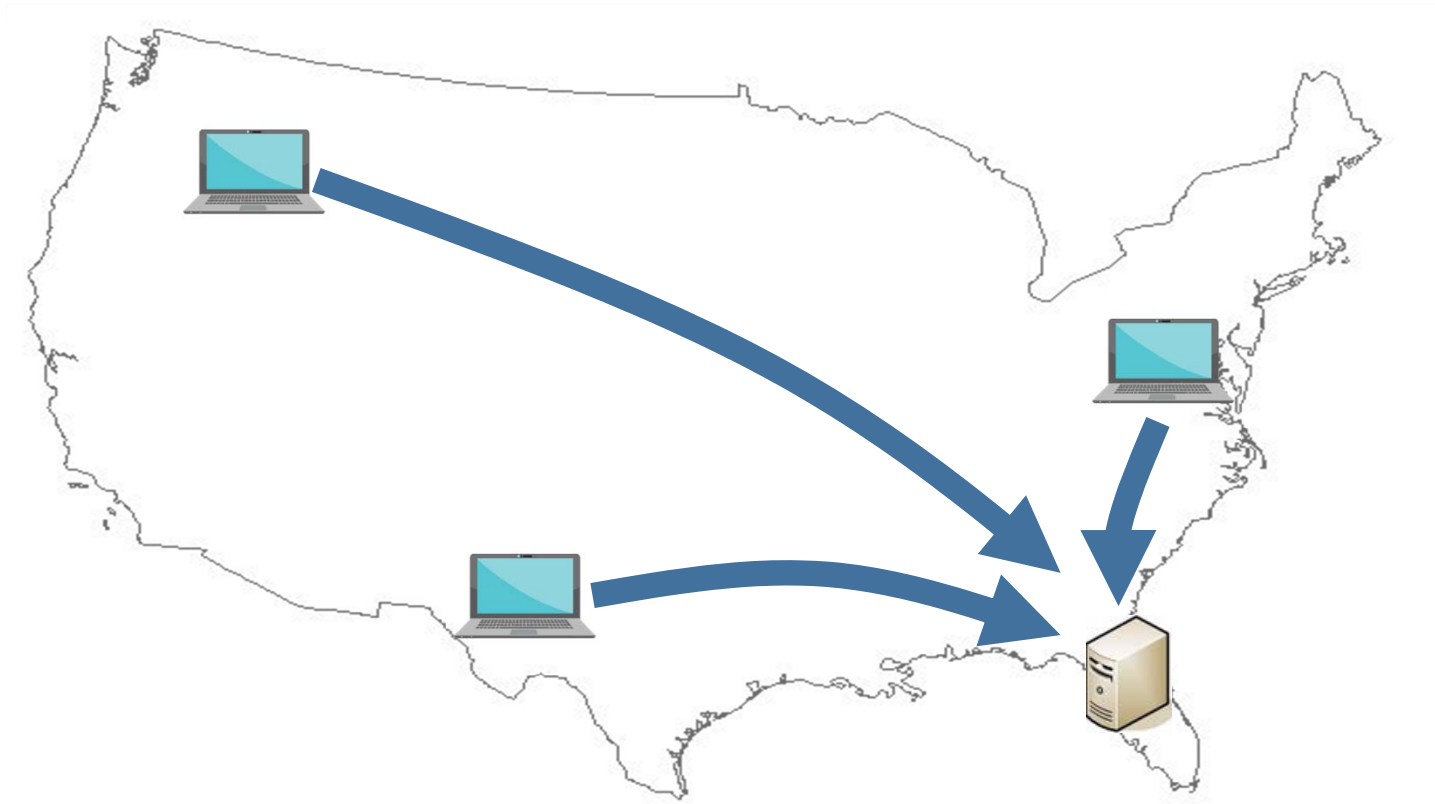


- Challenges: redundancy, low latency, etc
- Concepts:
 - Anycast/Load Balancing
 - NAT

Anycast and Load Balancing


- Direct users to servers
- Done for various reasons
 - For resiliency
 - To balance load
 - To give differential QOS, to target advertisements, to censor...
- Key approaches
 - Wide-area: DNS anycast
 - Wide- and local-area: IP anycast
 - Local-area: load balancing appliance

How to deploy a big Internet svc



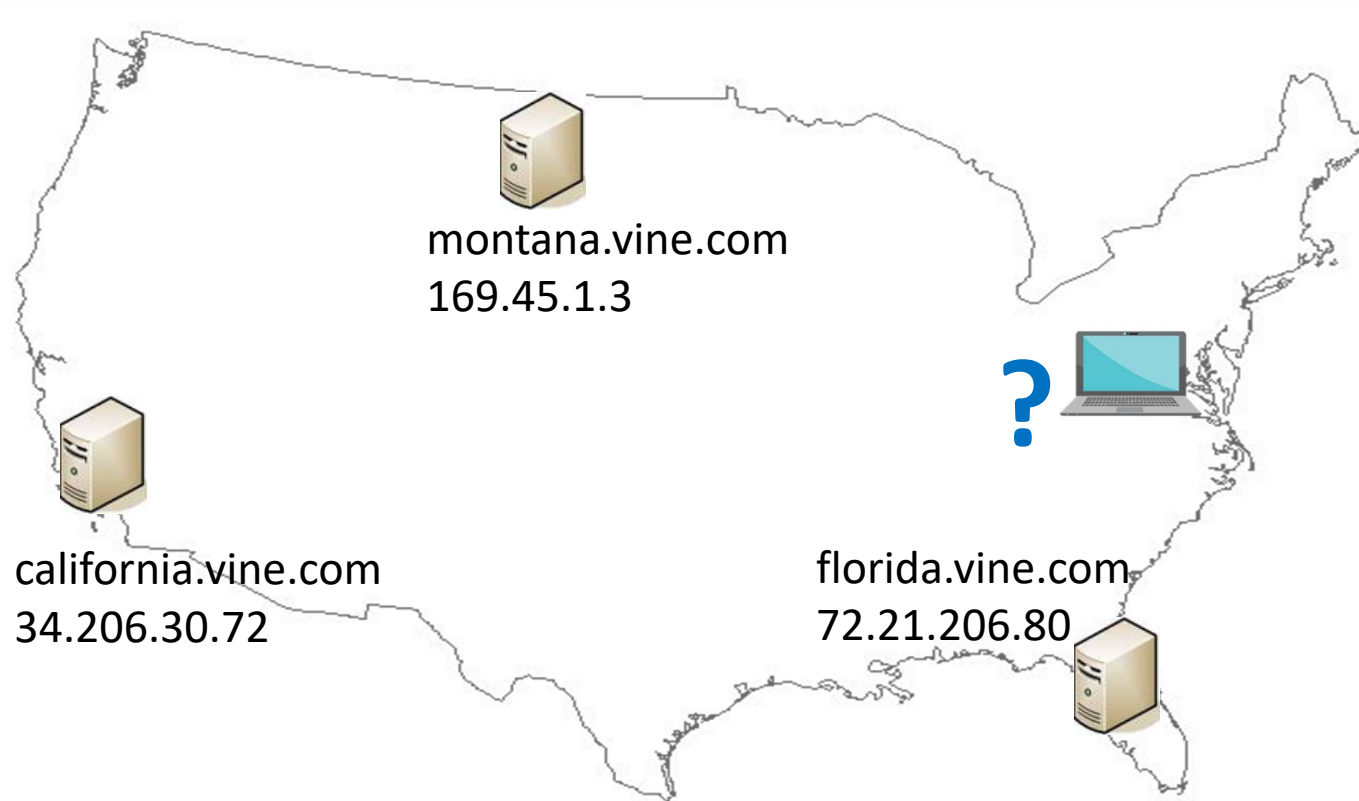
- Solution: Put down a server somewhere
- Problems: Too much load, Far away from some clients, Single point of failure, Easy to DoS

Better: Deploy multiple servers



- Solution: Put down several servers
- Benefits: Improve scaling, resilience, proximity
- Question: How to actually implement this

How about: Tell users where to go



- Solution: Tell users which servers to go to
- Problem: How are you going to enforce that
 - Could filter users who don't listen
 - How to know where users are located?

Better idea: Transparent direction

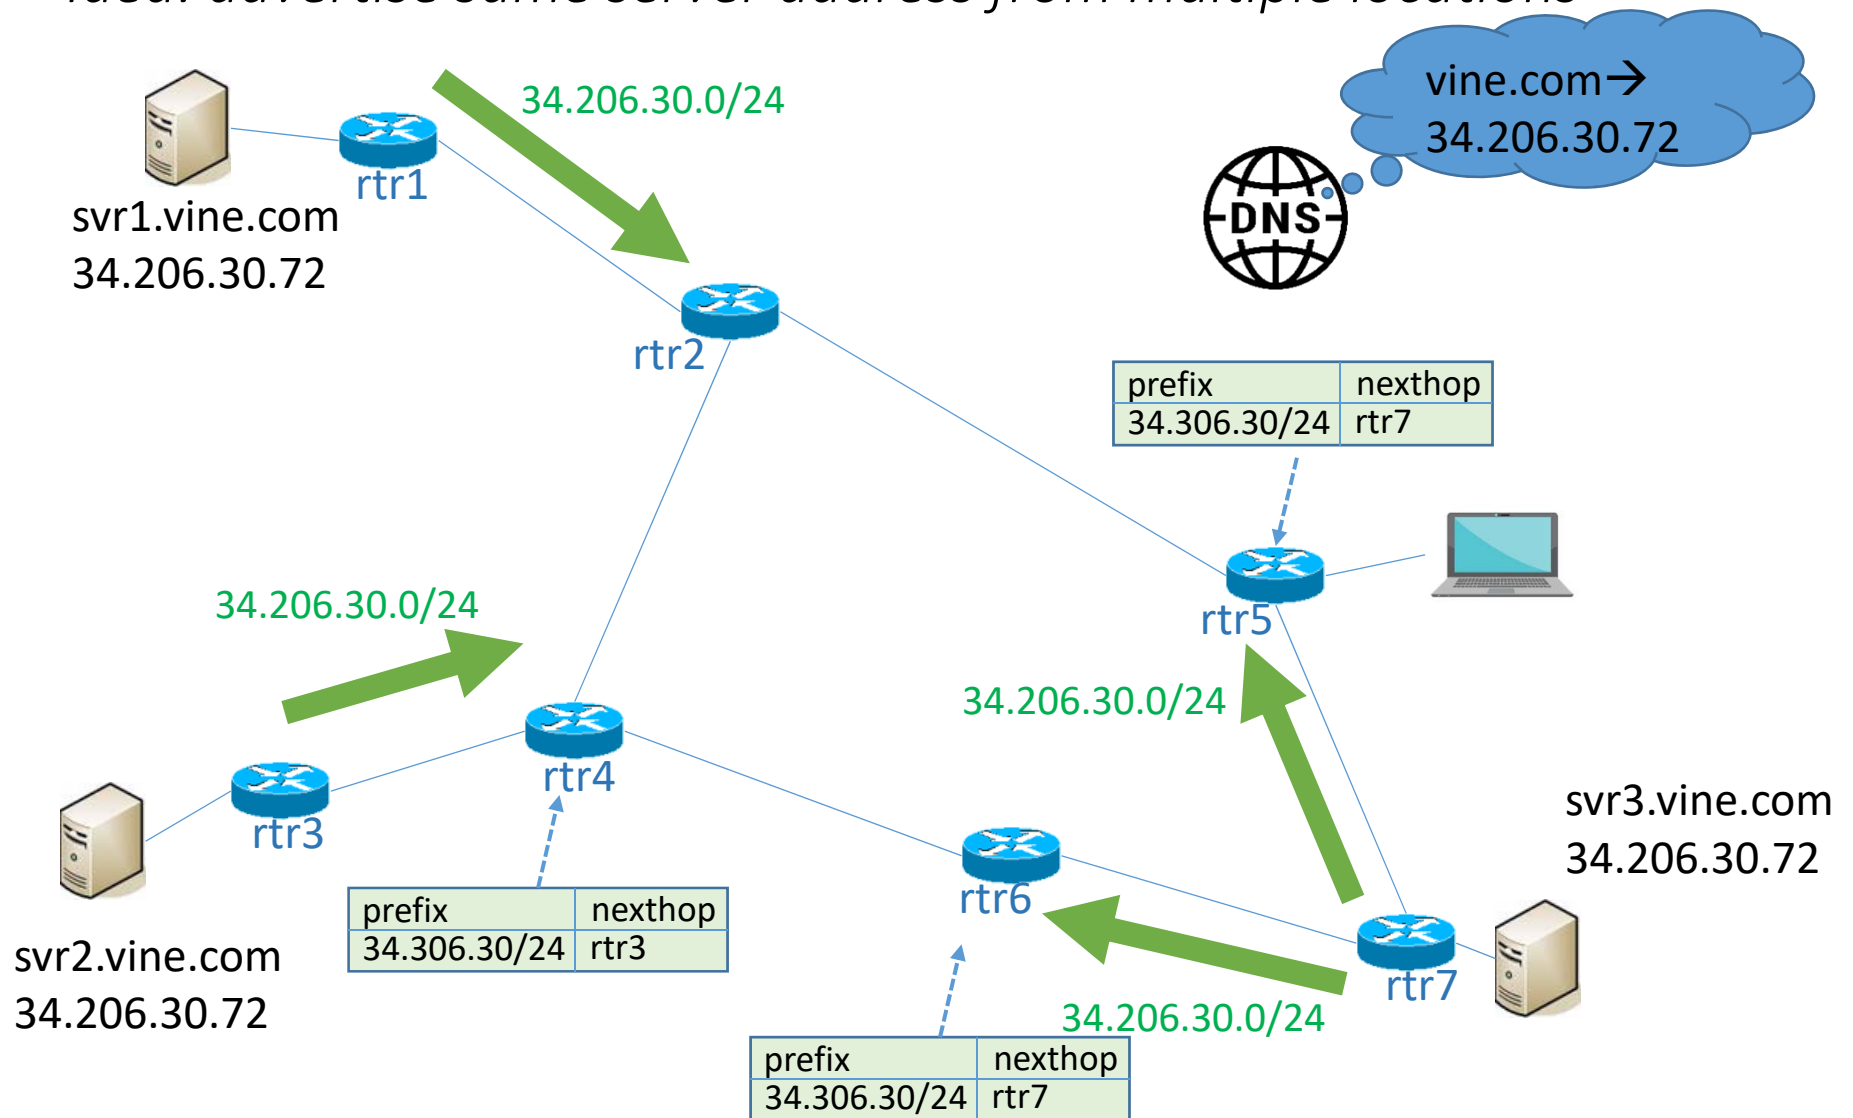
- Give the illusion of single server
- Benefits:
 - Don't have to coordinate with humans
 - Can force users to servers in desirable ways
- Approach: Anycast
 - DNS Anycast (typically wide-area)
 - IP Anycast (typically wide- and regional-area)
 - Load balancers (typically local-area)

Approach 1: IP Anycast

- Idea: advertise same server address from multiple locations
 - Configure upstream routers to advertise same/overlapping prefix
- Routers in between choose shortest path
 - Directs traffic towards “closest” server
 - Closest in terms of # of hops

Approach 1: IP Anycast

Idea: advertise same server address from multiple locations



IP Anycast: Pros and Cons

- IP Anycast has several nice properties
- Fully transparent to higher layers
- ISPs and network can control
- Fast failover for network failures
 - Networks on server events leads to churn
- Cons:
 - Imprecise control (can't always control how advertisements flow)

Approach 2: DNS Anycast

- DNS is a huge distributed service
 - But you control the name→address mapping for your domain
- Idea: modify Authoritative DNS to return different mappings for different clients
 - Goal: map clients to “nearby” servers
- Inputs to mapping function
 - Mapping of Internet topology
 - Useful for estimating latency between hosts
 - Server loads
 - Etc

How are users assigned to PoPs

- Through DNS
 - Resolving server can hand back an IP address “close” to requesting user

Chicago

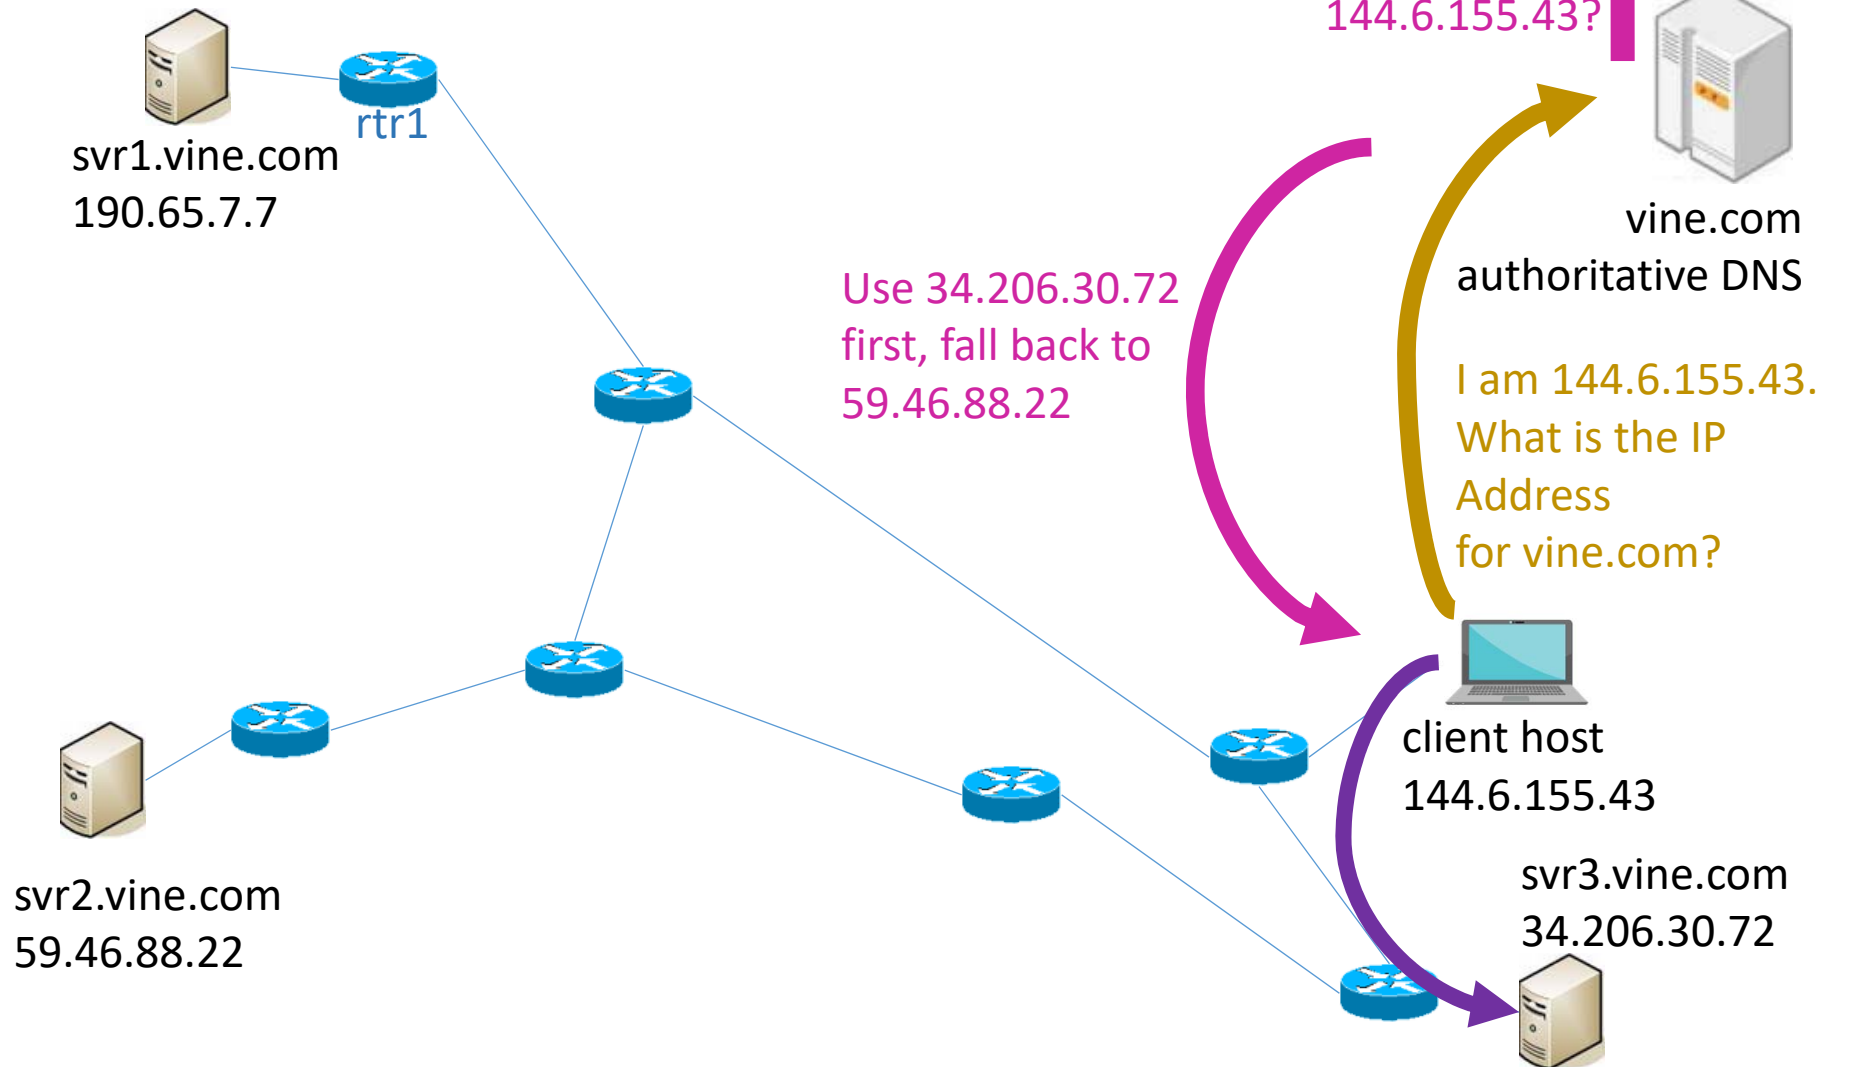
```
$ dig +short www.linkedin.com  
216.52.242.80
```

Spain

```
$ dig @109.69.8.51 +short www.linkedin.com  
91.225.248.80
```

Approach 2: DNS Anycast

Idea: Configure DNS to return subset of servers, based on network properties



Approach 2: DNS Anycast

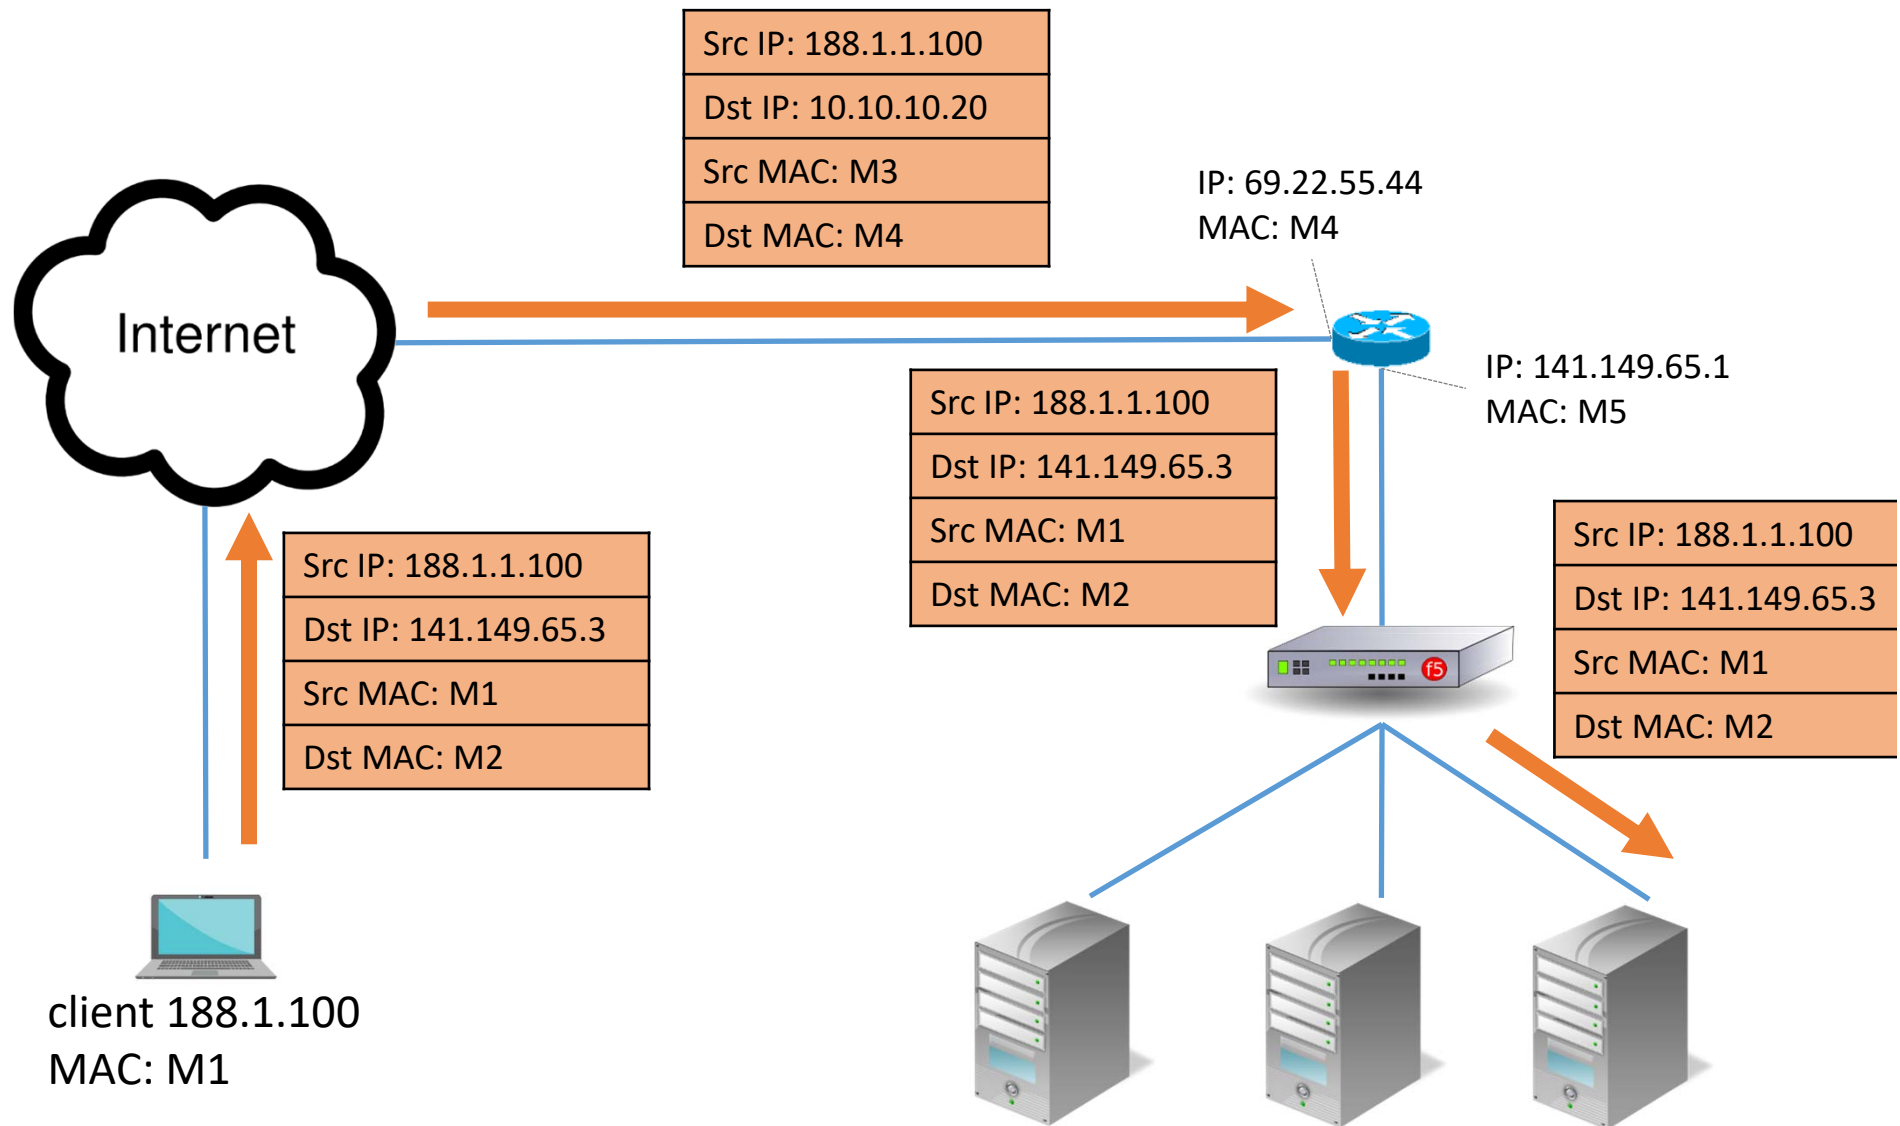
Config DNS to return subset of servers, based on network properties

- **Benefits**
 - Tighter control over host-server mappings than IP anycast
 - Can map based on latency, server load
 - No need to coordinate with network operator
 - Need to convince providers to advertise prefix
- **Cons**
 - Harder for ISPs to control traffic
 - May require additional infrastructure to probe/survey network
 - Need to coordinate with DNS provider
 - Clients may be far from their local DNS (where request actually comes from)
 - Some local DNS servers don't conform to specifications

Approach 3: Load-Balancing Devices


- In the local area, we have tighter control over forwarding
 - Can force ingress/egress path to be symmetric
 - Forwarding speeds low enough to deeply inspect packets
 - Network latency less of an issue
- Approach: install load-balancing hardware into the network
 - Device monitors servers, uses load balancing algorithm to distribute traffic

Approach 3: Load-Balancing Devices



DNS

DNS Lookup Process (Background)



Internet Names & Addresses

- Machine addresses: *e.g., 192.17.90.136*
 - router-usable labels for machines
 - conforms to network structure (the “where”)
- Machine names: *e.g., instr.engr.illinois.edu*
 - human-usable labels for machines
 - conforms to organizational structure (the “who”)
- The Domain Name System (DNS) is how we map from one to the other

DNS: History

- Initially all host-address mappings were in a hosts.txt file (in /etc/hosts):
 - Maintained by the Stanford Research Institute (SRI)
 - Changes were submitted to SRI by email
 - New versions of hosts.txt periodically FTP'd from SRI
 - An administrator could pick names at their discretion
- As the Internet grew this system broke down:
 - SRI couldn't handle the load; names were not unique; hosts had inaccurate copies of hosts.txt
- The Domain Name System (DNS) was invented to fix this
 - First name server implementation done by 4 students

Goals

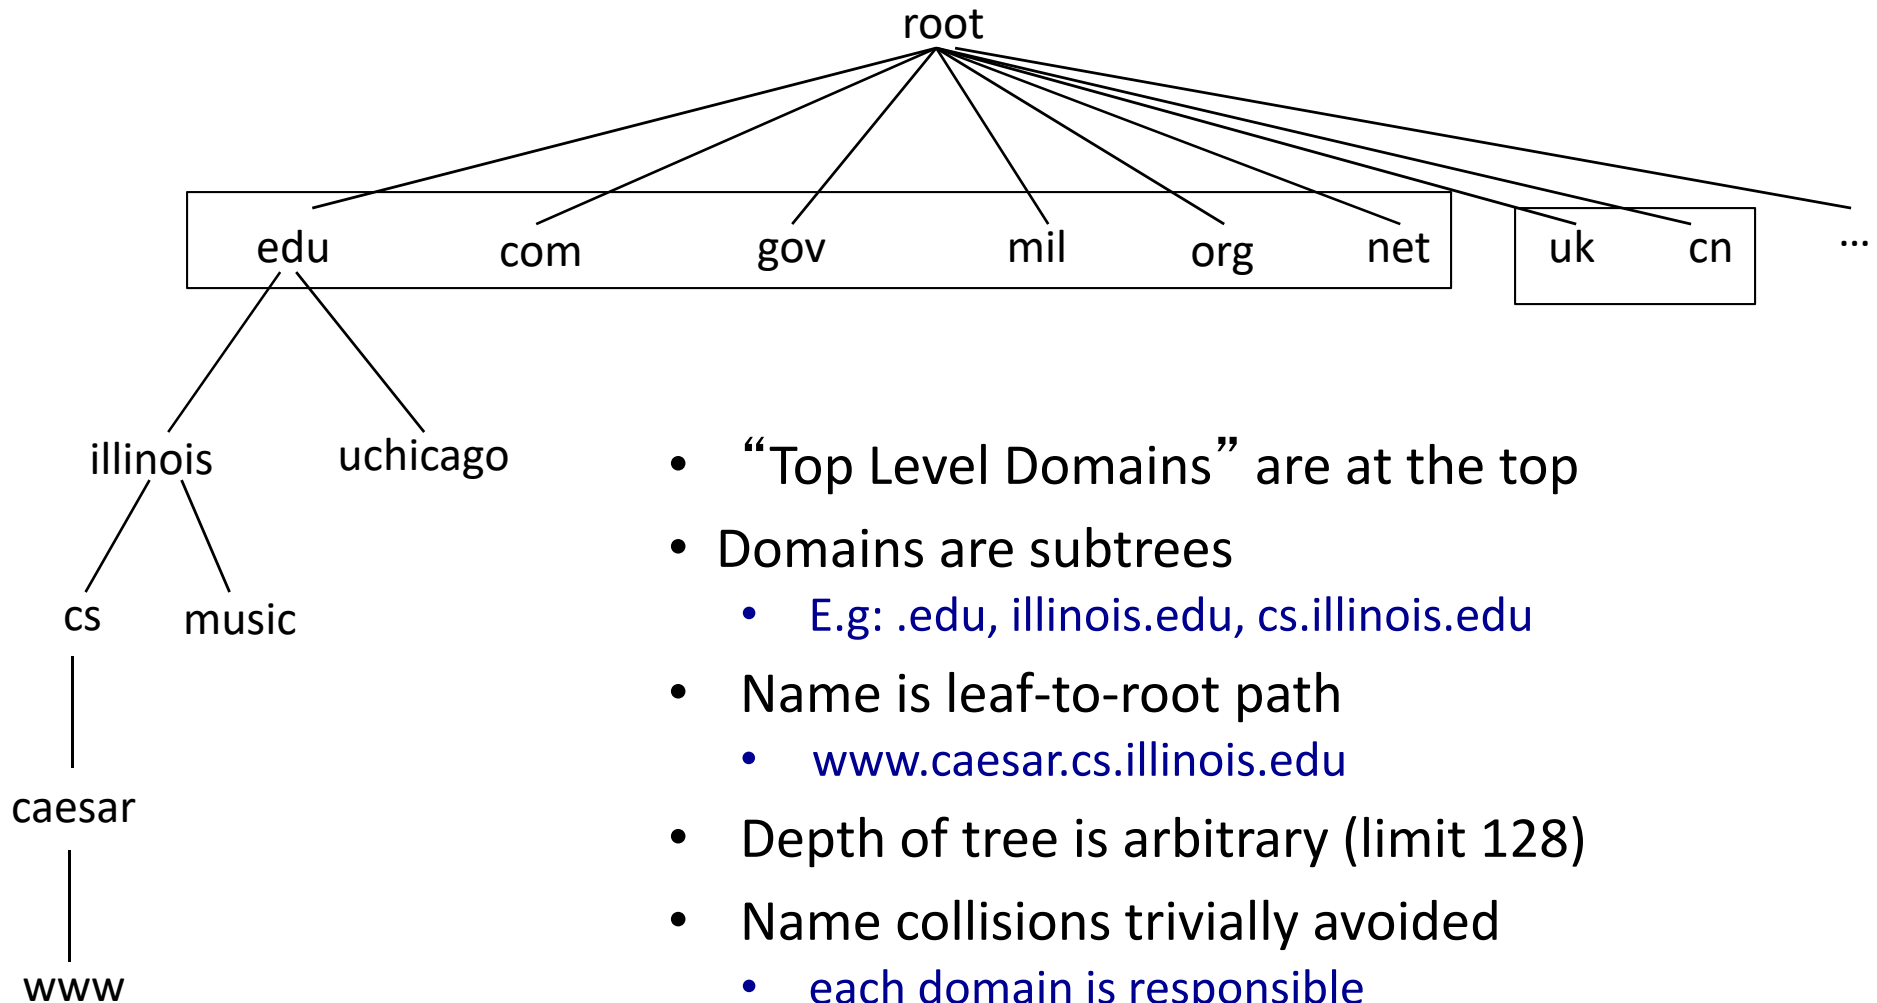
- No naming conflicts (uniqueness)
- Scalable
 - many names
 - (secondary) frequent updates
- Distributed, autonomous administration
 - Ability to update my own (machines') names
 - Don't have to track everybody's updates
- Highly available
- Lookups are fast
- Non-goal?: perfect consistency

Key idea: Hierarchy

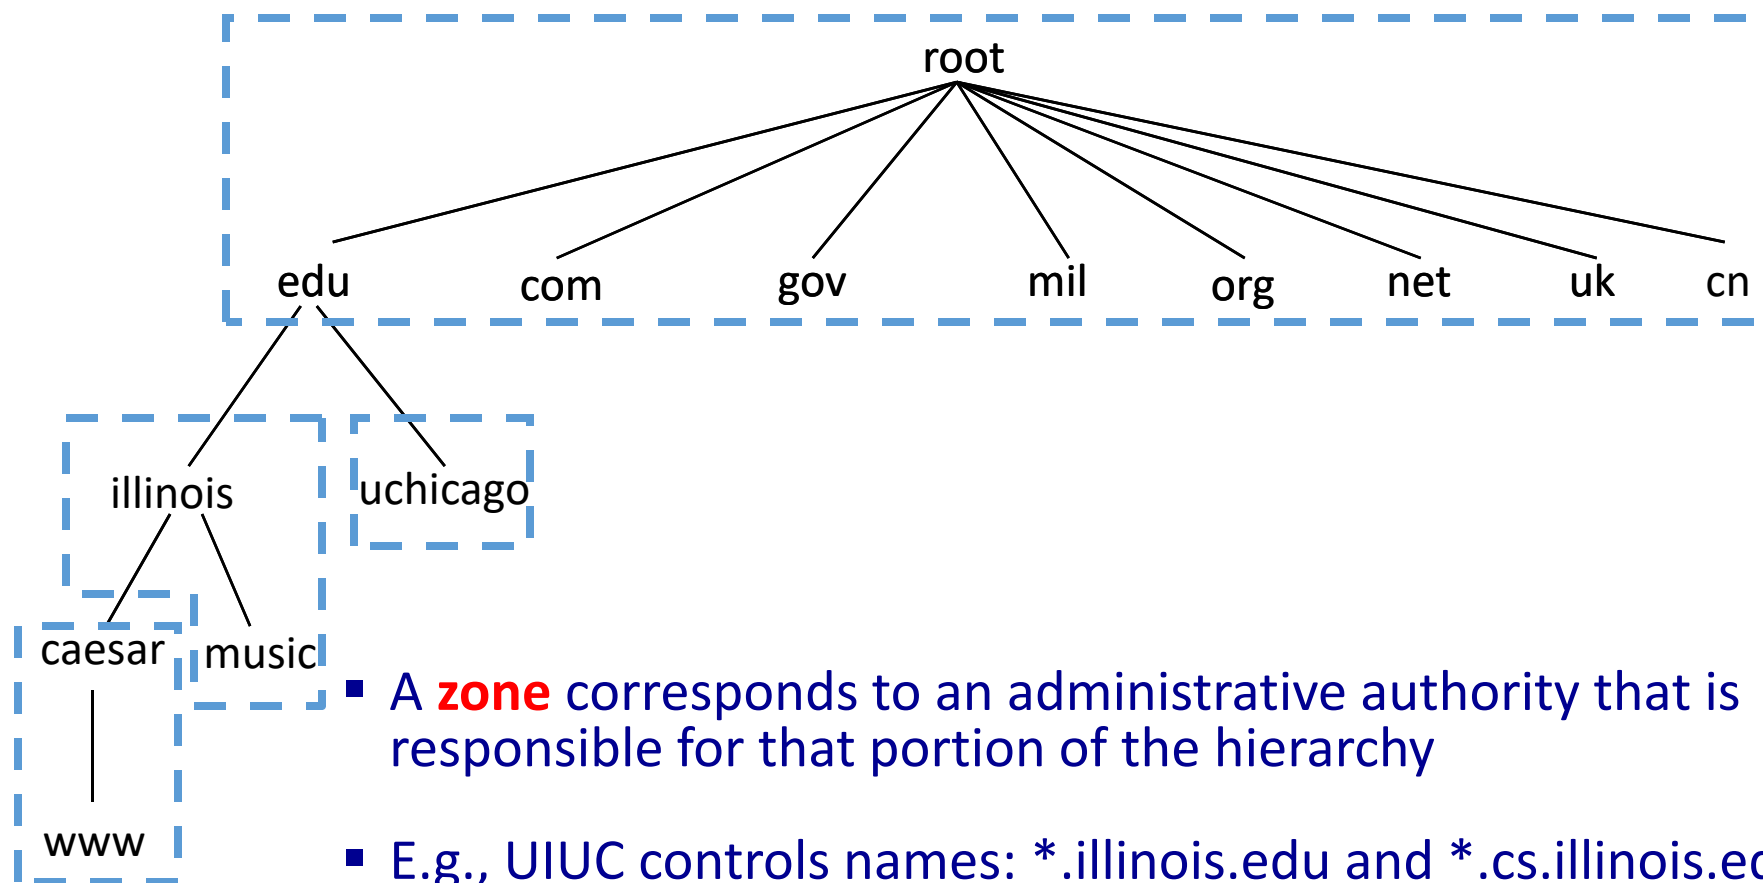
Three intertwined hierarchies

- Hierarchical namespace
 - As opposed to original flat namespace
- Hierarchically administered
 - As opposed to centralized
- (Distributed) hierarchy of servers
 - As opposed to centralized storage

Hierarchical Namespace



Hierarchical Administration



- A **zone** corresponds to an administrative authority that is responsible for that portion of the hierarchy
- E.g., UIUC controls names: *.illinois.edu and *.cs.illinois.edu
- E.g., CS controls names: *.cs.illinois.edu

Server Hierarchy

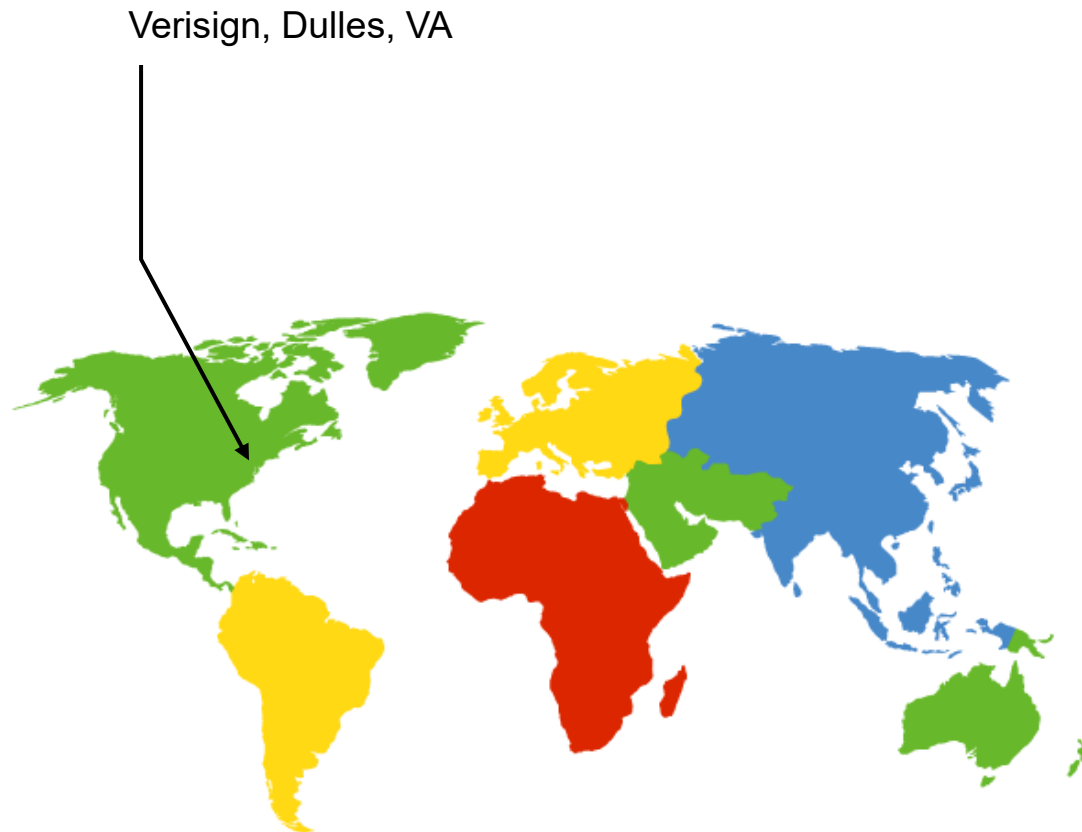
- Top of hierarchy: Root servers
 - Location hardwired into other servers
- Next Level: Top-level domain (TLD) servers
 - .com, .edu, etc.
 - Managed professionally
- Bottom Level: **Authoritative** DNS servers
 - Actually store the name-to-address mapping
 - Maintained by the corresponding administrative authority

Server Hierarchy

- Each server stores a (small!) subset of the total DNS database
- An authoritative DNS server stores “resource records” for all DNS names in the domain that it has authority for
- Each server needs to know other servers that are responsible for the other portions of the hierarchy
 - Every server knows the root
 - Root server knows about all top-level domains

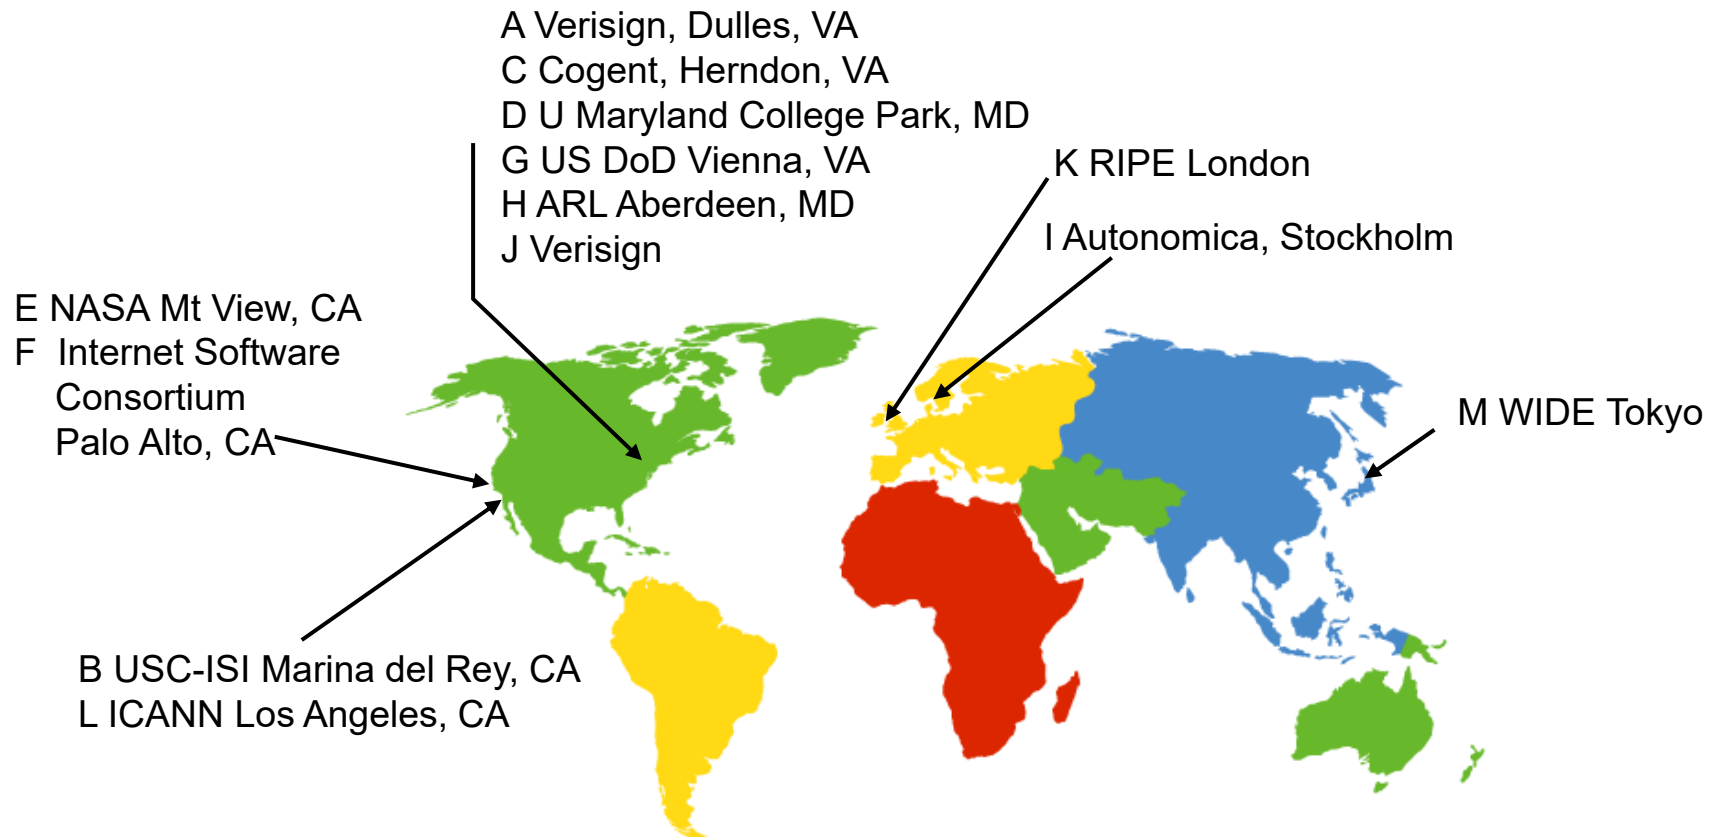
DNS Root

- Located in Virginia, USA
- How do we make the root scale?



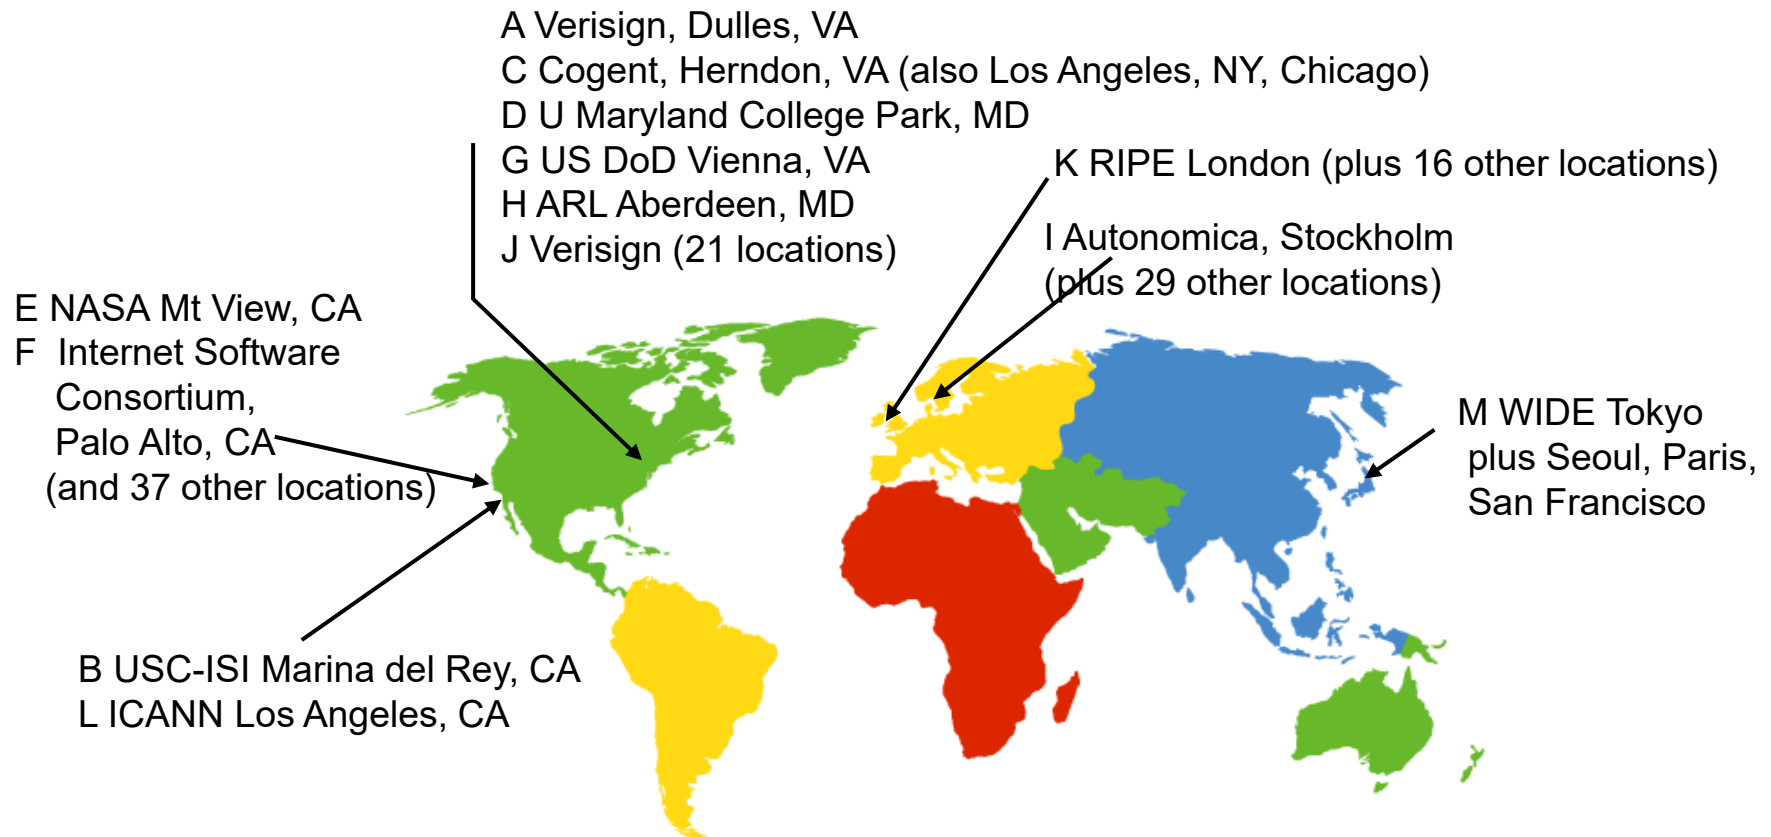
DNS Root Servers

- 13 root servers (labeled A-M; see <http://www.root-servers.org/>)



DNS Root Servers

- 13 root servers (labeled A-M; see <http://www.root-servers.org/>)
- Replicated via **any-casting**



Anycast in a nutshell

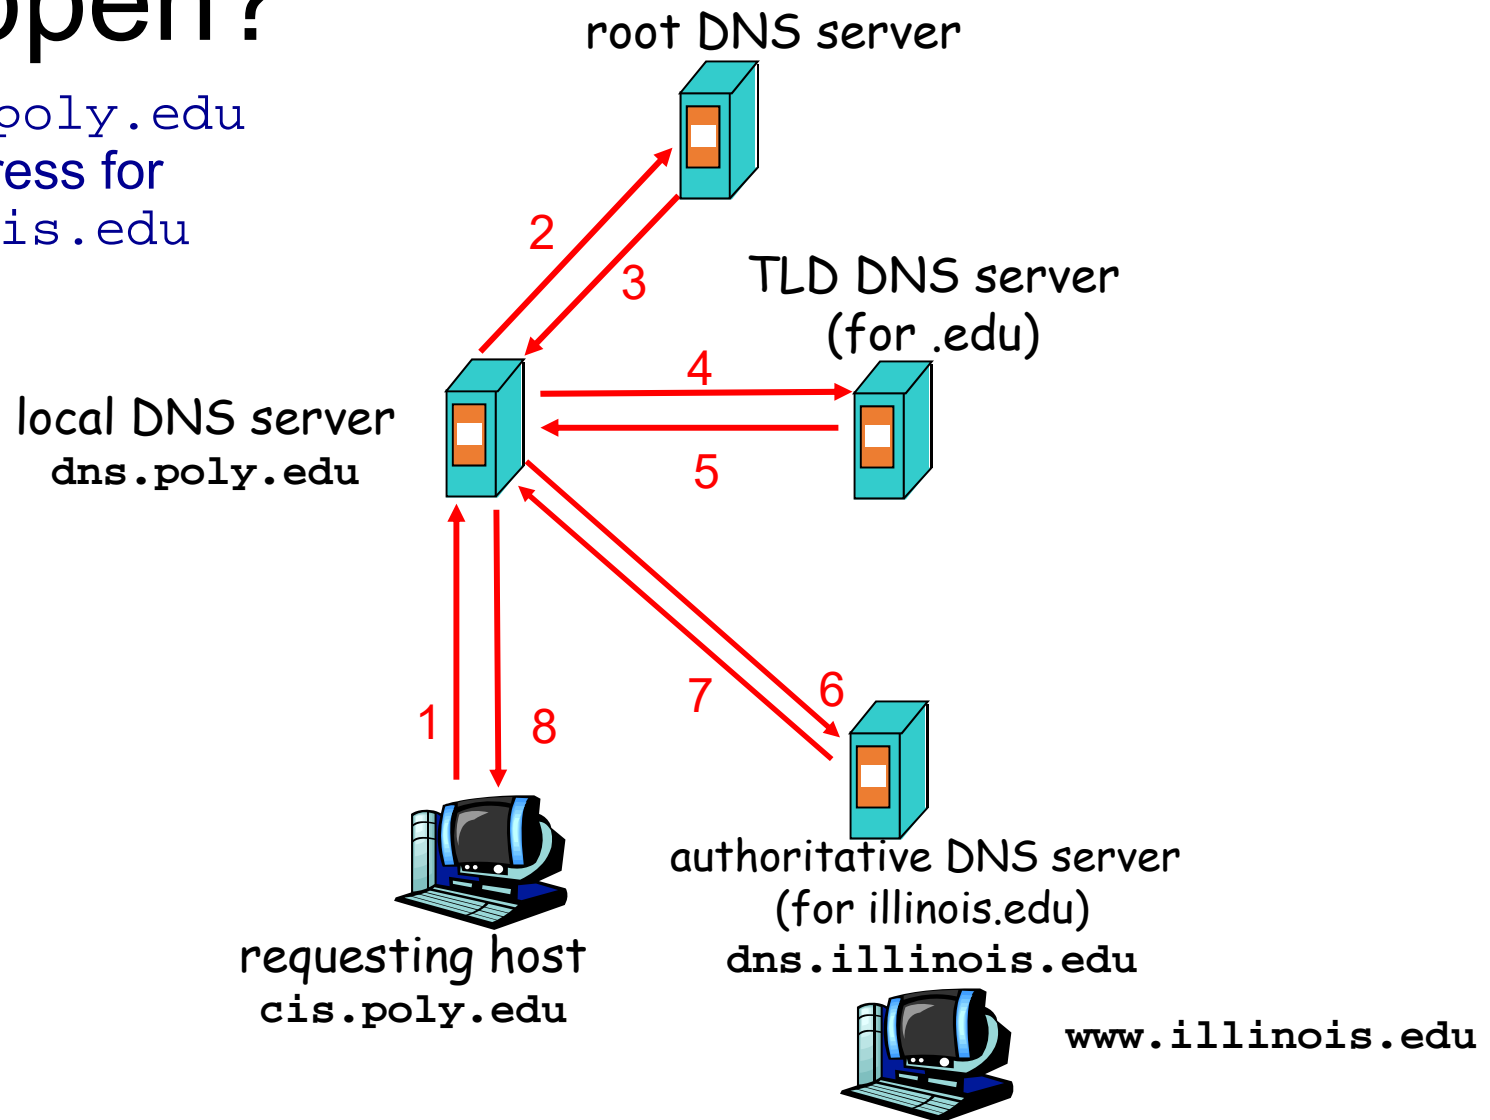
- Routing finds shortest paths to destination
- If several locations are given the same address, then the network will deliver the packet to the closest location with that address
- This is called “anycast”
 - Very robust
 - Requires no modification to routing algorithms

Using DNS (Client/App View)

- Two components
 - Local DNS servers
 - Resolver software on hosts
- Local DNS server (“default name server”)
 - Usually near the endhosts that use it
 - Hosts configured with local server address (e.g., /etc/resolv.conf) or learn server via a host configuration protocol (e.g., DHCP)
- Client application
 - Obtain DNS name (e.g., from URL)
 - Do **gethostbyname()** to trigger DNS request to its local DNS server


How Does Resolution Happen?

Host at `cis.poly.edu`
wants IP address for
`www.illinois.edu`



Recursive vs. Iterative Queries

- **Recursive** query
 - Ask server to get answer for you
 - E.g., request 1 and response 8
- **Iterative** query
 - Ask server who to ask next
 - E.g., all other request-response pairs
- Usually recursive is disabled (subject to DoS attacks)



DNS Records

- DNS info. stored as **resource records (RRs)**
 - RR is (name, value, type, TTL)
- Type = A: (\rightarrow Address)
 - name = hostname
 - value = IP address
- Type = NS: (\rightarrow Name Server)
 - name = domain
 - value = name of dns server for domain

DNS Records (cont'd)

- Type = CNAME: (*→ Canonical NAME*)
 - name = hostname
 - value = canonical name
- Type = MX: (*→ Mail eXchanger*)
 - name = domain in email address
 - value = canonical name(s) of mail server(s)

- Example:

NAME	TYPE	VALUE

bar.example.com.	CNAME	foo.example.com.
foo.example.com.	A	192.0.2.23

Inserting Resource Records into DNS

- Example: just created company “FooBar”
- Get a block of address space from ISP
 - Say 212.44.9.128/25
- Register **foobar.com** at registrar (e.g., Network Solutions)
 - Provide registrar with names and IP addresses of your authoritative name server (primary and secondary)
 - Registrar inserts RR pairs into the .com TLD server:
 - (**foobar.com**, **dns1.foobar.com**, NS)
 - (**dns1.foobar.com**, **212.44.9.129**, A)
- You store appropriate records in your server **dns1.foobar.com**:
 - e.g., type A record for **www.foobar.com**
 - e.g., type MX record for **foobar.com**

DNS Protocol

- Query and Reply messages; both with the same message format
 - header: identifier, flags, *etc.*
 - plus resource records
 - *see text/section for details*
- Client--server interaction on UDP Port 53
 - Spec supports TCP too, but not always implemented

Goals -- how are we doing?

- No naming conflicts (uniqueness)
- Scalable
- Distributed, autonomous administration
- Highly available
- Fast lookups

DNS Caching

- Performing all these queries takes time
 - And all this **before** actual communication takes place
 - E.g., 1-second latency before starting Web download
- **Caching** can greatly reduce overhead
 - The top-level servers very rarely change
 - Popular sites (e.g., www.cnn.com) visited often
 - Local DNS server often has the information cached
- How DNS caching works
 - DNS servers cache responses to queries
 - Responses include a “**time to live**” (TTL) field
 - Server deletes cached entry after TTL expires

Negative Caching

- Remember things that don't work
 - Misspellings like *www.cnn.comm* and *www.cnnn.com*
 - These can take a long time to fail the first time
 - Good to remember that they don't work
 - ... so the failure takes less time the next time around
- But: negative caching is optional
 - And not widely implemented

Reliability

- DNS servers are **replicated** (primary/secondary)
 - Name service available if at least one replica is up
 - Queries can be load-balanced between replicas
- Usually, UDP used for queries
 - Need reliability: must implement this on top of UDP
- Try alternate servers on timeout
 - **Exponential backoff** when retrying same server
- Same identifier for all queries
 - Don't care which server responds

Important Properties of DNS

Administrative delegation and hierarchy results in:

- Easy unique naming
- “Fate sharing” for network failures
- Reasonable trust model
- Caching lends scalability, performance

DNS provides Indirection

- Addresses can **change** underneath
 - Move `www.cnn.com` to `4.125.91.21`
 - Humans/Apps should be unaffected
- Name could map to **multiple** IP addresses
 - Enables
 - Load-balancing
 - Reducing latency by picking nearby servers
- **Multiple names** for the same address
 - E.g., many services (mail, www, ftp) on same machine
 - E.g., aliases like `www.cnn.com` and `cnn.com`
- But, this flexibility applies only within domain!

DNSSEC

- What if a DNS server gets compromised/intercepted?
 - Can redirect users to other sites
 - E.g., fake bank site that learns your password
 - “Cache poisoning” – filling DNS caches with bad information
- Solution: DNSSEC
 - Digitally sign IP-to-name mappings
 - Signature is cryptographic hash of DNS data
 - If signatures can’t be verified, error to user
 - Doesn’t mitigate DoS or provide data confidentiality

DNSSEC Lookup

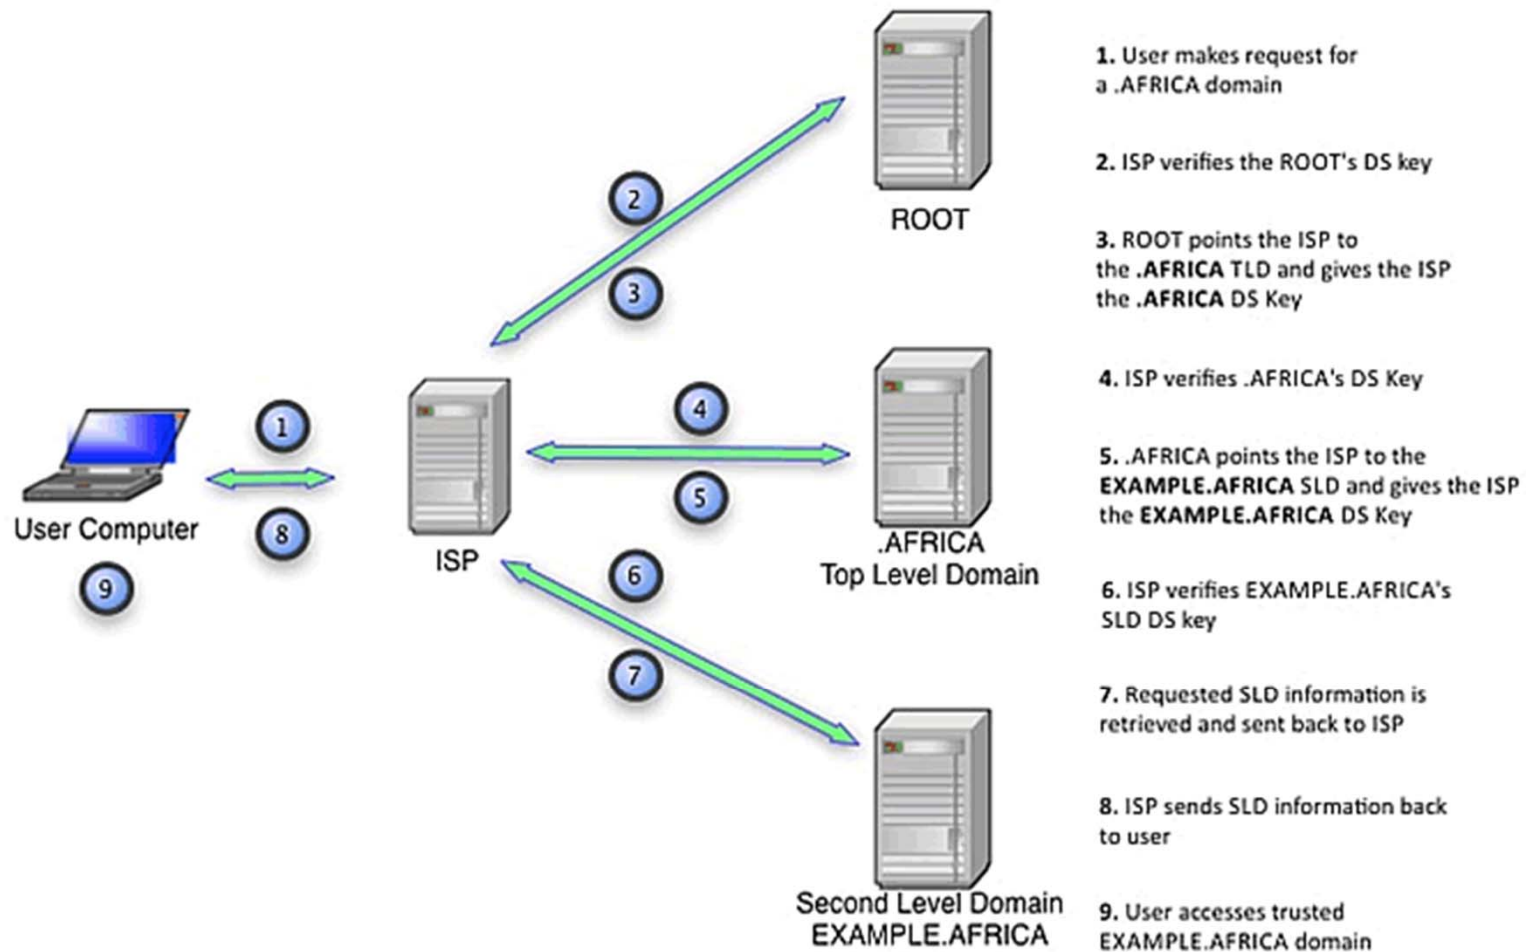
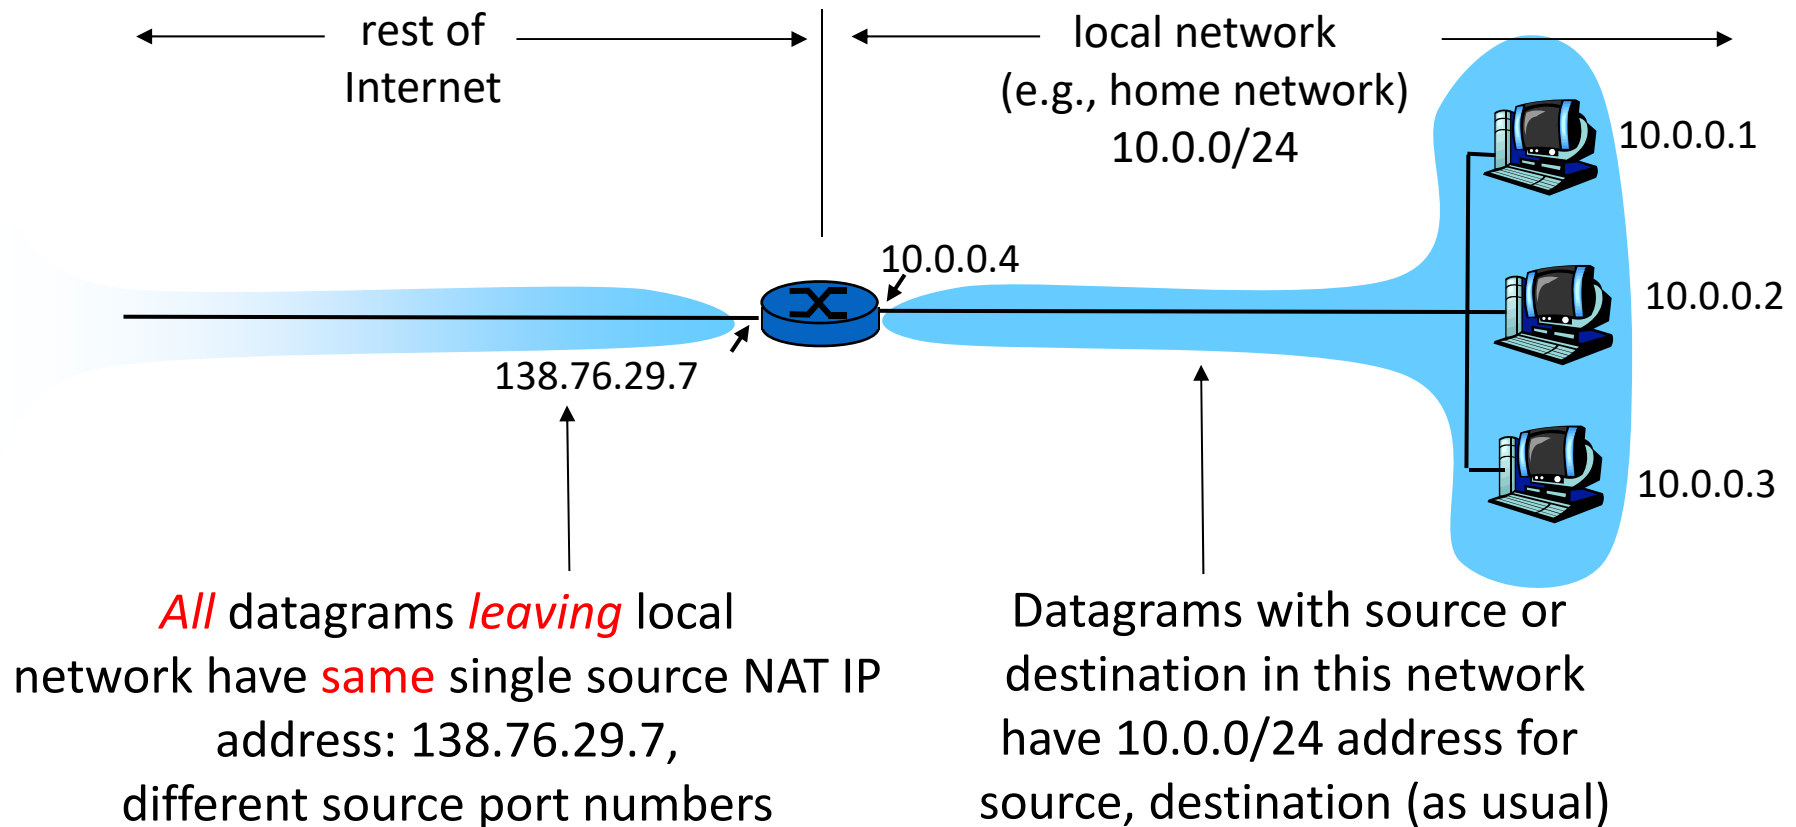


Figure 1: Example of Trust Chain Using DNSSEC

NAT

NAT: Network Address Translation



NAT: Network Address Translation

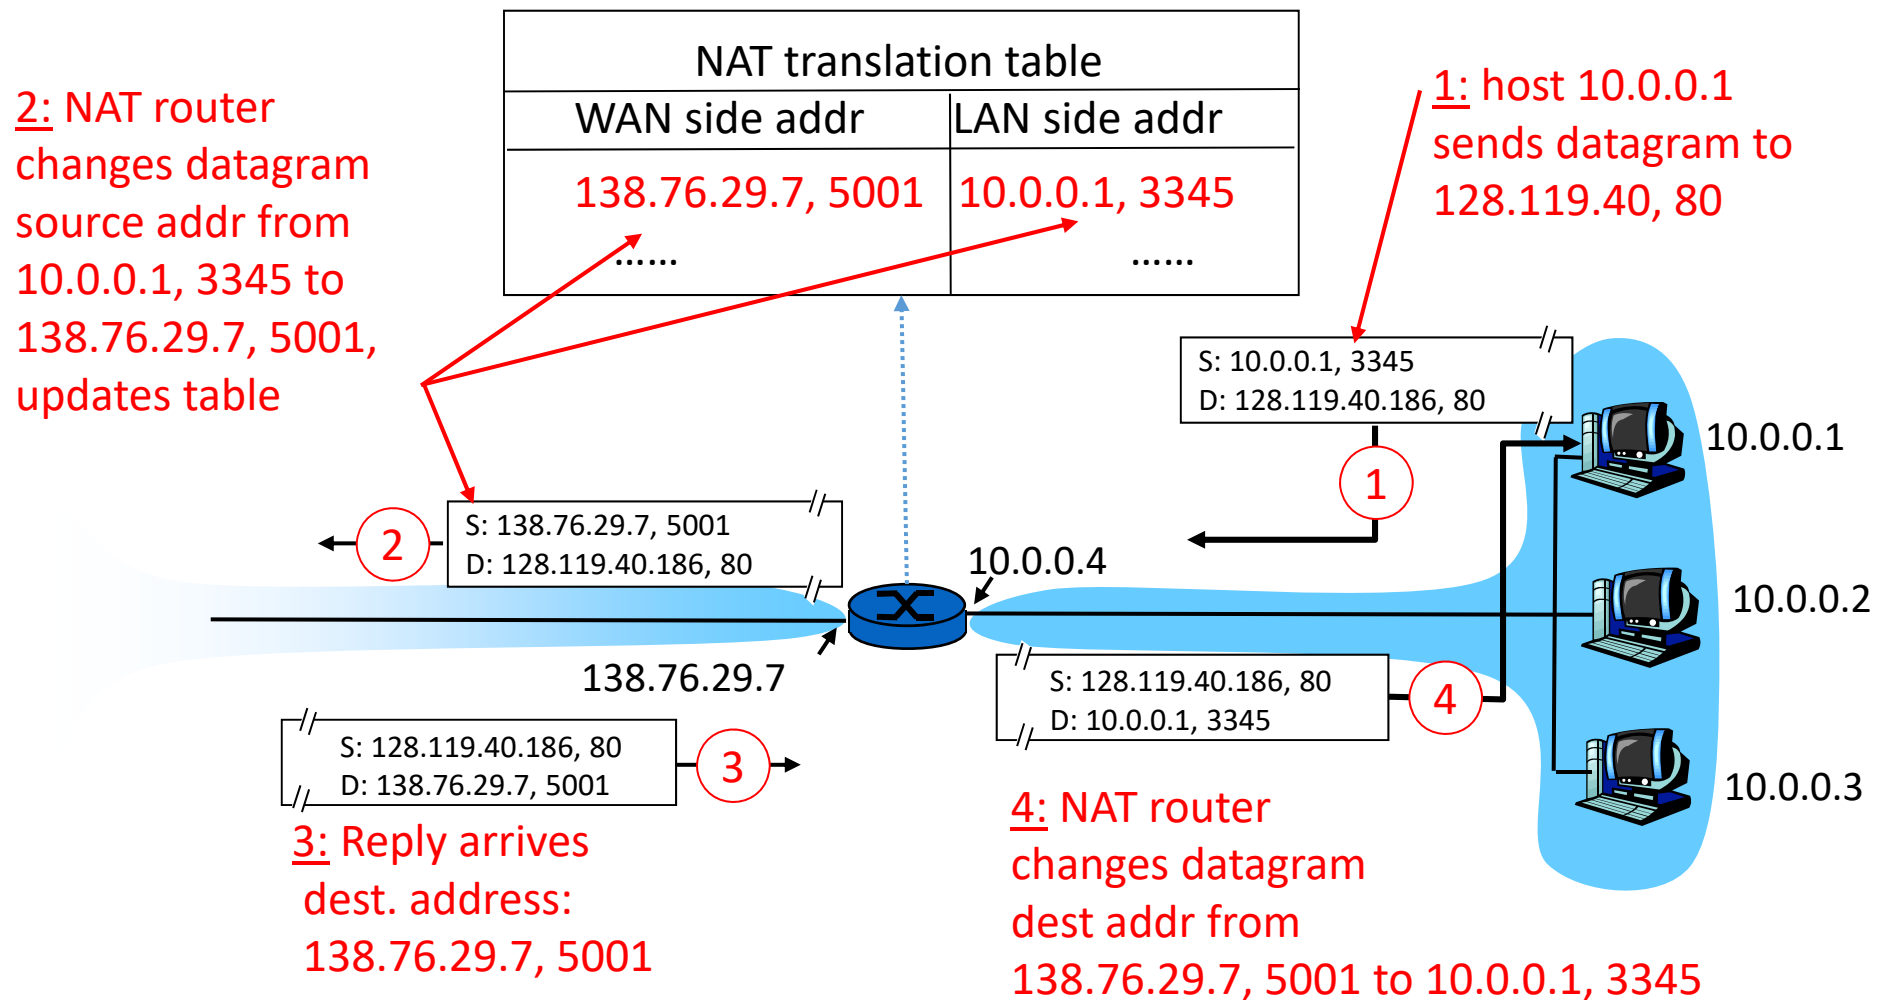
- **Motivation:** local network uses just one IP address as far as outside world is concerned:
 - no need to be allocated range of addresses from ISP: - just one IP address is used for all devices
 - can change addresses of devices in local network without notifying outside world
 - can change ISP without changing addresses of devices in local network
 - devices inside local net not explicitly addressable, visible by outside world (a security plus).

NAT: Network Address Translation

Implementation: NAT router must:

- *outgoing datagrams: replace* (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)
... remote clients/servers will respond using (NAT IP address, new port #) as destination addr.
- *remember (in NAT translation table)* every (source IP address, port #) to (NAT IP address, new port #) translation pair
- *incoming datagrams: replace* (NAT IP address, new port #) in dest fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

NAT: Network Address Translation



NAT: Network Address Translation

- 16-bit port-number field:
 - 60,000 simultaneous connections with a single LAN-side address!
- NAT is (was?) controversial:
 - routers should only process up to layer 3
 - violates end-to-end argument
 - NAT possibility must be taken into account by app designers, eg, P2P applications
 - address shortage should instead be solved by IPv6

Types of NAT

- PAT: Port-Address Translation
 - Uses port field in addition to address field
- Restricted NAT
 - An external host H_e can only contact an internal host H_i if H_i has previously sent packets to H_e
 - See also “hole-punching” source sends packet to create association, informs association to peer (may have to guess port assignment alg)
- Port forwarding
 - Install static mapping in NAT, to make services on internal network publicly reachable
 - E.g., port 80 on NAT address gets mapped to machine X, port 8080
- Carrier-grade NAT
 - NAT run by an ISP, as opposed to an enterprise/end-user
 - NAT444: cust private nw → provider private nw → public nw