

# CS 425 / ECE 428

## Distributed Systems

### Fall 2019

Indranil Gupta (Indy)

*Lecture 12: Time and Ordering*

# Why Synchronization?

- **You want to catch a bus at 6.05 pm, but your watch is off by 15 minutes**
  - What if your watch is Late by 15 minutes?
    - You'll miss the bus!
  - What if your watch is Fast by 15 minutes?
    - You'll end up unfairly waiting for a longer time than you intended
- **Time synchronization is required for both**
  - Correctness
  - Fairness

# Synchronization In The Cloud

- Cloud airline reservation system
- Server A receives a client request to purchase last ticket on flight ABC 123.
- Server A timestamps purchase using local clock 9h:15m:32.45s, and logs it. Replies ok to client.
- That was the last seat. Server A sends message to Server B saying “flight full.”
- B enters “Flight ABC 123 full” + its own local clock value (which reads 9h:10m:10.11s) into its log.
- Server C queries A’s and B’s logs. Is confused that a client purchased a ticket at A after the flight became full at B.
- This may lead to further incorrect actions by C

# Why is it Challenging?

- **End hosts in Internet-based systems (like clouds)**
  - Each have their own clocks
  - Unlike processors (CPUs) within one server or workstation which share a system clock
- **Processes in Internet-based systems follow an *asynchronous* system model**
  - No bounds on
    - Message delays
    - Processing delays
  - Unlike multi-processor (or parallel) systems which follow a *synchronous* system model

# Some Definitions

- An Asynchronous Distributed System consists of a number of **processes**.
- Each process has a **state** (values of variables).
- Each process takes **actions** to change its state, which may be an **instruction** or a communication action (**send**, **receive**).
- An **event** is the occurrence of an action.
- Each process has a local clock – events *within* a process can be assigned **timestamps**, and thus ordered linearly.
- But – in a distributed system, we also need to know the time order of events *across* different processes.

# Clock Skew vs. Clock Drift

- Each process (running at some end host) has its own clock.
- When comparing two clocks at two processes:
  - Clock **Skew** = Relative Difference in clock *values* of two processes
    - Like distance between two vehicles on a road
  - Clock **Drift** = Relative Difference in clock *frequencies (rates)* of two processes
    - Like difference in speeds of two vehicles on the road
- A non-zero clock skew implies clocks are not synchronized.
- A non-zero clock drift causes skew to increase (eventually).
  - If faster vehicle is ahead, it will drift away
  - If faster vehicle is behind, it will catch up and then drift away

# How often to Synchronize?

- Maximum Drift Rate (MDR) of a clock
- Absolute MDR is defined relative to Coordinated Universal Time (UTC). UTC is the “correct” time at any point of time.
  - MDR of a process depends on the environment.
- Max drift rate between two clocks with similar MDR is  $2 * \text{MDR}$
- Given a maximum acceptable skew  $M$  between any pair of clocks, need to synchronize at least once every:  $M / (2 * \text{MDR})$  time units
  - Since time = distance/speed

# External vs Internal Synchronization

- **Consider a group of processes**
- **External Synchronization**
  - Each process  $C(i)$ 's clock is within a bound  $D$  of a well-known clock  $S$  external to the group
  - $|C(i) - S| < D$  at all times
  - External clock may be connected to UTC (Universal Coordinated Time) or an atomic clock
  - E.g., Cristian's algorithm, NTP
- **Internal Synchronization**
  - Every pair of processes in group have clocks within bound  $D$
  - $|C(i) - C(j)| < D$  at all times and for all processes  $i, j$
  - E.g., Berkeley algorithm



# External vs Internal Synchronization (2)

- **External Synchronization with  $D \Rightarrow$  Internal Synchronization with  $2 \cdot D$**
- **Internal Synchronization does not imply External Synchronization**
  - In fact, the entire system may drift away from the external clock  $S$ !

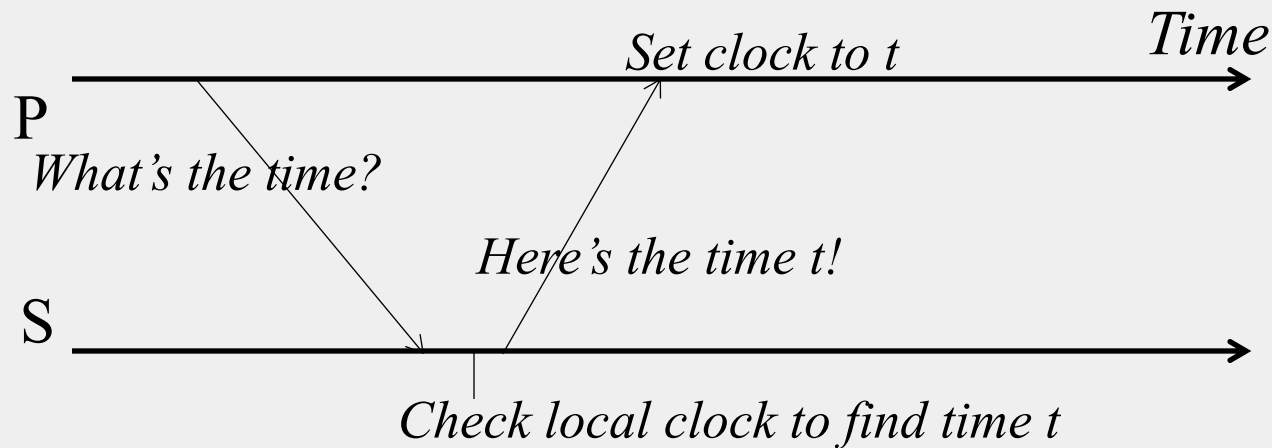
# Next

- Algorithms for Clock Synchronization

# Cristian's Algorithm

# Basics

- **External time synchronization**
- **All processes P synchronize with a time server S**

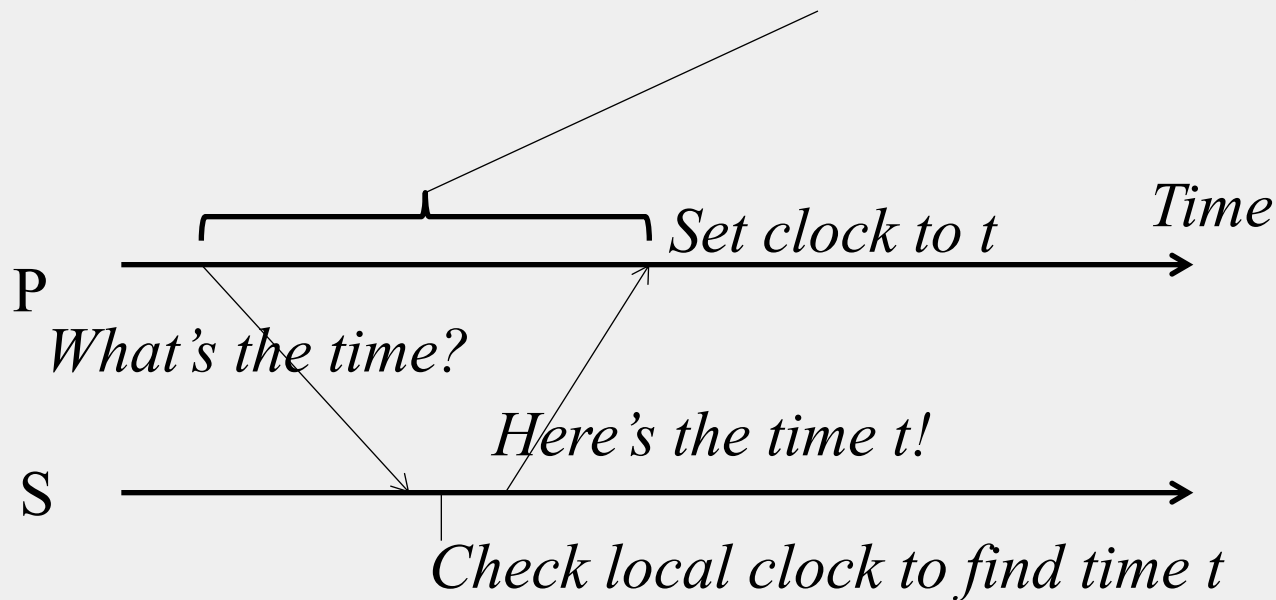


# What's Wrong

- By the time response message is received at P, time has moved on
- P's time set to  $t$  is inaccurate!
- Inaccuracy a function of message latencies
- Since latencies unbounded in an asynchronous system, the inaccuracy cannot be bounded


# Cristian's Algorithm

- P measures the round-trip-time RTT of message exchange




# Cristian's Algorithm (2)

- **P measures the round-trip-time RTT of message exchange**
- **Suppose we know the minimum  $P \rightarrow S$  latency  $\min1$**
- **And the minimum  $S \rightarrow P$  latency  $\min2$** 
  - $\min1$  and  $\min2$  depend on Operating system overhead to buffer messages, TCP time to queue messages, etc.



# Cristian's Algorithm (3)


- P measures the round-trip-time RTT of message exchange
- Suppose we know the minimum  $P \rightarrow S$  latency  $\min1$
- And the minimum  $S \rightarrow P$  latency  $\min2$ 
  - $\min1$  and  $\min2$  depend on Operating system overhead to buffer messages, TCP time to queue messages, etc.
- The actual time at P when it receives response is between  $[t+\min2, t+\text{RTT}-\min1]$





# Cristian's Algorithm (4)

- The actual time at P when it receives response is between  $[t + \min_2, t + \text{RTT} - \min_1]$
- P sets its time to halfway through this interval
  - To:  $t + (\text{RTT} + \min_2 - \min_1) / 2$
- Error is at most  $(\text{RTT} - \min_2 - \min_1) / 2$ 
  - Bounded!




# Gotchas

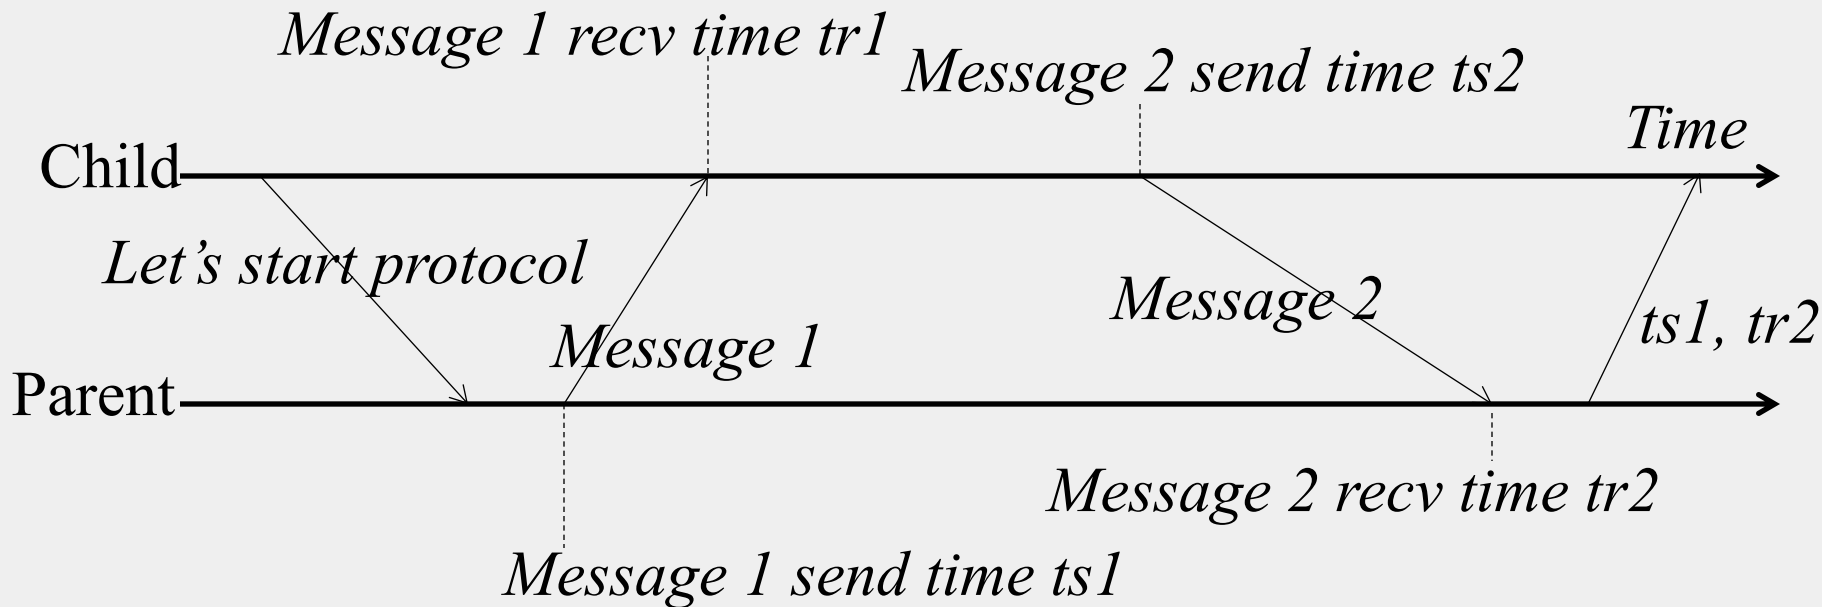
- **Allowed to increase clock value but should never decrease clock value**
  - May violate ordering of events within the same process
- **Allowed to increase or decrease speed of clock**
- **If error is too high, take multiple readings and average them**

# NTP = Network Time Protocol

- NTP Servers organized in a tree
- Each Client = a leaf of tree
- Each node synchronizes with its tree parent



# NTP Protocol



# What the Child Does

- Child calculates *offset* between its clock and parent's clock
- Uses *ts1*, *tr1*, *ts2*, *tr2*
- Offset is calculated as

$$o = (tr1 - tr2 + ts2 - ts1)/2$$

# Why $o = (tr1 - tr2 + ts2 - ts1)/2$ ?

- Offset  $o = (tr1 - tr2 + ts2 - ts1)/2$
- Let's calculate the error
- Suppose real offset is  $oreal$ 
  - Child is ahead of parent by  $oreal$
  - Parent is ahead of child by  $-oreal$
- Suppose one-way latency of Message 1 is  $L1$  ( $L2$  for Message 2)
- No one knows  $L1$  or  $L2$ !
- Then

$$tr1 = ts1 + L1 + oreal$$

$$tr2 = ts2 + L2 - oreal$$

# Why $o = (tr1 - tr2 + ts2 - ts1)/2$ ? (2)

- **Then**

$$tr1 = ts1 + L1 + o_{real}$$

$$tr2 = ts2 + L2 - o_{real}$$

- **Subtracting second equation from the first**

$$o_{real} = (tr1 - tr2 + ts2 - ts1)/2 + (L2 - L1)/2$$

$$\Rightarrow o_{real} = o + (L2 - L1)/2$$

$$\Rightarrow |o_{real} - o| < |(L2 - L1)/2| < |(L2 + L1)/2|$$

– Thus, the error is bounded by the round-trip-time

# And yet...

- **We still have a non-zero error!**
- **We just can't seem to get rid of error**
  - Can't, as long as message latencies are non-zero
- **Can we avoid synchronizing clocks altogether, and still be able to order events?**



# Lamport Timestamps

# Ordering Events in a Distributed System

- To order events across processes, trying to sync clocks is one approach
- What if we instead assigned timestamps to events that were not *absolute* time?
- As long as these timestamps obey *causality*, that would work

If an event A causally happens before another event B, then  $\text{timestamp}(A) < \text{timestamp}(B)$

Humans use causality all the time

E.g., I enter a house only after I unlock it

E.g., You receive a letter only after I send it


# Logical (or Lamport) Ordering

- Proposed by Leslie Lamport in the 1970s
- Used in almost all distributed systems since then
- Almost all cloud computing systems use some form of logical ordering of events

# Logical (or Lamport) Ordering(2)

- Define a logical relation *Happens-Before* among pairs of events
- *Happens-Before* denoted as  $\rightarrow$
- Three rules
  1. On the same process:  $a \rightarrow b$ , if  $time(a) < time(b)$  (using the local clock)
  2. If p1 sends  $m$  to p2:  $send(m) \rightarrow receive(m)$
  3. (Transitivity) If  $a \rightarrow b$  and  $b \rightarrow c$  then  $a \rightarrow c$
- Creates a *partial order* among events
  - Not all events related to each other via  $\rightarrow$


# Example



While P1 and P3 each have an event labeled E, these are different events as they occur at different processes.




# Happens-Before



- $A \rightarrow B$
- $B \rightarrow F$
- $A \rightarrow F$

# Happens-Before (2)




- P3.
- $H \rightarrow G$
  - $F \rightarrow J$
  - $H \rightarrow J$
  - $C \rightarrow J$

# In practice: Lamport timestamps

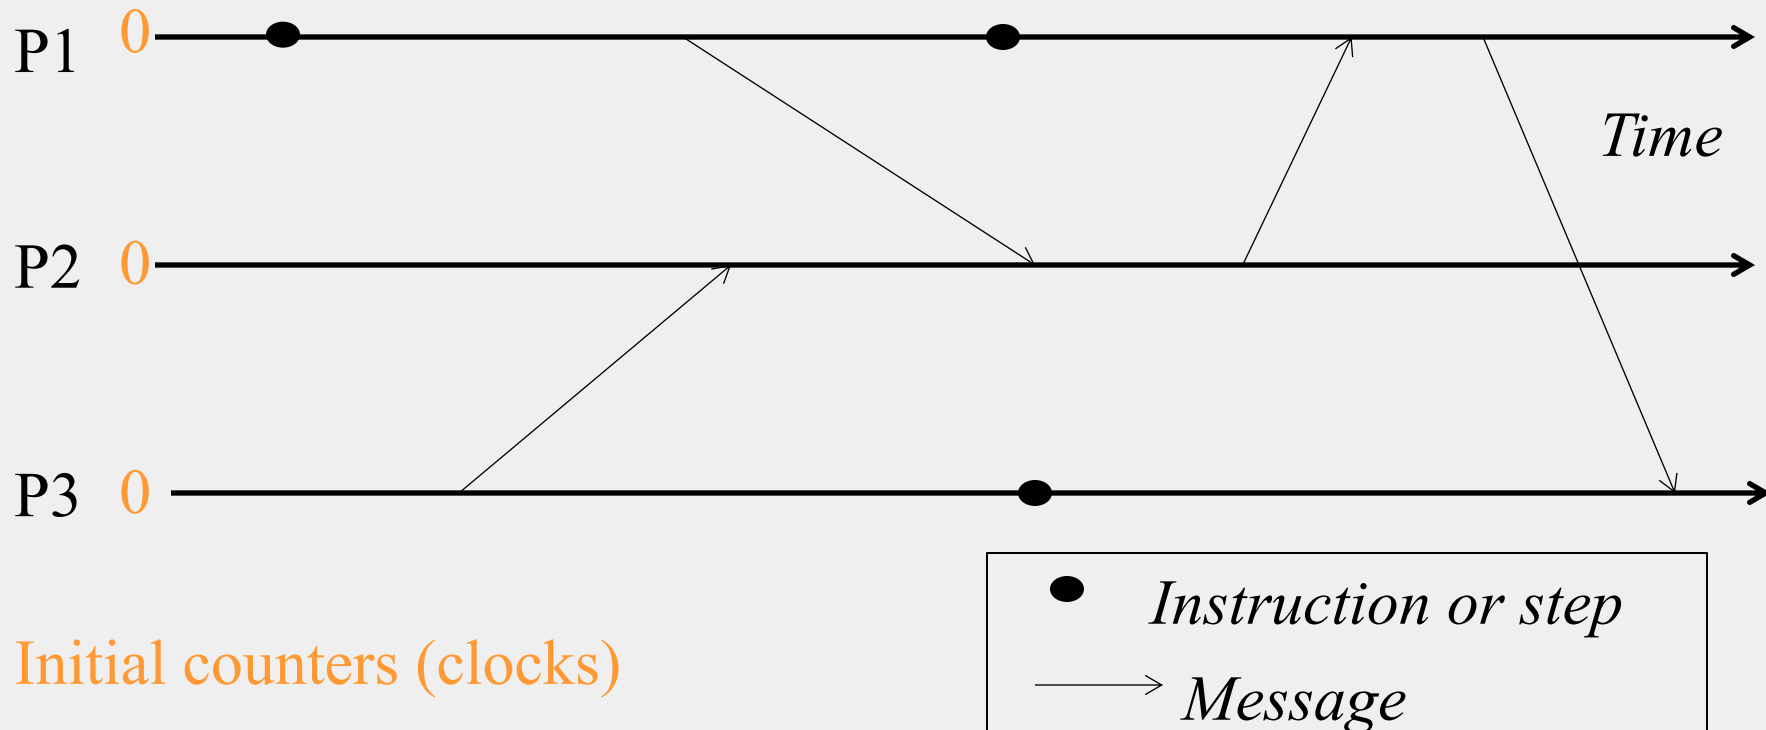
- **Goal: Assign logical (Lamport) timestamp to each event**
- **Timestamps obey causality**
- **Rules**
  - Each process uses a local counter (clock) which is an integer
    - initial value of counter is zero
  - A process increments its counter when a **send** or an **instruction** happens at it. The counter is assigned to the event as its timestamp.
  - A **send (message)** event carries its timestamp
  - For a **receive (message)** event the counter is updated by
$$\max(\text{local clock}, \text{message timestamp}) + 1$$



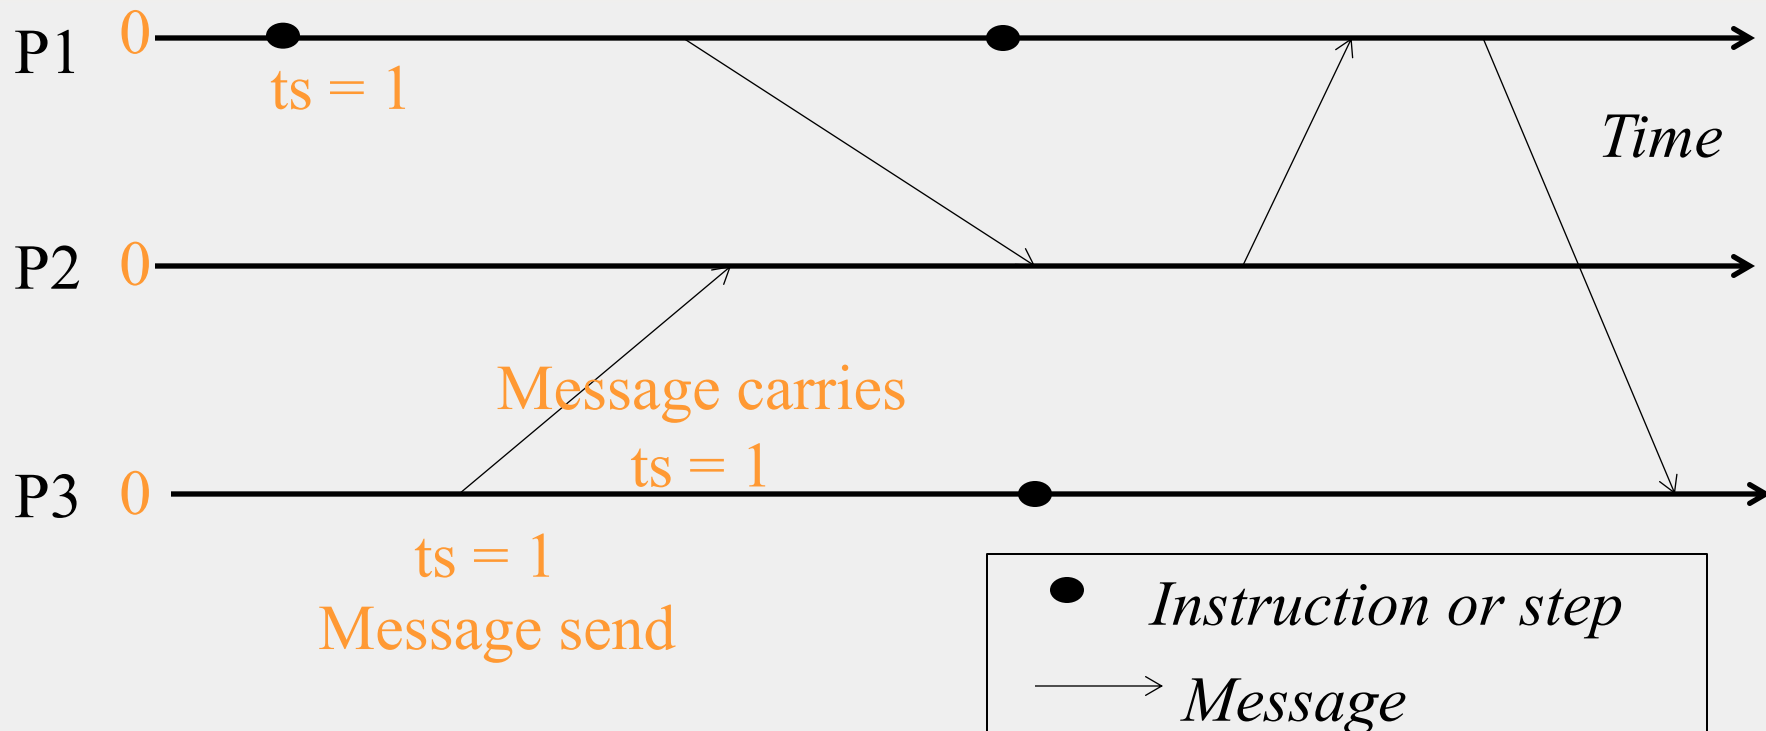
# Example



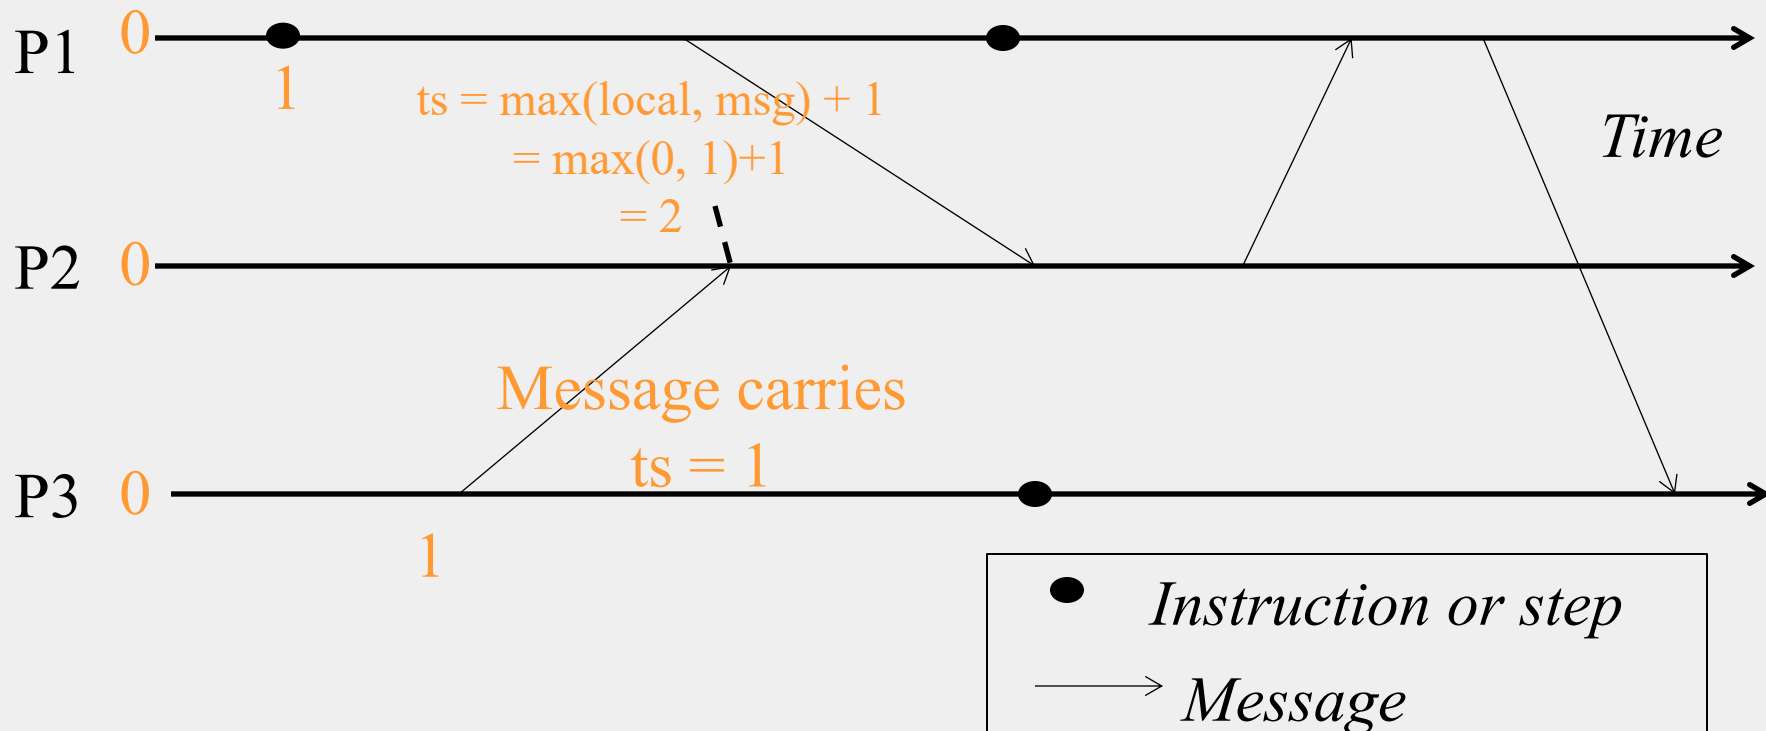
# Lamport Timestamps



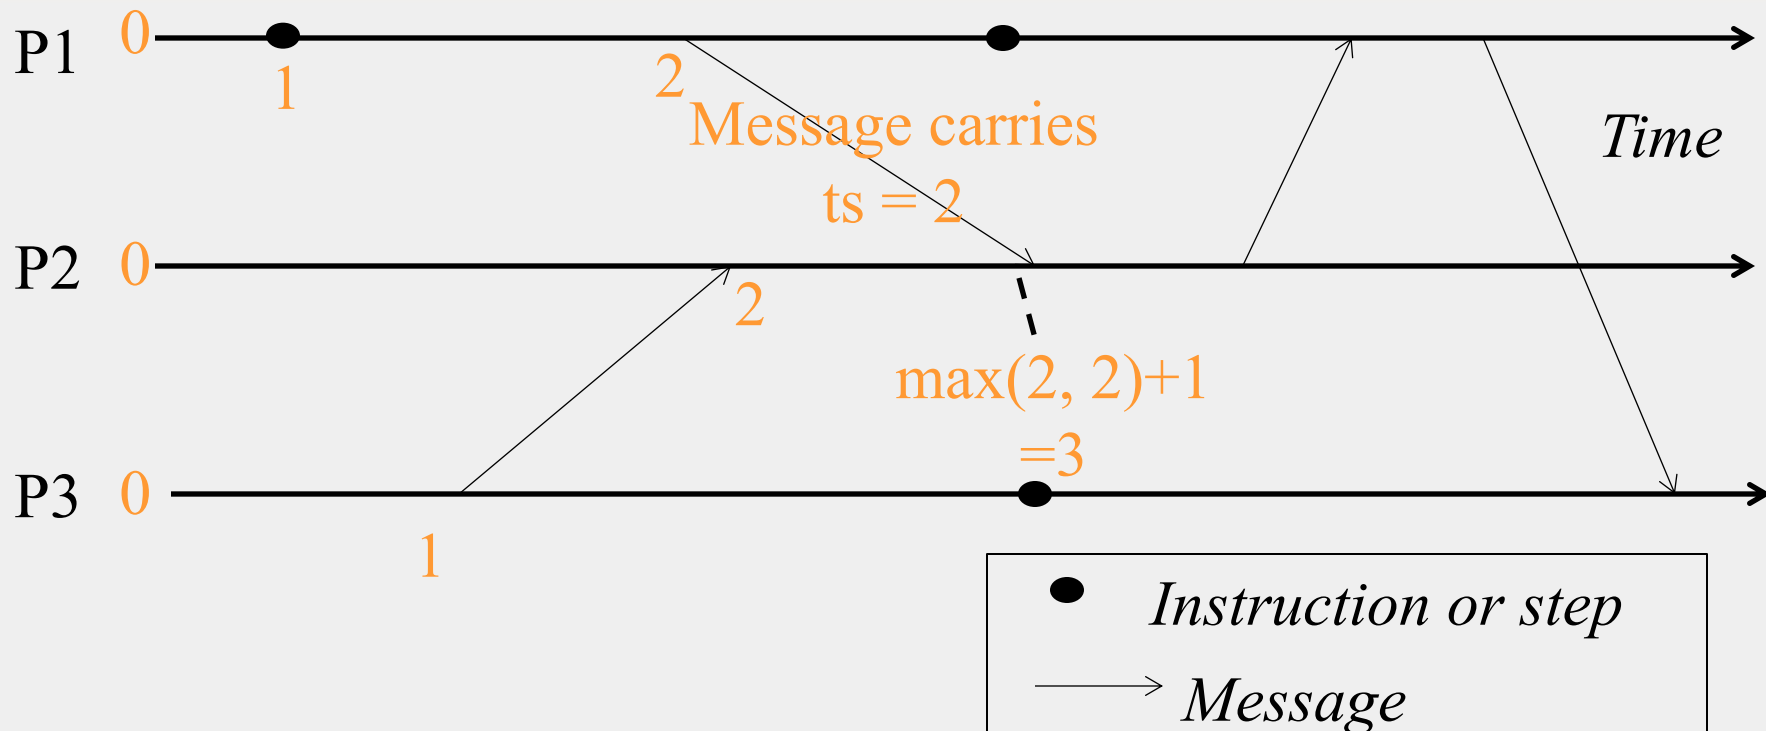
# Lamport Timestamps



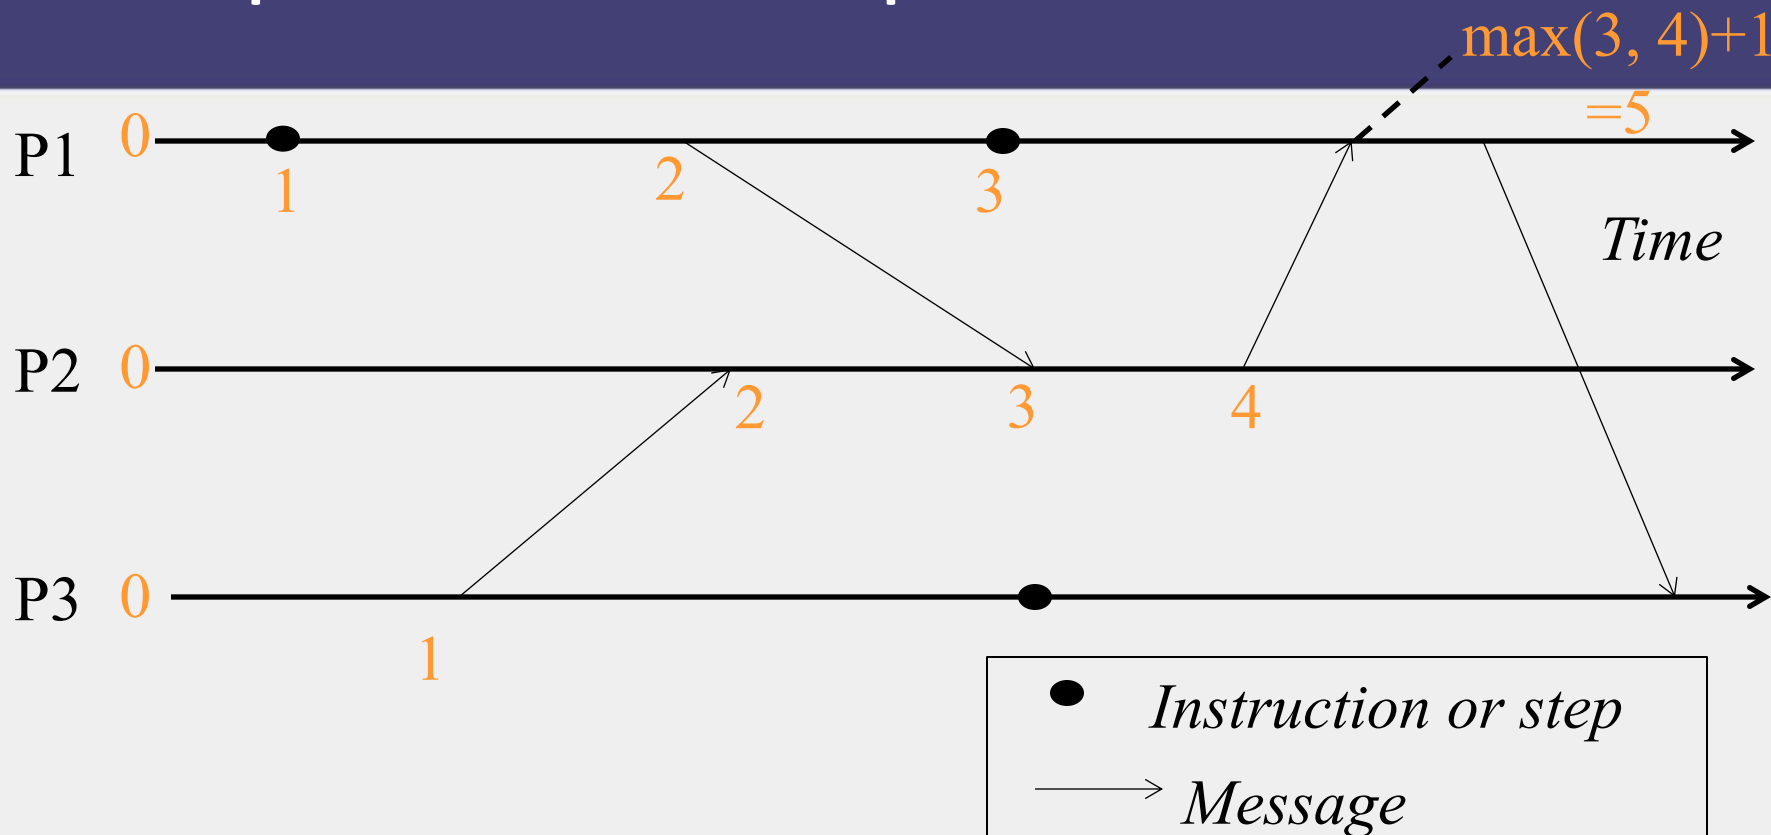
# Lamport Timestamps




# Lamport Timestamps




# Lamport Timestamps



# Lamport Timestamps




# Obeying Causality






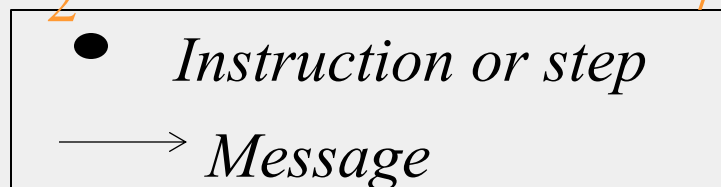
# Obeying Causality (2)



# Not always implying Causality



- $? C \rightarrow F ? :: 3 = 3$
- $? H \rightarrow C ? :: 1 < 3$
- (C, F) and (H, C) are pairs of concurrent events



# Concurrent Events

- **A pair of concurrent events doesn't have a causal path from one event to another (either way, in the pair)**
- **Lamport timestamps not guaranteed to be ordered or unequal for concurrent events**
- **Ok, since concurrent events are not causality related!**
- **Remember**

$E1 \rightarrow E2 \Rightarrow \text{timestamp}(E1) < \text{timestamp}(E2)$ , **BUT**

$\text{timestamp}(E1) < \text{timestamp}(E2) \Rightarrow$

$\{E1 \rightarrow E2\} \text{ OR } \{E1 \text{ and } E2 \text{ concurrent}\}$

# Next

- Can we have causal or logical timestamps from which we can tell if two events are concurrent or causally related?

# Vector Timestamps

- Used in key-value stores like Riak
- Each process uses a vector of integer clocks
- Suppose there are  $N$  processes in the group  $1 \dots N$
- Each vector has  $N$  elements
- Process  $i$  maintains vector  $V_i[1 \dots N]$
- $j$ th element of vector clock at process  $i$ ,  $V_i[j]$ , is  $i$ 's knowledge of latest events at process  $j$


# Assigning Vector Timestamps

- Incrementing vector clocks
  1. On an instruction or send event at process  $i$ , it increments only its  $i$ th element of its vector clock
  2. Each message carries the send-event's vector timestamp  $V_{\text{message}}[1 \dots N]$
  3. On receiving a message at process  $i$ :


$$V_i[i] = V_i[i] + 1$$

$$V_i[j] = \max(V_{\text{message}}[j], V_i[j]) \text{ for } j \neq i$$

# Example




# Vector Timestamps




Initial counters (clocks)




# Vector Timestamps




# Vector Timestamps



# Vector Timestamps



# Vector Timestamps



# Causally-Related ...

- $VT_1 = VT_2$ ,  
*iff* (if and only if)  
 $VT_1[i] = VT_2[i]$ , for all  $i = 1, \dots, N$
- $VT_1 \leq VT_2$ ,  
*iff*  $VT_1[i] \leq VT_2[i]$ , for all  $i = 1, \dots, N$
- Two events are **causally related** *iff*  
 $VT_1 < VT_2$ , i.e.,  
*iff*  $VT_1 \leq VT_2$  &  
there exists  $j$  such that  
 $1 \leq j \leq N$  &  $VT_1[j] < VT_2[j]$

# ... or Not Causally-Related


- Two events  $VT_1$  and  $VT_2$  are **concurrent**

*iff*

$\text{NOT } (VT_1 \leq VT_2) \text{ AND NOT } (VT_2 \leq VT_1)$


We'll denote this as  $VT_2 \parallel VT_1$

# Obeying Causality



- $A \rightarrow B :: (1,0,0) < (2,0,0)$
- $B \rightarrow F :: (2,0,0) < (2,2,1)$
- $A \rightarrow F :: (1,0,0) < (2,2,1)$


# Obeying Causality (2)



- $H \rightarrow G :: (0,0,1) < (2,3,1)$
- $F \rightarrow J :: (2,2,1) < (5,3,3)$
- $H \rightarrow J :: (0,0,1) < (5,3,3)$
- $C \rightarrow J :: (3,0,0) < (5,3,3)$



# Identifying Concurrent Events



- C & F :: (3,0,0) ||| (2,2,1)
- H & C :: (0,0,1) ||| (3,0,0)
- (C, F) and (H, C) are pairs of concurrent events

# Logical Timestamps: Summary

- **Lamport timestamps**
  - Integer clocks assigned to events
  - Obey causality
  - Cannot distinguish concurrent events
- **Vector timestamps**
  - Obey causality
  - By using more space, can also identify concurrent events

# Time and Ordering: Summary

- **Clocks are unsynchronized in an asynchronous distributed system**
- **But need to order events, across processes!**
- **Time synchronization**
  - Cristian's algorithm
  - NTP
  - Berkeley algorithm
  - But error a function of round-trip-time
- **Can avoid time sync altogether by instead assigning logical timestamps to events**

# Reminders (10/1)

- (4 cr students) MP2 due 10/6, Demos on Monday 10/7
  - Signup sheet (soon) on Piazza
- (All) HW2 due 10/10
- No lecture this Thursday 10/3, but video lectures on website (mandatory)
- No Indy office hours this week 10/1 and 10/3 (resumes next week 10/8)